Morgan Stanley | RESEARCH

FOUNDATION

July 6, 2026 09:00 PM GMT

Kingsoft Cloud Holdings | Asia Pacific

Transition to AI Cloud Bearing Fruit; MaaS Brings Further Upside

We like Kingsoft Cloud's early and firm transition from a mid-tier commodity cloud player to AI cloud, with rapidly accelerating AI revenue and improving profitability, anchored by powerful ecosystem support from Xiaomi and Kingsoft Group. We initiate at Overweight.

Key Takeaways

- KC's timely transition and early all-in AI strategy have positioned it well in the AI era among the ten major public cloud players in China.
- It is demonstrating strong pricing power amid the global chip shortage; more importantly, favorable cash flow will ease balance sheet constraints.
- The Xiaomi + Kingsoft ecosystem provides differentiated AI use cases, while revenue contributions from other major customers are also growing rapidly.
- In our base case, we expect KC to follow industry leaders' rising margin trajectory with a stable gap in the next 3-5 years, reaching a low teens EBIT margin.
- Our bull case assumes higher asset yields and margins, supported by KC's rapidly growing MaaS business built on open-source LLMs.

We initiate coverage on Kingsoft Cloud (KC) with an Overweight rating and US $15 price target. KC is a pure and elastic AI cloud play, and is the preferred name in our Greater China IT services & software coverage. Over 2025-28, we forecast a 35% revenue CAGR and a 79% adjusted EBITDA CAGR driven by the AI cloud business (AI contributed 34% of revenue in 4Q25, and we forecast this to rise to >40% in 2026 and >60% by 2028). The adjusted OP margin upside (6.9% in 2028 vs. -1.6% in 2025) is backed by favorable unit economics of GPU cloud and potentially MaaS.

Valuation: We derive our price target from a 5.5x 2027e EV/EBITDA multiple, at a discount to the US neocloud peer median of 10x to reflect greater uncertainty around supply chain availability. We prefer EBITDA for comparability across companies given differences in depreciation policies and limited EBIT.

Key downside risks include supply-side restrictions leading to procurement shortfall, higher-than-expected interest rate leading to higher cost of funding, and slower-than-expected AI model developments in China.

Key catalysts include a longer-than-expected useful life for servers, pick-up in MaaS contribution and breakthroughs from Xiaomi's AI ecosystem.

MORGAN STANLEY ASIA LIMITED+

Yang Liu
Equity Analyst
Yang.Liu@morganstanley.com +852 2239-1911

Tom Tang
Equity Analyst
Tom.Tang@morganstanley.com +852 3963-1860

Andy Meng, CFA
Equity Analyst
Andy.Meng@morganstanley.com +852 2239-7689

Gary Yu
Equity Analyst
Gary.Yu@morganstanley.com +852 2848-6918

Kingsoft Cloud Holdings (KC.O, KC US)

Greater China IT Services and Software | China

Stock Rating	**Overweight**
Industry View	**In-Line**
Price target	US$15.00
Up/downside to price target (%)	65
Shr price, close (Jul 2, 2026)	US$9.10
52-Week Range	US$18.52-8.35
Sh out, dil, curr (mn)	299
Mkt cap, curr (mn)	Rmb18,425.7
EV, curr (mn)	Rmb23,417.4
Avg daily trading value (mn)	US$22

Fiscal Year Ending	12/25	12/26e	12/27e	12/28e
EPS (Rmb)**	(2.01)	(1.88)	0.13	2.74
EPS (Rmb)§	(2.55)	(1.55)	(0.88)	(0.36)
Revenue, net (Rmb mn)	9,559	13,094	17,902	23,418
EBITDA (Rmb mn)	1,707	4,673	8,615	12,950
ModelWare net inc (Rmb mn)	(936)	(876)	(285)	486
EV/rev	2.8	2.6	2.3	2.0
EV/EBITDA	15.6	7.2	4.8	3.7
P/E	NM	NM	NM	37.9
P/BV	2.3	2.2	2.3	2.1

Unless otherwise noted, all metrics are based on Morgan Stanley ModelWare framework
** = Based on consensus methodology
§ = Consensus data is provided by Refinitiv Estimates
e = Morgan Stanley Research estimates

Morgan Stanley does and seeks to do business with companies covered in Morgan Stanley Research. As a result, investors should be aware that the firm may have a conflict of interest that could affect the objectivity of Morgan Stanley Research. Investors should consider Morgan Stanley Research as only a single factor in making their investment decision.

For analyst certification and other important disclosures, refer to the Disclosure Section, located at the end of this report.

+= Analysts employed by non-U.S. affiliates are not registered with FINRA, may not be associated persons of the member and may not be subject to FINRA restrictions on communications with a subject company, public appearances and trading securities held by a research analyst account.

Financial Summary

Exhibit 1: Kingsoft Cloud: Financial summary

INCOME STATEMENT

Years Ending December 31 Rmb mn	2025	2026E	2027E	2028E
Public cloud service	6,633	9,935	14,585	19,968
Enterprise cloud service	2,925	3,159	3,317	3,450
Total consolidated revenue	**9,559**	**13,094**	**17,902**	**23,418**
Cost of revenue	8,055	11,122	14,917	19,194
Gross profit	**1,503**	**1,972**	**2,985**	**4,224**
S&M expense	551	584	625	669
G&A expense	915	869	895	922
R&D expenses	810	827	851	885
Total operating expenses	2,276	2,280	2,372	2,476
Adjusted EBITDA	**2,336**	**5,035**	**8,987**	**13,333**
Non-opearting items	(175)	(575)	(905)	(1,185)
Profit before tax	(948)	(883)	(292)	563
Net profit	(944)	(883)	(292)	479
Net loss to ordinary shareholders	**(936)**	**(876)**	**(285)**	**486**
Diluted EPS (reported)	(3.45)	(2.96)	(0.98)	1.60
MW EPS	**(2.01)**	**(1.88)**	**0.13**	**2.74**

BALANCE SHEET

Years Ending December 31 Rmb mn	2025	2026E	2027E	2028E
Cash	6,117	3,482	1,859	1,343
Accounts receivable	1,740	2,210	2,801	3,384
Current Assets	11,023	9,586	9,354	10,190
Prepayments and other assets	140	193	259	333
Other assets	98	98	98	98
Non-current asset	15,706	26,028	34,867	42,191
Total assets	**26,729**	**35,614**	**44,221**	**52,381**
S/T borrowings	3,348	4,348	5,348	5,848
Account payables and notes payables	2,014	2,681	3,461	4,281
Other ST liabilities	114	114	114	114
Total current liabilities	9,421	12,776	16,064	18,857
LT debt	3,024	3,024	3,024	3,024
Other LT liabilities	2,759	6,024	8,930	11,597
Total non-current liabilities	7,995	14,407	20,019	24,907
Total liabilities	**17,416**	**27,184**	**36,083**	**43,764**
Common shares	(0)	(0)	(0)	(0)
Additional paid-in capital	24,073	24,073	24,073	24,073
Total shareholders' equity	**9,313**	**8,430**	**8,138**	**8,617**
Total liabilities and equity	**26,729**	**35,614**	**44,221**	**52,381**

CASH FLOW STATEMENT

Years Ending December 31 Rmb mn	2025	2026E	2027E	2028E
Net profit	(944)	(883)	(292)	479
Depreciation and amortization	2,480	4,981	8,001	11,202
Other non-cash adjustments	57	344	585	803
Net change in working capital	1,761	2,682	2,096	1,729
Operating cash flow	**3,801**	**7,436**	**10,713**	**14,545**
Capex	(4,742)	(9,150)	(10,065)	(11,071)
Purchases of intangible assets	-	-	-	-
Net change in investment	-	-	-	-
Other investing cash flow	213	-	-	-
Investing cash flow	**(4,530)**	**(9,150)**	**(10,065)**	**(11,071)**
Issuance/(repayment) of debts	2,617	1,000	1,000	500
Net change in equity	31	-	-	-
Other financing cash flow	2,355	-	-	-
Financing cash flow	**4,183**	**(921)**	**(2,271)**	**(3,990)**
Other adjustments	(67)	-	-	-
Net change in cash	3,387	(2,635)	(1,623)	(517)
Beginning cash balance	2,730	6,117	3,482	1,859
Ending cash balance	6,117	3,482	1,859	1,343

KPI & FINANCIAL RATIOS

Years Ending December 31 Rmb mn	2025	2026E	2027E	2028E
Margin				
Gross margin	16%	15%	17%	18%
EBITDA margin	24%	38%	50%	57%
Operating margin	-8%	-2%	3%	7%
Net margin	-10%	-7%	-2%	2%
Normalize net margin	-10%	-7%	-2%	2%
Growth % YoY				
Public cloud service	32%	50%	47%	37%
Enterprise cloud service	5%	8%	5%	4%
Revenue	23%	37%	37%	31%
Operating profit	-6%	-60%	-299%	185%
Net income	-52%	-6%	-67%	-264%
Normalized net income	-52%	-6%	-67%	-271%
KPI				
IDC as % of revenue	69.4%	75.9%	81.5%	85.3%
Other VAS as % of revenue	0.0%	0.0%	0.0%	0.0%
Capex	4,742	9,150	10,065	11,071
FCF	(1,680)	(2,413)	(302)	2,295

Source: Company data, Morgan Stanley Research estimates.

Risk Reward – Kingsoft Cloud Holdings (KC.O)

Benefiting from computing power shortage

PRICE TARGET US$15.00

Base case, derived from 5.5x 2027e EV/EBITDA multiple. This multiple is at a discount to the US neocloud peer average of 13x and median of 10x to reflect lower asset yields and greater uncertainty around supply chain availability. We prefer EBITDA for comparability across companies given differences in depreciation policies and limited EBIT.

Consensus Price Target Distribution

US$15.00 (MS PT) — US$20.45 (Mean) — US$25.02

◆ Mean ◆ Morgan Stanley Estimates

Source: Refinitiv, Morgan Stanley Research

RISK REWARD CHART

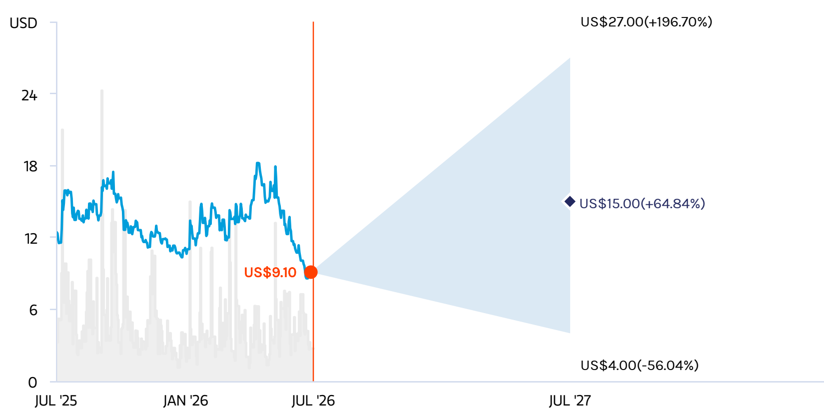

Key: — Historical Stock Performance ● Current Stock Price ◆ Price Target

Source: Refinitiv, Morgan Stanley Research

OVERWEIGHT THESIS

- We forecast AI revenue to grow 109% in 2026, driven by strong computing power demand and total server investments of Rmb15bn.
- We expect gross margin to expand on the back of its more profitable AI business.
- Our PT implies 5.5x 2027e EV/EBITDA, at a discount to US neocloud peers to reflect greater uncertainty around supply chain availability.

Consensus Rating Distribution

●	100%	Overweight
	0%	Equal-weight
	0%	Underweight

● MS Rating

Source: Refinitiv, Morgan Stanley Research

Risk Reward Themes

New Data Era: *Positive*
Secular Growth: *Positive*

View descriptions of Risk Rewards Themes here

BULL CASE US$27.00

7.0x 2027e EV/EBITDA

We assume stronger demand for high-end GPU driving faster EBTIDA growth and positive development from the Mimo models.

BASE CASE US$15.00

5.5x 2027e EV/EBITDA

We expect KC to continue to benefit from high-end GPU demand by the hyperscalers.

BEAR CASE US$4.00

4x 2027e EV/EBITDA

We assume a slowdown in Chinese AI development.

Risk Reward – Kingsoft Cloud Holdings (KC.O)

KEY EARNINGS INPUTS

Drivers	NA 2025	NA 2026e	NA 2027e	NA 2028e
Public cloud revenue (Rmb)	6,633	9,935	14,585	19,968
Enterprise cloud revenue (Rmb)	2,925	3,159	3,317	3,450
EBITDA margin (%)	24.4	38.5	50.2	56.9

INVESTMENT DRIVERS

- AI revenue growth
- Gross margins

GLOBAL REVENUE EXPOSURE

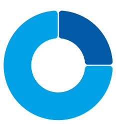

- 20-30% APAC, ex Japan, Mainland China and India
- 70-80% Mainland China

Source: Morgan Stanley Research Estimate
View explanation of regional hierarchies here

RISKS TO PT/RATING

RISKS TO UPSIDE

- Stronger-than-expected AI GPU demand
- Lower-than-expected interest rate driving better return profile
- Better-than-expected Xiaomi's AI model development

RISKS TO DOWNSIDE

- Supply-side restrictions leading to procurement shortfall
- Higher-than-expected interest rates leading to higher cost of funding
- Slower-than-expected AI model developments in China

OWNERSHIP POSITIONING

HF Sector Long/Short Ratio	2.1%
HF Sector Net Exposure	29.5x

Source: Morgan Stanley Prime Brokerage. Includes certain hedge fund exposures held with MSPB. Information may be inconsistent with or may not reflect broader market trends. Long/Short Ratio = Long Exposure / Short exposure. Sector % of Total Net Exposure = (For a particular sector: Long Exposure - Short Exposure) / (Across all sectors: Long Exposure – Short Exposure).

MS ESTIMATES VS. CONSENSUS

FY 2026e

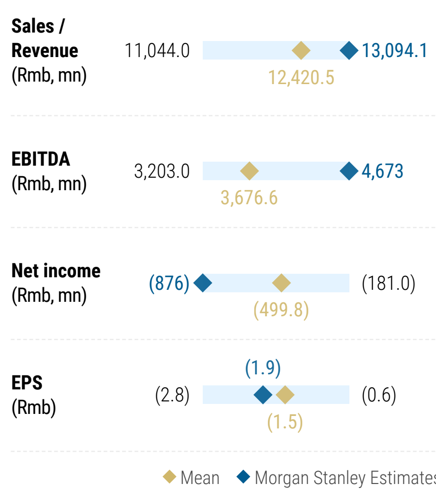

Source: Refinitiv, Morgan Stanley Research

Executive Summary

Kingsoft Cloud's early and firm transition to neocloud is helping it capture upside in the AI era. Previously a mid-tier commodity cloud provider, KC started the transition in 2022 by massively scaling down the low-margin CDN business, merging it into Xiaomi's ecosystem and going all-in on AI-related business and the technology stack despite limited R&D and financial resources in 2024. KC successfully tapped into fast-rising training demand from hyperscalers and AI labs. These initiatives have borne fruit, with three years of revenue growth acceleration, the EBITDA margin trending up and operating profit turning around in 2025. AI cloud contributed 49% of public cloud revenue in 4Q25 and is set to further increase to 61% in 2026 and 78% in 2028 based on our estimates. We believe sizable upside from AI inference will follow soon for KC, especially from the MaaS business built on open-source LLMs, given KC's strong model optimization capabilities.

Xiaomi + Kingsoft ecosystem. As shareholder and ecosystem partners, Xiaomi and Kingsoft Group jointly provide various smart and general compute use cases, covering smartphones, AIoT home appliances, electric vehicles, office software, and games. Also, Xiaomi entered the AI model market with the debut of MiMo and has committed to high AI investment on a sustained basis in 2026. This directly led to 27-47% upward revisions to the connected transaction volume cap between Xiaomi and KC in 2026-27, paving the way for KC's faster revenue growth, though KC management indicated there would be even faster growth from external customers.

Margin upside and improved cash flow arrangement. Chinese public cloud players demonstrated good pricing power during the recent round of chip shortages by hiking cloud service prices for the first time in decades. In some scenarios (equipment in bad shortage), the price hike more than covers the cost inflation. In our base case, we assume KC's OPM will follow industry leader AliCloud's trend and keeps their current gap in the next three years, reaching low teens by 2028. In our bull case, we assume KC to leverage open-source LLMs to offer MaaS service at scale, which should help to narrow its margin gap. Specifically for KC, it used its advantage to get a much better cash flow arrangement from a customer with a sizable prepayment. This can ease balance sheet constraints for this capital-intensive business, and raise the server lifecycle unleveraged IRR by 5ppt, assuming the same price.

We think KC's recent share price weakness (down 23% for the month of June), alongside declines in the Heng Sang Tech Index (-8%) and Xiaomi's share price (-23%) offers an attractive entry point. The market has largely treated KC as a higher-beta proxy for Xiaomi, as revenue from the Xiaomi ecosystem increased from 14% in 2022 to 31% in 1Q26. However, we expect this correlation to moderate over time as the industry enters a period of constrained computing power supply, where access to high-quality capacity becomes more important to KC's growth than demand from any single large customer. In addition, we believe KC's role within the Xiaomi ecosystem, combined with its pricing power, should limit the impact of volume and margin pressure in Xiaomi's consumer electronics business.

Investment positives

- Strong demand that guarantees negotiation power vs the customers and a decent IRR on project basis
- Prepayment and finance lease improves the capital structure
- China's low interest rate and loose credit environment supports overall cash flow

Investment concerns

- Supply chain constraints may create uncertainty around new revenue generation
- Still-high capital intensity nature could introduce some concerns over financing
- Legacy cloud businesses such as Content Delivery Network (CDN) could dilute overall growth

Key downside risks

- If KC meets server supply chain disruptions, growth could face downside risk due to lack of capacity. On the other hand, if the supply/availability of next-generation servers increase significantly or hit the market much earlier than expected, rental pricing for KC's legacy servers could come under significant pressure.
- If AI development in China slows due to weaker token consumption or reduced capex from Chinese hyperscalers, cloud players could face softer demand.
- If leading LLMs in China choose to go closed-source in the AI model updates (currently open-sourced), this would limit KC's future margin upside brought by the MaaS business.
- If key players in AI LLM decide to exit the market due to downstream consolidation, this could negatively affect overall AI training demand, at least in the near term.

Key upside risks

- If KC's servers remain operational beyond their four- or five-year depreciable life, margins could see meaningful upside following the depreciation cycle. We estimate that an additional year of service life, assuming flat pricing, could increase equipment IRR by 5ppt.
- If Xiaomi's AI ecosystem develops ahead of market expectations, particularly with respect to MiMo's capabilities, it could bring additional upside to KC's business volume.
- Given that hyperscalers directly utilize AI capacity from neocloud providers, any upward revision to their AI investment (opex) should support demand for KC's services.

See more reports in our series on China's AI path

Data Centers: Inflection and Disruption in Remote Areas (21 Jan 2026)

China's AI Path: Owning the Full AI Stack via In-house Chips (11 Mar 2026)

China's AI Path: Monetizing Surging Token Use via AI Cloud (16 Mar 2026)

China's AI Path: More Bang For The Buck (27 Apr 2026)

Benefitting From A Surging Cloud Market

AI Cloud

We define AI cloud as IaaS + MaaS. We exclude AI SaaS from our market sizing given the fragmented nature of the application layer and the difficulty of sizing it accurately. Moreover, when we refer to cloud players, we generally mean infrastructure providers, which supply the underlying capabilities to application developers. That said, we recognize that AI applications could represent a much larger TAM than the infrastructure layer over time, driven by the development in agentic and AI-native applications.

According to IDC, the GenAI IaaS + MaaS layer is expected to expand from Rmb15bn in 2024 to Rmb218bn by the end of 2029, representing a CAGR of 72%. As a percentage of IaaS + PaaS, GenAI is expected to rise from 6% in 2024 to 39% by the end of 2029.

Our survey of industry CIOs (link) indicates a similar trend, with the share of public cloud spending related to AI expected to rise from 6% in 1H26 to 10% over the next three years.

Exhibit 2: How the AI era differs from traditional cloud

	Cloud era	AI era
Datacenter	CPU based Tier 1 cities	GPU based Remote area
IaaS	Computing Storage Networking	Model training Model inferencing
PaaS/MaaS	Database Middleware Operating system	Vector database MaaS
SaaS	ERP Office software CRM	Chatbot Coding agent

Source: Company data, Morgan Stanley Research.

Exhibit 3: AI to take a rising share of IaaS + PaaS in China

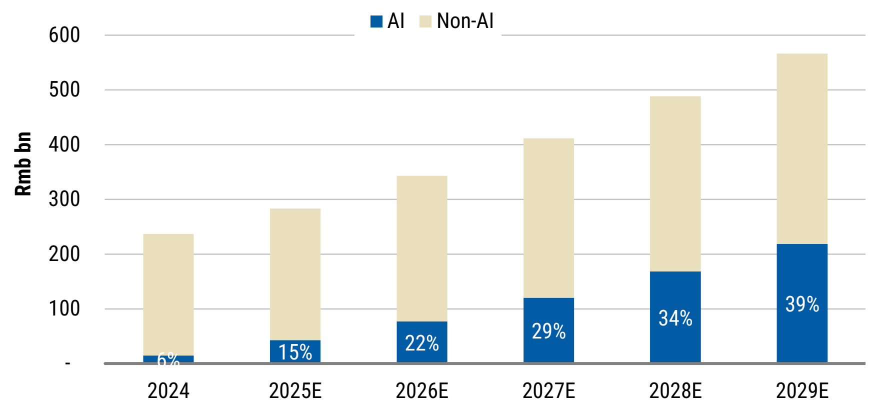

Source: IDC estimates, Morgan Stanley Research.

Exhibit 4: Public cloud spending related to AI (mix by CIOs)

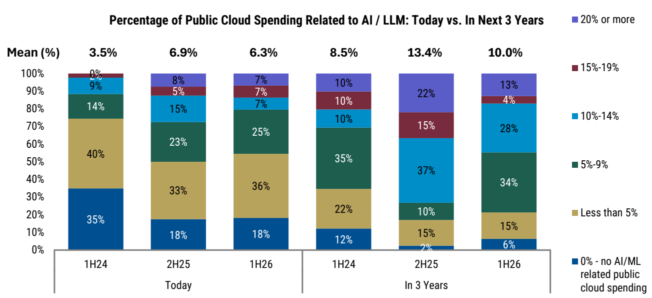

Source: AlphaWise, Morgan Stanley Research.

Exhibit 5: Public cloud spending related to AI (weighted average)

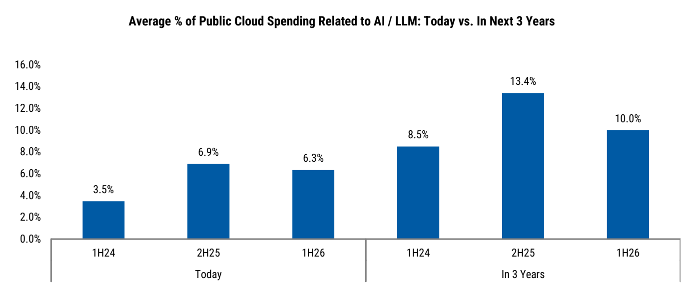

Source: AlphaWise, Morgan Stanley Research.

AIaaS – GenAI is becoming a material contributor to growth: According to IDC, GenAI IaaS revenue in China reached Rmb16.7bn in 1H25 (up 220% YoY), accounting for 17% of the total domestic IaaS market (which grew 17% YoY). In terms of total growth contribution, GenAI contributed 13ppt, vs. 3.7ppt for other IaaS revenue, demonstrating that GenAI is becoming the most important growth driver of the IaaS market in China. Infrastructure is the layer showing the most visible growth, which is in line with the early stage of the adoption cycle – but we think that with future application innovations, the AI SaaS layer should catch up and drive a surge in token usage.

AI PaaS/MaaS: MaaS has been a special layer between the infrastructure and platform layers during the AI era. MaaS provides a range of models to users through API, and the hyperscalers can charge based on token usage. Major cloud players have launched their respective MaaS products, including Baidu's Qianfan, ByteDance's Volcanic Ark, and Alibaba's Bailian platform. There are also some emerging players such as SiliconFlow that are gaining more market share within the segment.

AI SaaS is still at an early stage, and the market is waiting for the emergence of a killer app. Most current applications are in the form of chatbots, agents, and video generation tools. We are encouraged by the recent emergence of OpenClaw and its growing popularity among both developers and consumers. While we believe AI SaaS/agents could ultimately represent a significantly larger TAM than the infrastructure layer, the market remains at an early stage, making it difficult to measure at present.

Exhibit 6: IaaS market breakdown

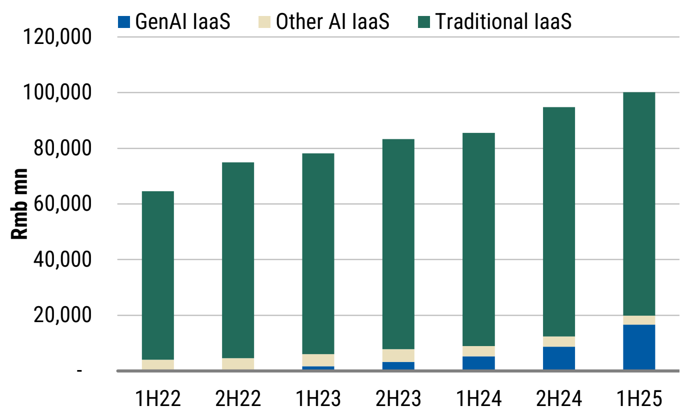

Source: IDC, Morgan Stanley Research.

Exhibit 7: Growth contribution: GenAI is the major growth driver of China's IaaS market

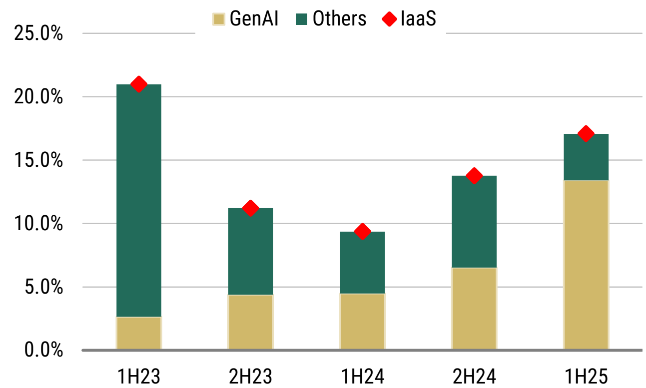

Source: IDC, Morgan Stanley Research.

Key demand driver – rising token usage

Inference is the primary demand driver in China: We expect model inference demand to be the principal driver of GenAI IaaS adoption going forward. According to IDC, training workloads accounted for 76% of the GenAI IaaS market in 2024, but this share is expected to decline to 23% by the end of 2029 as inference workloads grow at a significantly faster pace. IDC forecasts training demand to expand at a 26% CAGR over 2024-29E, vs. 103% for inference demand. We think strong inference demand is driven by:

1. More advanced and accessible LLMs;
2. Emergence of AI applications going forward and the integration of GenAI into daily workloads; and
3. Improvements in multi-modal models, which could require even more tokens compared to the standard text-to-text LLMs.

We think the key verticals driving token growth and AI adoption will primarily include internet companies, AI-native developers, and the EV/auto industry.

Token usage continues to rise...: A token is the underlying unit that drives AI demand. We note that token usage in China has been rising, led by ByteDance's Doubao (an AI-powered chatbot), where daily average token use rose from 4trn in December 2024 to 50trn by the end of 2025, according to ByteDance. Furthermore, key AI apps in China have seen rapid growth in time spent after the market volatility that followed the release of DeepSeek in February 2025. The rising adoption of OpenClaw is another example of how new use cases can lead to a non-linear increase in token usage. MiniMax noted that its ARR (annualized recurring revenue) increased 50% from US$100mn in December 2025 to US $150mn in February 2026, while token usage increased 6x, driven by OpenClaw. That said, we note that >70% of its revenue comes from overseas markets.

In general, we think several factors support the surge in token usage in China:

1. 2C application usage (Qwen, Quark, Yuanbao, Doubao, etc.);
2. Broader 2B adoption (such as coding agent); and
3. Emerging new use cases such as the recent OpenClaw.

…which could drive higher public cloud adoption vs. the last cycle: We previously noted that a primary reason for the smaller public cloud market in China relative to the US is the preference by enterprises for private cloud or on-premise deployment in China. Our latest CIO survey suggests early signs of change: the intention to adopt public cloud rose from 44% in the previous survey to 54%, while the intention in three years time also rose from 19% to 35%.

Why use AI cloud? In terms of downstream users, model providers account for the largest share (23%) of GenAI IaaS usage followed by industry, such as auto and finance. Some reasons for companies to use IaaS rather than buy their own infrastructure or adopt private cloud include:

- **More elasticity:** Customers can have more flexibility in terms of workload usage and do not need to overcommit on hardware purchases ahead of when they actually generate revenue, which could enhance their payback period. Large purchases of GPU servers can be more burdensome for companies (other than hyperscalers) because of the large amount of capex required.
- **Chipset constraints:** Given the tight supply of chipsets, especially in China, it makes sense for hyperscalers to procure AI chips (better prices through bulk purchases and also more access to the supply chain) and for smaller companies to access computing power through cloud services.
- **Rapid iteration of underlying technology:** Chipsets have shown huge performance improvements year on year and the useful life of a chip is usually only 4-5 years. Thus, it does not make economic sense for smaller enterprises to purchase AI servers.

- **Better model bundling:** Top hyperscalers come with their own frontier models, which could also provide better service and deployment if the enterprises are interested in adopting those models instead of hiring talent to adopt and fine-tune them.

Reasons not to use it – data security is the biggest pushback: We believe the business community will eventually recognize that adopting AI cloud may require sharing proprietary data with cloud vendors, potentially giving those vendors some visibility into internal data processing activities. As a result, organizations may choose to migrate certain workloads to private cloud or on-premise environments when the value of safeguarding proprietary data exceeds the cost benefits of public cloud deployment. This trend could hamper the long-term TAM of AI cloud. Our CIO survey also suggests that a hybrid deployment model may be the optimal long-term approach.

Exhibit 8: GenAI IaaS by workload type

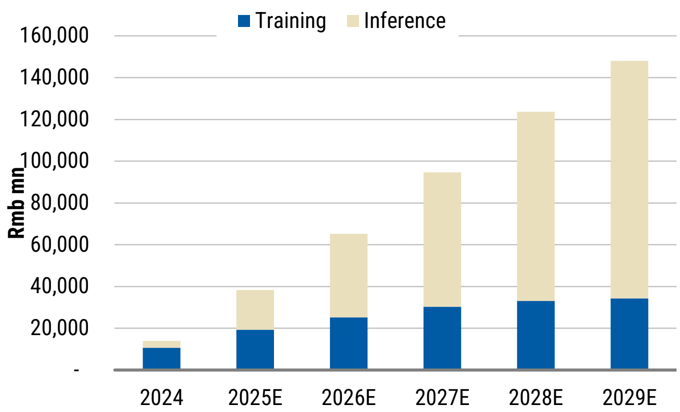

Source: IDC, Morgan Stanley Research estimates.

Exhibit 9: GenAI IaaS by downstream user

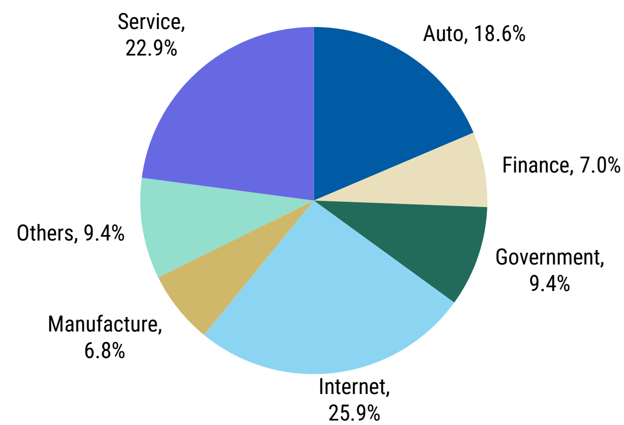

Source: IDC, Morgan Stanley Research.

Exhibit 10: Vendors to assist AI deployment, by vendor type

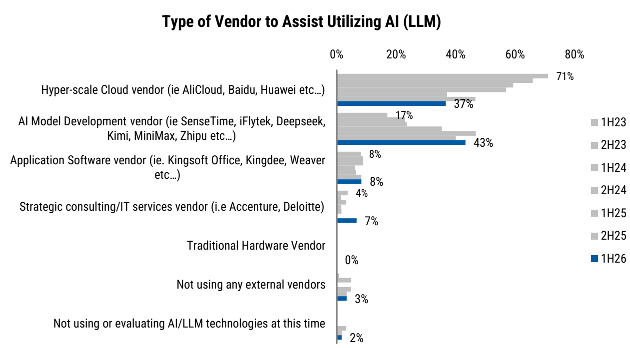

Source: AlphaWise, Morgan Stanley Research.

Exhibit 11: Vendors to assist AI deployment, by vendor name

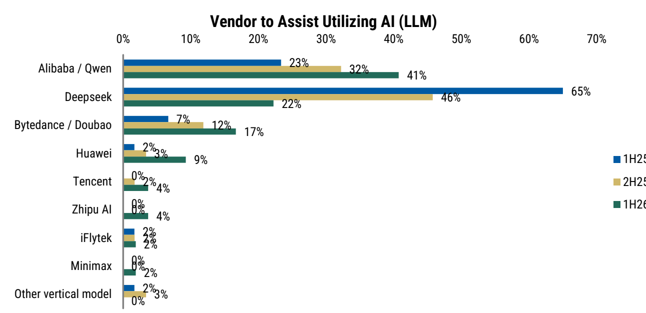

Source: AlphaWise, Morgan Stanley Research.

Exhibit 12:
Public deployment to pick up over the next 12 months

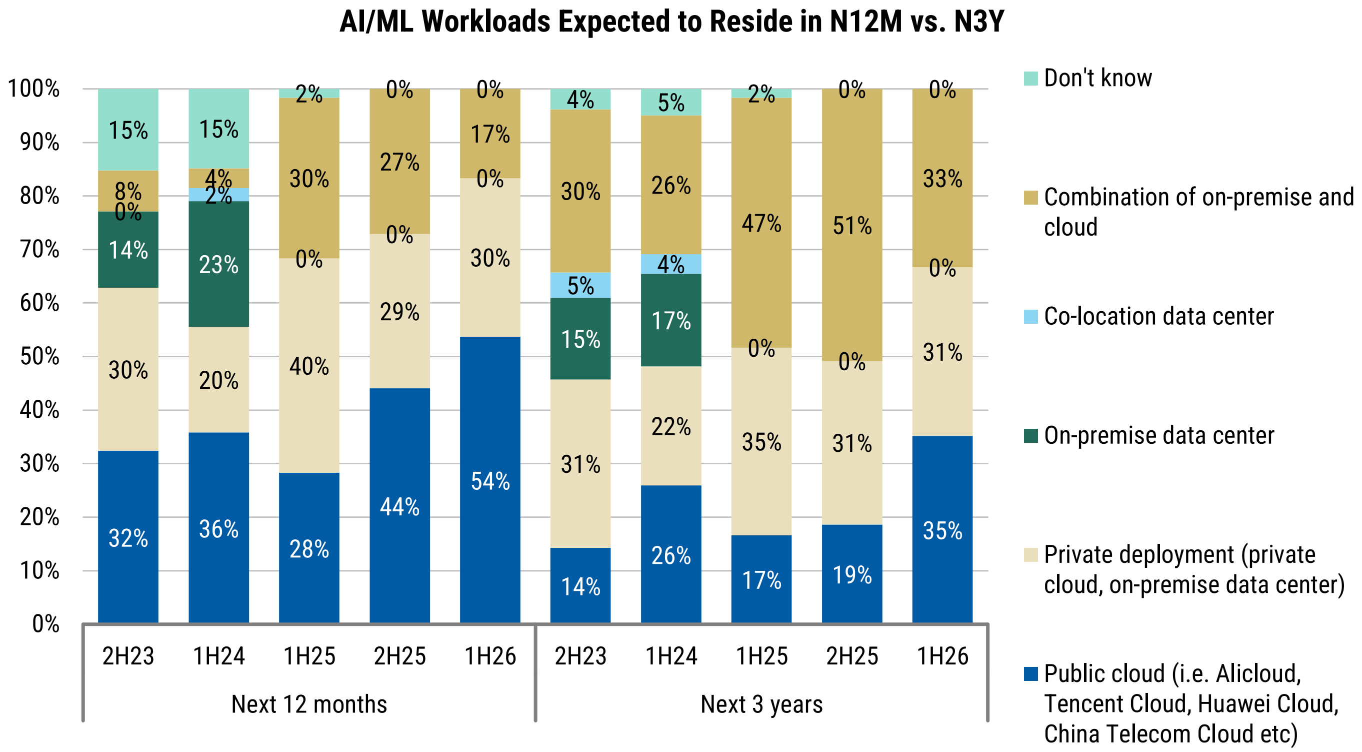

Source: AlphaWise, Morgan Stanley Research.

Exhibit 13: Key AI apps: Time spent still growing rapidly

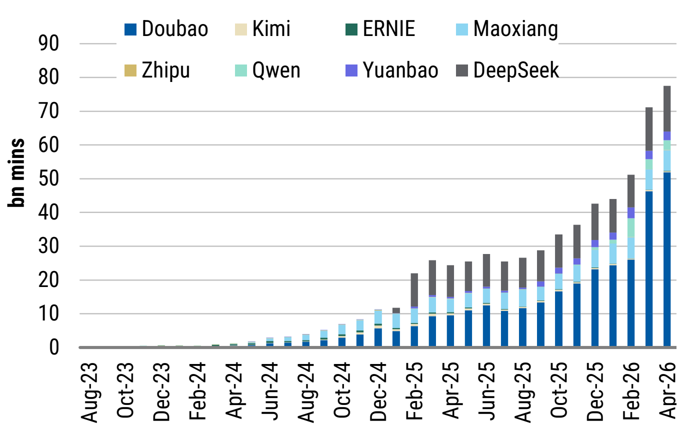

Source: QuestMobile, Morgan Stanley Research.

Exhibit 14: Major data points of token usage in China and the US

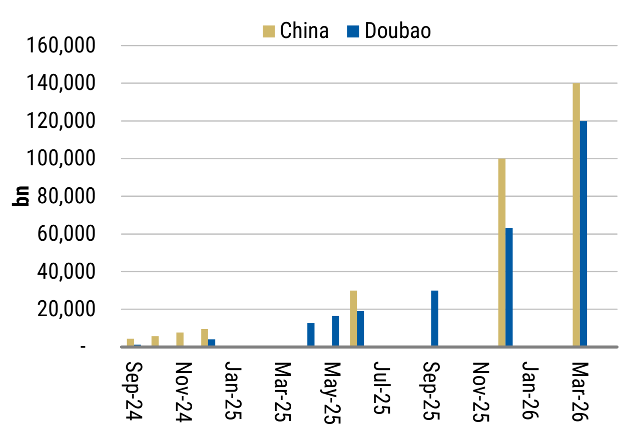

Source: Company data, Morgan Stanley Research.

Supply side – chips remain the biggest constraint

Availability of chipsets is the key constraint: We note that AI demand in China is still significantly capped by the availability of AI chips. We expect this constraint to ease gradually from 2H26 onwards with the expansion in domestic AI wafers as well as the potential re-entrance of Nvidia H200 and equivalent chips for training demand. Our Greater China Semiconductors team expects domestic AI capacity to almost double in 2026-27, which could support the development of domestic AI technologies. Our talks with data center vendors highlighted that hyperscalers are expecting to ramp up their domestic AI capacity from 2H26 when they receive chip shipments.

Growth is capex-driven, in our view: In the early stage of the cycle, we believe growth is mostly supply-driven and hence capex could be the leading indicator of cloud revenue growth. Based on company reported numbers, hyperscaler capex growth in China bottomed in 1Q23 and started to accelerate rapidly in 2024 thanks to strong AI investment. As a result, public cloud revenue growth bottomed in late 2023 and started to pick up in 4Q24, mostly led by Alibaba's cloud revenue growth, which accelerated from single digits in 2024 to 34% in calendar 3Q25. We think ongoing investment in AI will further support cloud growth, and cloud growth could continue to accelerate given more efficient use of infrastructure and the adoption of more powerful chipsets in the future. We forecast hyperscaler capex to grow 41% YoY, to Rmb602bn or US$86bn, which is 11-12% of the US hyperscaler level. We forecast Alibaba's and Tencent's capex intensity to increase to about 15%, which still lags Google and Meta at 40-50% and Amazon and Microsoft at around 20%.

Optimization at the infrastructure level drives higher efficiency: Although China faces greater constraints at the chipset level, infrastructure providers, including model developers and cloud providers, continue to improve AI chipset utilization efficiency. These optimization efforts can help providers support higher workload/usage despite limited chip supply. Using MiniMax as an example, in its prospectus the company noted that it could achieve MFU (model flop utilization) of >75%, vs. the industry average of 40-50%, enabling lower inference costs and enhanced performance.

Exhibit 15: Cloud capex vs. growth

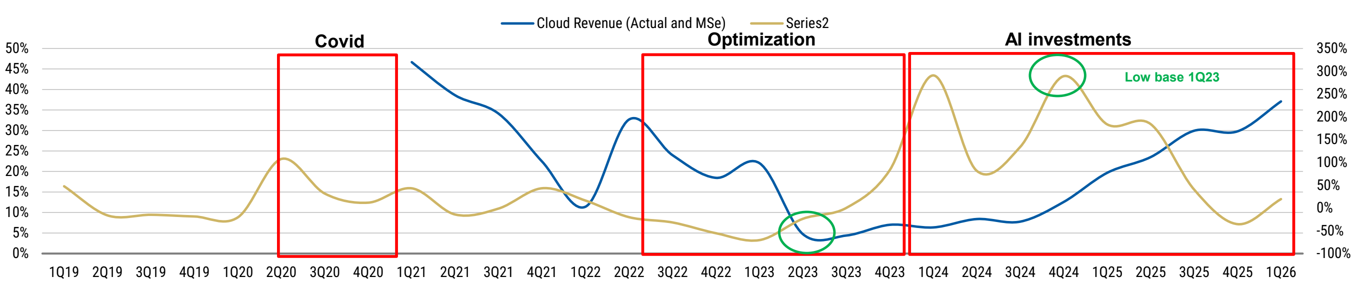

Source: Company data, Morgan Stanley Research estimates.

Exhibit 16: We forecast hyperscaler capex to reach Rmb602bn in 2026

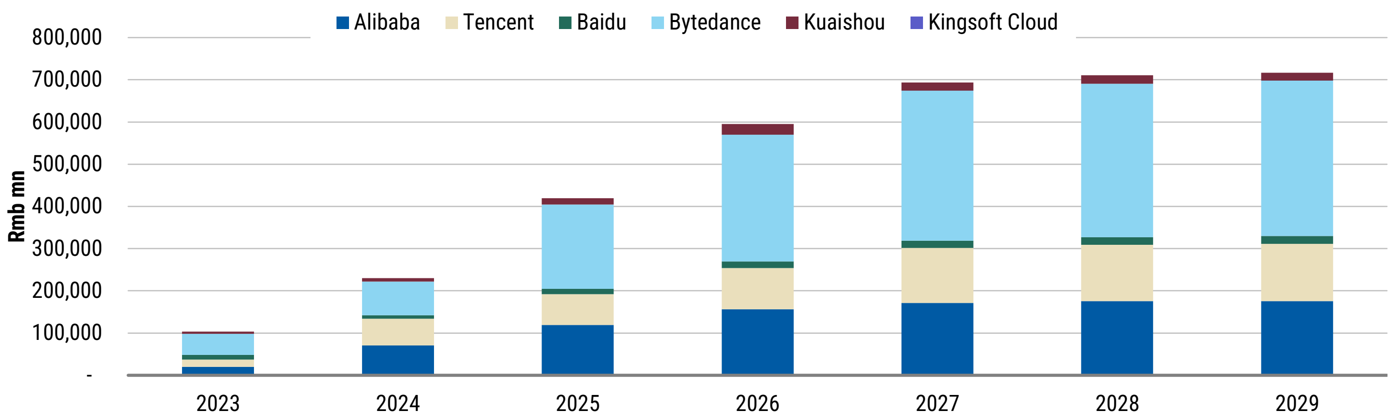

Source: Company data, Morgan Stanley Research estimates.

Exhibit 17: China – server spending and shipment growth

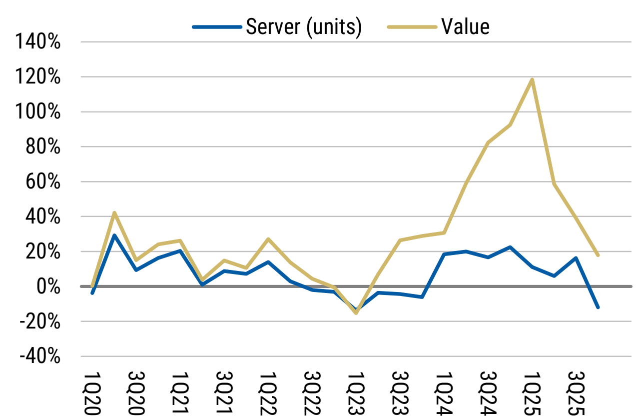

Source: IDC, Morgan Stanley Research.

Exhibit 18: Our semiconductors team forecasts SMIC's AI capacity to double in 2026

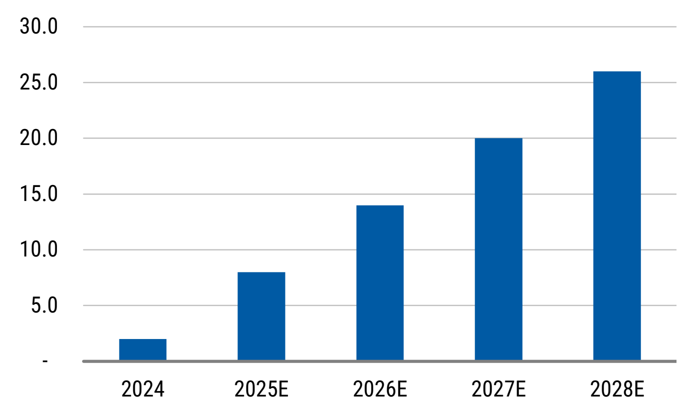

Source: Company data, Morgan Stanley Research (E) estimates.

Kingsoft Cloud's Transition to Neocloud

Foundation of neoclouds

A neocloud is a specialized cloud provider that focuses on delivering infrastructure and services specifically designed for AI workloads. More specifically, a neocloud is a type of cloud company that operates data centers built for AI workloads, offering faster or easier access to high-performance GPUs used to train and run AI to companies seeking this specialized infrastructure. Unlike traditional hyperscale cloud providers (like AWS, Azure, Google Cloud, AliCloud) that offer a broad range of general-purpose services, neoclouds are purpose-built for the demanding requirements of AI — particularly those involving GPUs. Global leading neocloud players are CoreWeave and Nebius.

The emergence of neoclouds addresses the growing demand for AI-specific hardware and the challenges enterprises face in accessing and affording the necessary resources. Hyperscalers and tech firms are usually the core customers of neocloud.

Key characteristics:

- **AI-First infrastructure:** Neoclouds design their entire infrastructure — compute, storage, and networking — to optimize for AI-specific tasks.
- **GPU-as-a-Service:** The core offering is GPUaaS, providing access to powerful, often latest-generation GPUs on demand. Customers rent GPUs from neoclouds to train AI models or run AI inference.
- **Capex to opex shift:** Companies can rent GPU compute capacity instead of acquiring expensive AI hardware themselves, shifting from capital expenditure to opex.
- **Optimized for AI workloads:** They excel at providing computational power for tasks like training and fine-tuning LLMs, complex analytics, and real-time simulations.
- **High-performance networking:** Neoclouds prioritize high-speed interconnects and low-latency networking between GPU clusters to ensure efficient communication and distributed training.
- **Specialized services:** Beyond GPUs, neoclouds often provide AI-optimized storage, data pipeline support, AI model training and fine-tuning services, and monitoring tools.
- **Flexible pricing:** Neoclouds typically employ pay-as-you-go pricing models, often with transparent pricing ranging from hourly price to GPU life-cycle price, eliminating the need for large upfront investment.
- **Agility:** Neoclouds are often leaner and more agile than larger cloud providers, quickly adopting the latest technology.

The core value proposition offered by neocloud companies include:

- Accessing the computing power and other related infrastructure/equipment/component including data centers, memory, switch, networking. In China, chipset sourcing is way much more important than others, while in US, the data center and related power sourcing are big hurdles.

- Software stack capabilities including cluster orchestration, model training and fine-tuning, inference optimization, developer tooling. These are more important to serve startups and enterprises, relative less important when serving hyperscalers.
- Capital support.

Kingsoft Cloud's transition from traditional cloud to neocloud

Kingsoft Cloud has undergone a significant strategic and competitive repositioning as the Chinese cloud industry has evolved from the traditional cloud computing era to an AI-driven era.

1. Market Position in the Cloud Era (before 2024)

In the traditional cloud era, Kingsoft Cloud was a mid-tier player in China's public cloud market. It showed strong early growth, with revenues increasing by 51.3% annually between 2019 and 2021. It was a relatively independent player compared with cloud peers such as AliCloud, Huawei Cloud, Tencent Cloud, and SOE telcos. We estimate Kingsoft Cloud's market share in China's public cloud IaaS/PaaS market was 3-4% in 2025. Rather than competing directly with hyperscalers across all segments, Kingsoft Cloud has focused on specific industry verticals, particularly gaming and video, where it holds a top-three position in cloud services. The company also relied heavily on Content Delivery Network (CDN) services, which contributed as high as half of its revenue. However, CDN is a commoditized low-margin business that has weighed on profitability.

We note the IaaS value chain did not face meaningful shortages across most resources/equipment/components during that period. As a result, cloud providers primarily relied on unit cost advantages and a more favorable customer mix, which enabled greater resource utilization and reuse, to drive profitability.

Overall, Kingsoft Cloud faced several challenges, including price competition, limited ecosystem support, cost disadvantage stemming from its relative small scale vs. peers, and balance sheet constraints. The company has remained loss-making and cash-burning throughout the cloud cycle, leading investors to question the sustainability of its business model.

Exhibit 19: Public cloud vendors in China (logos)

Source: Morgan Stanley Research.

Exhibit 20: Kingsoft Cloud revenue and net loss, 2017-25

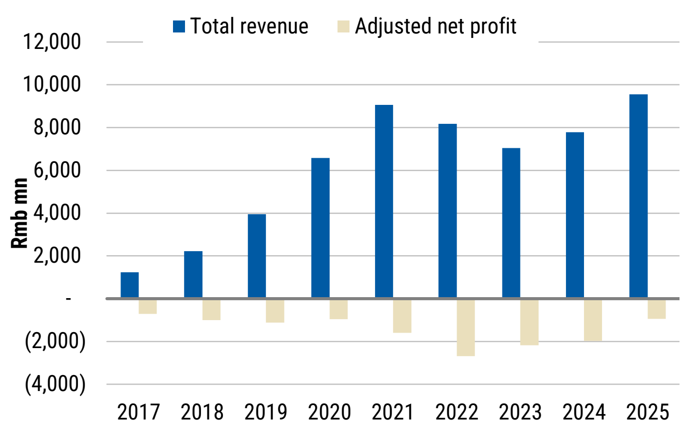

Source: Company data, Morgan Stanley Research estimates.

Exhibit 21: Kingsoft Cloud GPM, adjusted EBITDA margin, adjusted OPM, 2017-25

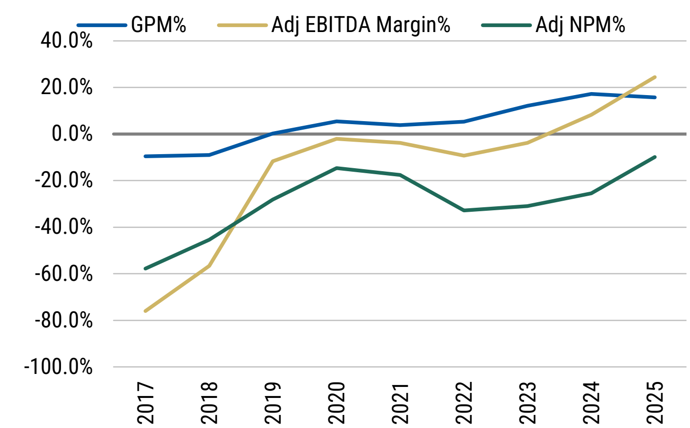

Source: Company data, Morgan Stanley Research estimates.

2. Market Position in the AI Era (after 2024)

Strategic Transformation

Kingsoft Cloud has undergone a two-step strategic transformation to reposition itself in the AI era, 1) shifting from a commodity cloud provider to AI-driven high-performance computing servers (neocloud), and 2) shifting from an independent cloud provider serving third parties to prioritizing serving the Xiaomi ecosystem.

- The first step started with the CEO change in 2H22. The company adopted a "high quality and sustainable development" strategy, embracing AI technology to enhance profitability and maintain a competitive edge. New CEO Mr. Tao ZOU gave up the previous revenue scale focus (business serving big internet companies, especially the CDN contribution dropping from half to teens %) to profitability. As a result, KC went through two years of negative revenue growth, but GPM increased dramatically from LSD in 1H22 to the high teens in 2H24, and adjusted EBITDA/operating profit both turned positive. This solved the survival issue for KC.

- The second step started with Xiaomi's determination to explore AI opportunities, which happened after Xiaomi EV's successful debut in 1H24. A key milestone was when Xiaomi and Kingsoft Cloud signed a 3-year strategic framework of connected party transactions agreement in Nov 2024. KC became the strategic cloud partner for the Xiaomi ecosystem, supporting Xiaomi's AI, AIoT, and EVs. Xiaomi would provide financial support for KC. This brought significant new growth opportunities for KC.

Key AI Products

Since 2024, Kingsoft Cloud has shifted all of its R&D and financial resources to AI cloud. In 2025, Kingsoft Cloud launched the StarFlow Platform — a turnkey AI training and inference platform built for foundation AI models and embodied intelligence, providing full lifecycle management from model development and training to inference. The StarFlow platform's Model API service offers high-availability, easy-to-integrate model calling, supporting high-concurrency inference and multi-model management. It is a neutral platform able to host essentially all open-source models. The company also offers Galaxy Stack for private deployment, providing heterogeneous GPU management, network, and intelligent container scheduling capabilities with full-stack localization. KC cooperates nearly all the Large Language Model ventures, providing large-scale computing clusters for model training.

As a result, KC revenue reaccelerated since 2H24 to 23% in 2025 with AI cloud contributing all of the incremental revenue. Together with peers' various strategy in AI, we see AliCloud, Vocano Engine, Baidu Cloud, and Kingsoft Cloud are gaining momentum in the AI era, while Tencent Cloud, Huawei Cloud, SOE telcos, and UCloud are losing momentum.

We note that one key change in the AI era is that high-end chipsets have become a scarce resource and cloud players that have supply chain capabilities are all demonstrating good pricing power and improved profitability.

What is new for cloud players in AI era?

1. Two groups of fast-rising big customers. The first group is LLM players, like DeepSeek, Minimax, Zhipu, and Kimi. They are using high-end GPUs intensively for LLM training, and most of them lease computing power rather than buying. Their good development and successful financing laid the foundation for the long-term development of AI training. They are also sophisticated in software stack capabilities and so they primarily use bare metal leasing service from neoclouds (big volume but relative low margin). The second group of rising customers are startups engaging in robotics and autonomous driving. They are relatively small in demand size vs. LLM players and require more software stack services, which means better margin for neoclouds.

2. More differentiated services make hyperscaler clouds start to integrate neoclouds. This is very similar to the US, where Microsoft and Meta are big customers of CoreWeave and Nenius. Alibaba management also mentioned in its earnings call that it is converting some AI capex to opex, by leasing computing power from neoclouds, though those neoclouds are also in competition with AliCloud.

These two new trends leave neoclouds with a lot of room to develop in the AI era as an important complement to hyperscalers and enterprises.

Exhibit 22: Kingsoft Cloud StarFlow Platform

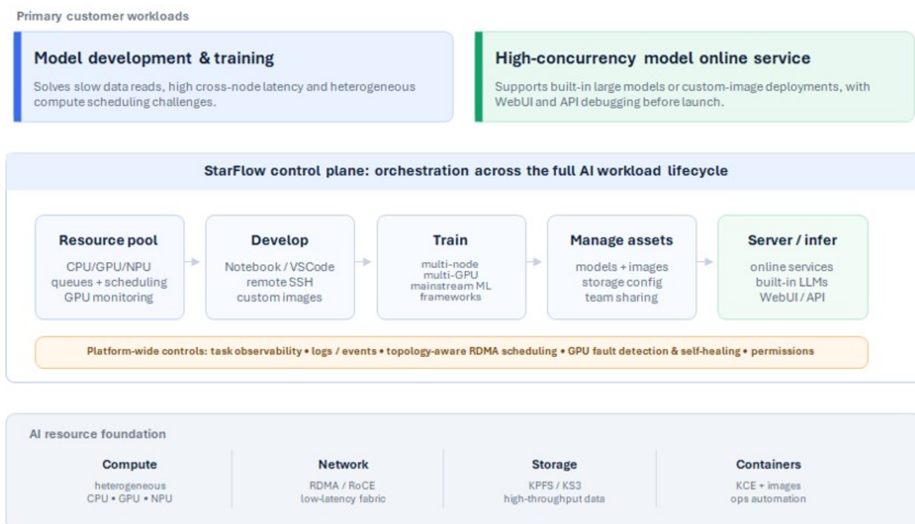

Source: Morgan Stanley Research.

Exhibit 23: Leading cloud players' revenue growth: Alicloud, Baidu Cloud, Tencent Cloud, and Kingsoft Cloud accelerated substantially in 2025, ByteDance remained elevated from a low base, while Huawei and telcos slowed down

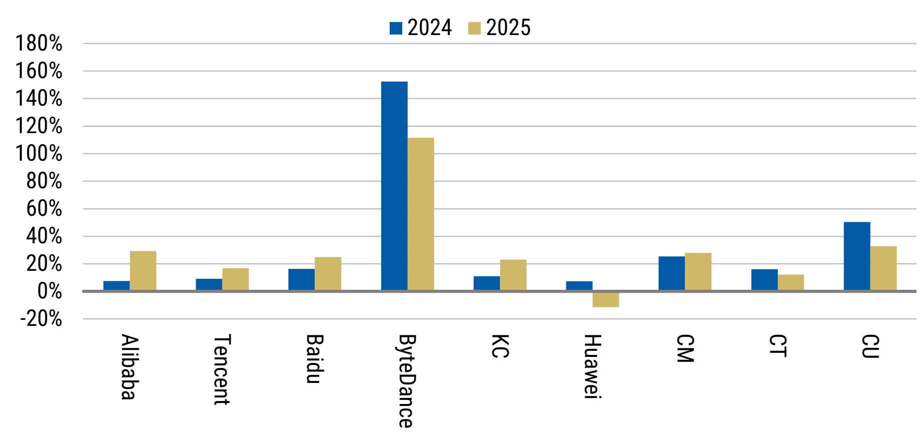

Source: IDC, company data, Morgan Stanley Research. Note: We use IDC numbers for ByteDance, Huaiwei, and the three telcos.

Exhibit 24: Kingsoft Cloud's AI revenue as a % of revenue

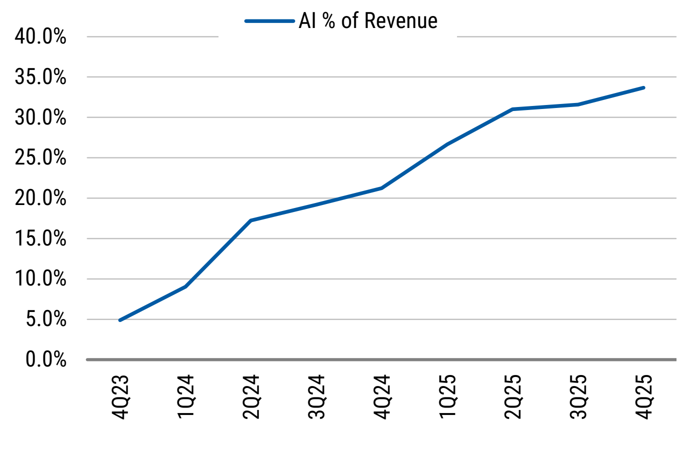

Source: Company data, Morgan Stanley Research estimates.

Exhibit 25: Kingsoft Cloud revenue breakdown

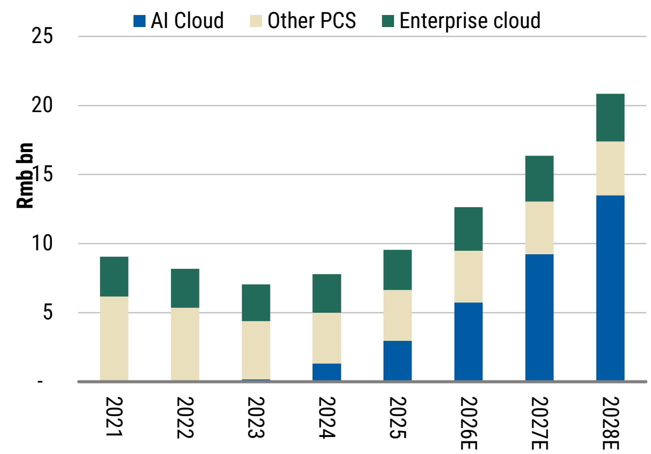

Source: Company data, Morgan Stanley Research estimates.

Exhibit 26:

Comparison between traditional and AI cloud era

	Cloud Era	AI Era
Key Products	CDN cloud storage cloud delivery basic compute	StarFlow Platform Galaxy Stack MaaS AI inference clusters
Market Share	~3–4% of China public cloud IaaS/PaaS	~3.8% overall, with rapidly growing AI segment share
Revenue Growth	+10.5% YoY (2023–2024)	AI cloud high double digit growth
Margin	Low GPM	Moderate GPM; long term 15%+
Key Verticals	Gaming Video streaming Internet	AI/LLM Cloud Autonomous driving Robotics
Ecosystem Role	Cloud services provider to Xiaomi/Kingsoft	Training/Inference to ecosystem External LLM customers
Competitive Positioning	Niche player vs. hyperscalers CDN-heavy	Neutral AI cloud platform Bridging the computing shortage

Source: Company data, Morgan Stanley Research.

Inference and Price Hikes Driving Margin Improvement

AI cloud margins set to improve vs traditional workloads

We identify three factors on the revenue side and three on the cost side that determine AI cloud margins.

On the revenue front:

- **Price per unit of computing power:** This is the price listed on the website or through bulk purchase, which is the unit cost of using the GPUs/ASICs provided by the cloud service provider. We think price per token is a function of the underlying cost of service as well as the general supply and demand environment. *Additional VAS offerings* such as bundled MaaS services or inference could increase revenue and hence enhance the margin profile.

- **Maximum computational power:** This is determined by the amount of underlying AI infrastructure, which usually comes in GPU hours for AI computing power or token generation for inference workloads.
- **Utilization rate:** This also includes the percentage of computing power that is directed to power self-owned apps or free use cases (for example, Doubao). Margins should scale faster than utilization given operating leverage.

On the cost front:

- **Capex or unit capex:** This is the amount spent on purchasing the underlying computing infrastructure (mainly AI chips and servers). We think adopting ASICs for inferencing could result in a cost advantage, given that third-party ASIC providers usually have a gross margin of 90% for chipsets and 60-70% for servers, meaning that in-house developed ASICs could reduce unit capex by 50%+ on a server basis.
- **Depreciable life:** US hyperscalers have already extended their servers' depreciable life to six years; meanwhile, Chinese cloud providers are still depreciating its servers over 3-5 years. Depreciation policy could have an impact on margins.
- **Other operating costs:** These include power and data center costs. The hyperscalers are already compressing this cost component by moving data centers to remote areas, which are known for their lower rental costs and cheaper power tariffs (see our report here).

Exhibit 27: Key factors determining the EBIT margin

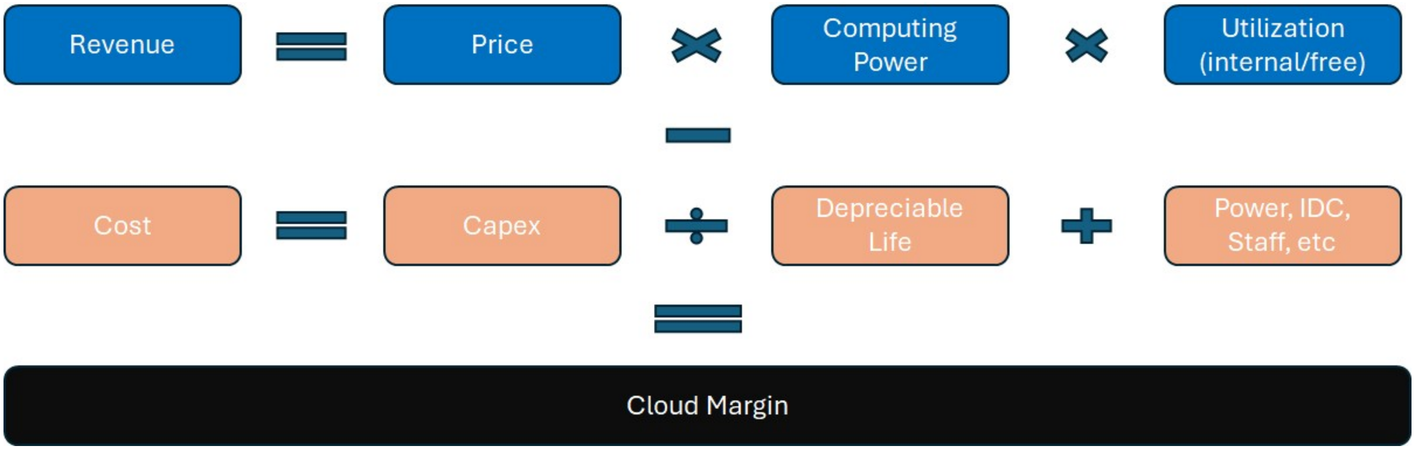

Source: Morgan Stanley Research.

From training to inference: We think that training workloads are more commoditized (offered and priced at GPU hours) – meaning that the pricing mainly reflects the supply dynamics. This means that pricing could fluctuate more depending on the availability of chipsets.

On the other hand, we see more opportunity for KC to improve margins for inference workloads. Inference workloads are billed on a token/API basis and the computing power from a single server can be allocated to multiple workloads. This brings more flexibility to margins through:

- **Better pricing mechanism and flexibility:** The pricing of training workloads is relatively fixed and commoditized, but for inference workloads the pricing can be more flexible. Hyperscalers can adjust pricing by adding more VAS such as MaaS and other platform services to drive higher revenue per GPU utilized. KC can also become the proprietary service provider for Xiaomi MiMo's close source APIs in the future. We think MaaS incremental margin could be in the range of 30-50%, which is much higher than the company's blended GPM.
- **Architecture/software innovation:** Inference workloads have more room to compress costs and improve utilization through architecture innovation such as KV cache, which can improve the FLOPS needed to generate tokens and hence improve the tokens generated per second for given computing resources. Architecture innovation can improve tokens generation by more than two folds with the same amount of hardware investment. Minimax noted that their revenue per server could be 2x higher than their peers in average, mainly due to more tokens generation per FLOPS.
- **More elastic resource allocation:** Training workload might have to occupy the entire server whereas inference workloads allow more elastic allocation of the computing resources. This means that multiple workloads can be processed on a single server allowing concurrent scheduling/batching, which maximizes the resource utilization and hence leads to higher margin. Filling workloads utilization during night time and weekends could effectively increase the margin of the cloud services revenue.

MaaS is a rapidly growing revenue stream. According to IDC, the MaaS revenue in China grew by 332% YoY to Rmb3.1bn in 2025 and is set to grow by another 500% YoY to Rmb18.6bn in 2026. This is primarily driven by surging token usage, which grew by 17x in 2025 to 1,900trn and another 20x to 40,000trn in 2026. Token prices on the other hand should be more than halved to increase adoption and affordability. IDC also notes that the future development of MaaS should be driven by cost efficiency (architecture to increase throughput), performance and the accompanying tools and application ecosystem.

We believe MaaS also serves as an important margin enhancement tool for cloud providers. As we noted above, MaaS can generate very high margins because IaaS infrastructure represents the largest cost component required to deliver MaaS services. Once cloud providers establish the underlying infrastructure, incremental MaaS revenue can contribute meaningfully to profitability.

Using MiniMax and Zhipu as examples, we estimate their cloud API margins at 60% and 21%, respectively, as of 2026. We note that MiniMax enjoys a higher margin, thanks mainly to overseas revenue and its multi-modal product (typically higher than foundational LLM).

Assuming hyperscalers monetize their APIs at similar pricing levels, we estimate this could increase their revenue by 25-150%. We believe most of the incremental revenue would contribute to margins, as cloud infrastructure represents the largest component of COGS. In practice, we expect domestic hyperscalers to trend toward the lower end of this range, given their focus on the China market and their ability to offer multimodal capabilities.

Exhibit 28: How training differs from inference workloads

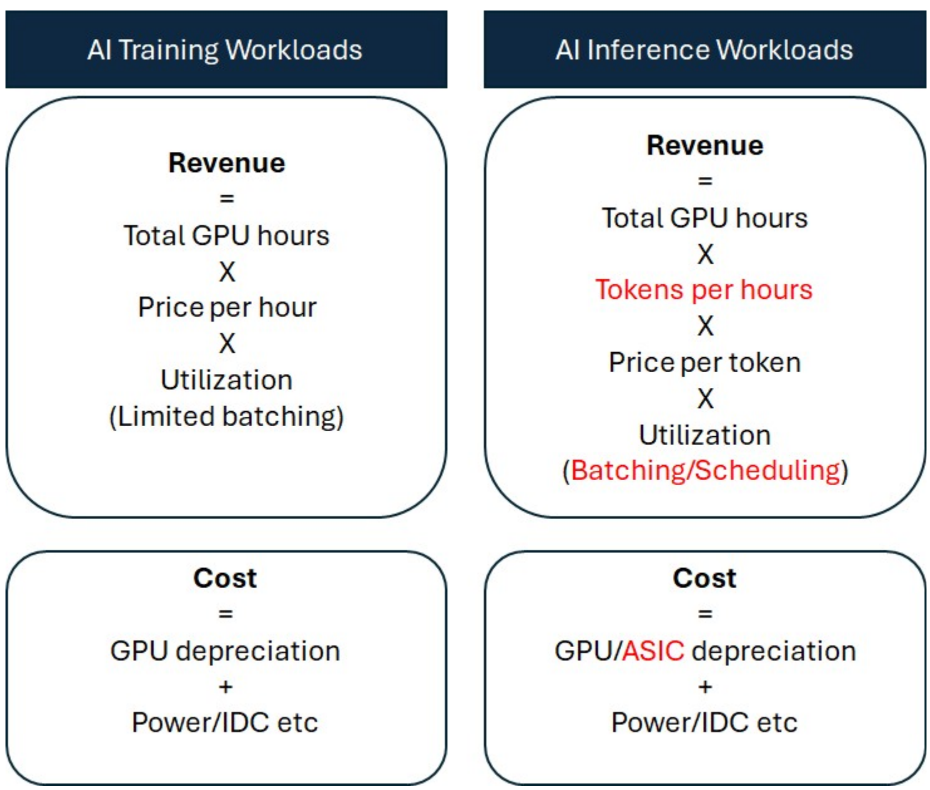

Source: Morgan Stanley Research.

Exhibit 29: China MaaS share by revenue

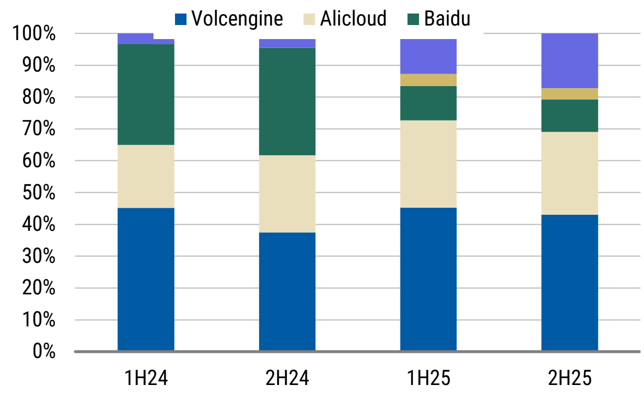

Source: IDC, Morgan Stanley Research.

Exhibit 30: China MaaS share by token

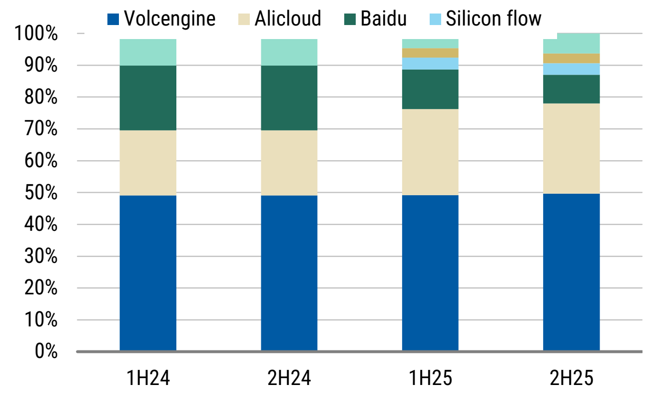

Source: IDC, Morgan Stanley Research.

Exhibit 31: MiniMax – MaaS margin

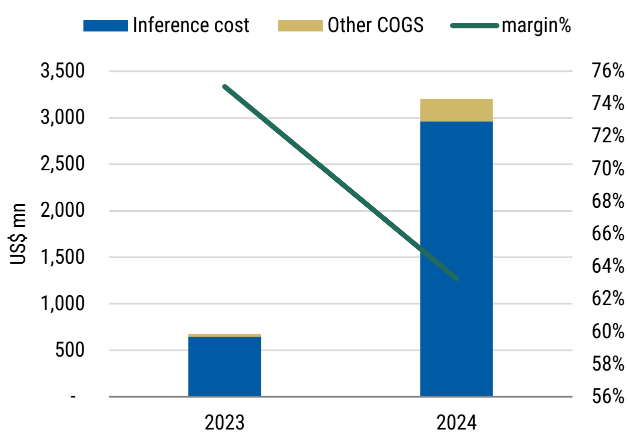

Source: Company data, Morgan Stanley Research.

Exhibit 32: Key AI hardware COGS as a percentage of revenue

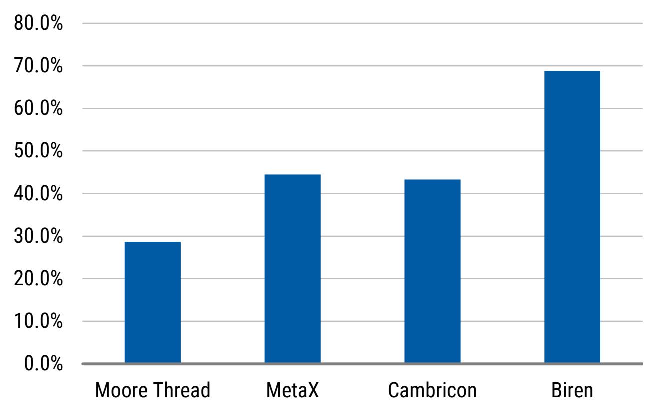

Source: Company data, Morgan Stanley Research.

Price hikes set to benefit pure cloud players

We have witnessed a series of price hikes in China and abroad (see, China's Emerging Frontiers: China's AI Path: Monetizing Surging Token Use via AI Cloud (16 Mar 2026)).

Cloud price hikes started with global hyperscalers

- On January 4, 2026, AWS increased the prices of its EC2 Capacity Blocks for Machine Learning by approximately 15%. The hikes, which affect high-performance instances like the p5e.48xlarge (NVIDIA H200 GPUs), are driven by high demand for AI resources, with prices rising from US$34.61 to US$39.80 per hour in many regions.
- On January 27, 2026, Google Cloud (GCP) announced that it would be implementing significant price increases effective in 2026 for networking, storage, and AI infrastructure, driven by rising demand and infrastructure costs. Data transfer rates for CDN and peering are rising to 100% in some regions, while the cost of specific storage and AI-focused services is also increasing. The new pricing will be effective from May 2026.

- On May 21, Nebius announced that effective June 1, the price for on-demand capacity will rise by an average of 29%, while preemptible capacity will increase by 51%, aimed at addressing the ongoing strong demand for GPU capacity while maintaining competitive pricing in the market.

Mainly driven by a sharp price hikes upstream:

- **Memory:** Price hikes started last year and the momentum continued into 2026, with companies such as SanDisk raising SSD prices again in January 2026 owing to component constraints.
- **CPUs:** In the beginning of the year, we have also seen component price hikes spill over to CPUs.
- **Other costs are stable:** Meanwhile, third-party IDC costs look relatively stable.

Exhibit 33: Silicon Data H100 Rental Index

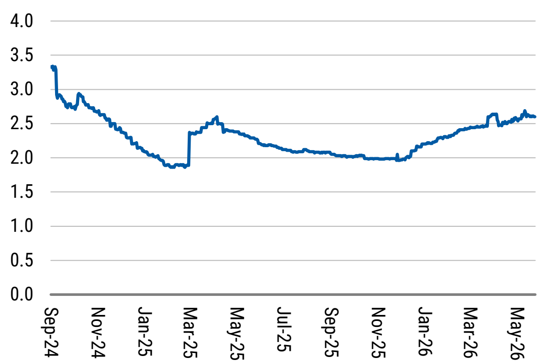

Source: Bloomberg, Morgan Stanley Research.

Exhibit 34: Silicon Data B200 Rental Index

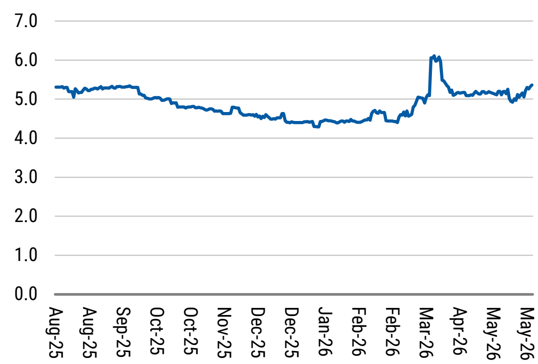

Source: Bloomberg, Morgan Stanley Research.

Following these developments, Chinese cloud and related products also started price hikes

- Wangsu announced 35-40% price hikes for its CDN offerings, effective February 1, 2026 (Link).
- Ucloud (not covered) also announced that it would hike prices for its cloud product line starting from March 1, 2026, for all contract renewals and new customers (Link).
- Tencent's AI platform also hiked the price of its in-house model in March as much as 400%, which we think can be attributed to the growing cost of underlying infrastructure (Link).
- Alicloud announced a price hike for its AI computing product by 5-34% in Mar 26 effective on Apr 18 (Link).
- Baidu cloud announced a price hike for AI computing related services by 5-30% following Alicloud's announcement (Link).
- Kingsoft Cloud announced it will raise prices for AI computing and file storage services by 15-50%, effective July 12, 2026. Management attributed the price increase to strong global demand for AI infrastructure and rising hardware costs (Link).

We note that price hikes have been gradually passed through from underlying semi/ hardware to service providers. The pricing mechanism has been smooth, which indicates an imbalanced supply/demand structure in the AI computing field.

Exhibit 35: Recent price hikes globally

Time	Product /Company	Key events
Jan-26	AWS	AWS raised EC2 Machine Learning Capacity Block prices by ~15%, mainly affecting AI-heavy workloads.
	Google Cloud	Announced higher data transfer prices effective from May 2026.
	Wangsu	Announced price hikes for CDN products by 35-40% effective from February 1, 2026.
Feb-26	Ucloud	Announced price hikes for cloud product line effective from March 1, 2026, for all contract renewals and new signings.
Mar-26	Tencent	Raised token prices for a series of models on its MaaS platform.
	Alicloud	Raised prices for AI-computing-related offerings by 5-34% effective from April 18, 2026.
	Baidu	Raised AI-computing-related offerings by 5-30% effective from April 18, 2026.
Apr-26	Tencent	Raised prices for AI-computing-related products effective from May 9, 2026.
	Alicloud	Raised prices for some server products effective from May 15, 2026.
Jun-26	Kingsoft Cloud	Raised prices for AI-computing-related products by 15-50% and for storage-related products by 30-50% effective from July 12, 2026.

Source: Company data, Morgan Stanley Research.

Exhibit 36: Taobao retail price of NVIDIA 4090 and 5090 graphic cards

Source: Taobao, Morgan Stanley Research.

What could be the terminal margin for KC?

We expect KC to achieve a blended gross margin of at least 20% over the longer term, vs. 16% in 2025. AI-related cloud services already generates 20-25% gross margins, and continued growth in this revenue stream should help margins expand toward the 20% level. Alibaba has also indicated a long-term margin target of 20%, vs. 9% in 1Q26, implying >10ppt upside potential. Similarly, we believe KC could deliver >10ppt of OPM expansion over the next five years.

We think the primary drivers of margin expansion for KC include

- **Architecture improvements:** KC focuses on more effective token generation and improving the overall utilization of computing power. This could be done through selling redundant resources at night and weekends to certain customers such as humanoid/autonomous driving as well as improving throughput by adopting FP4, etc.
- **KV cache:** Once the computing hits the KV cache, the margin profile could be much better as it does not use memory bandwidth anymore. KC has a high hit ratio for KV cache, which effectively reduces the cost of compute and hences improves the margin.
- **Price hike:** While KC did not announce any headline price hike, our checks suggest that they might narrow discounts given on certain contracts to customers. This should also result in an effective price hike and hence improve the overall margin of the business.
- **Higher margin product offerings:** On top of pure computing power, KC can also provide MaaS services to the customers by directly offering model capabilities through tokens. As we have discussed above, adding an additional layer is pure margin accretive. Other than that, customers aside from cloud and model players lack capabilities in building the network, which allows KC to earn additional margins. Those are typically customers from autonomous driving and robotics.

Looking at the COGS breakdown, we expect server depreciation to become a larger cost component going forward due to increased capex and financing lease obligations related to server procurement. However, we expect price increases for server rental services to largely offset this impact, as pricing typically tracks underlying procurement costs. Beyond server depreciation, we expect the remaining cost components including IDC and fulfilment expenses, to benefit from operating leverage, as their growth rates should remain below overall revenue growth.

Exhibit 37: KC's GPM profile

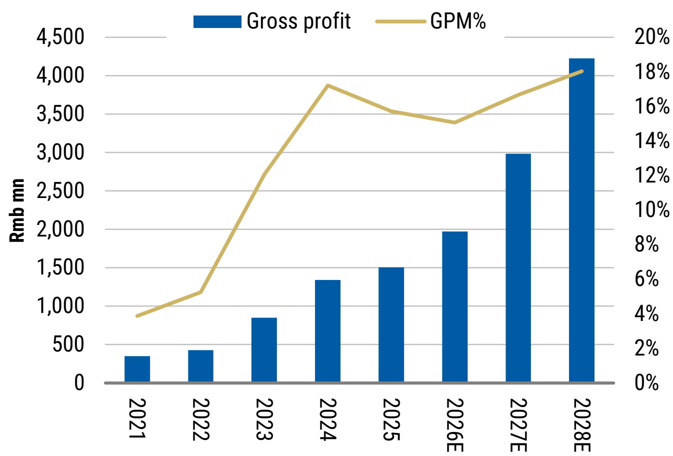

Source: Company data, Morgan Stanley Research estimates.

Exhibit 38: COGS breakdown

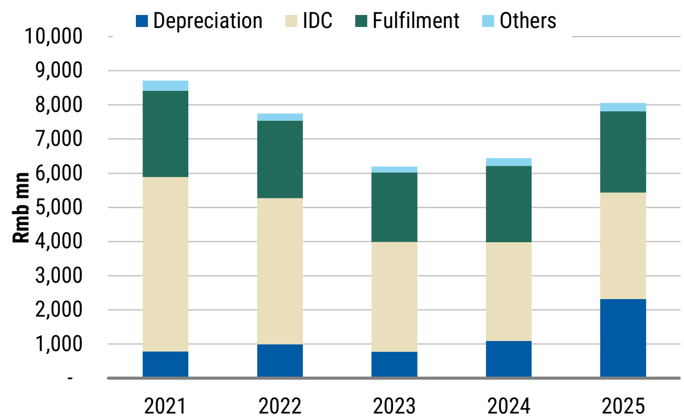

Source: Company data, Morgan Stanley Research.

Read-across from AliCloud's recent tone on cloud margins

In the March quarter results, AliCloud noted that AI-related revenue accounted for 30% of external cloud revenue and MaaS-related ARR reached Rmb8bn and is set to exceed Rmb10bn in the following quarter and Rmb30bn by the end of the year, indicating strong demand for AI-related cloud services.

In terms of the margin of the business, Alicloud previously guided to a 20% EBITA margin for cloud over the long term, largely driven by MaaS revenue contribution. This indicates that MaaS is likely to come with a higher margin compared to the pure GPU cloud business. This is in line with our thesis above that the adoption of ASIC and MaaS could drive a higher margin for the business. We note that pure MaaS margin can be as high as 60-70% if we look at the margins of Minimax and Zhipu.

We think that KC will likely generate a smaller contribution from MaaS revenue and lower MaaS margins, as it does not develop proprietary foundation models and therefore has less scope to optimize model performance relative to original model developers. As a result, the MaaS opportunity will depend on the commercialization progress of MiMo, where KC could serve as the dedicated MaaS provider for Xiaomi's large models. Beyond MaaS, we expect KC's overall margin profile to gradually converge toward that of AI cloud services as AI-related revenue accounts for a larger share of total revenue..

Exhibit 39: AI Models on KC's StarFlow Platform

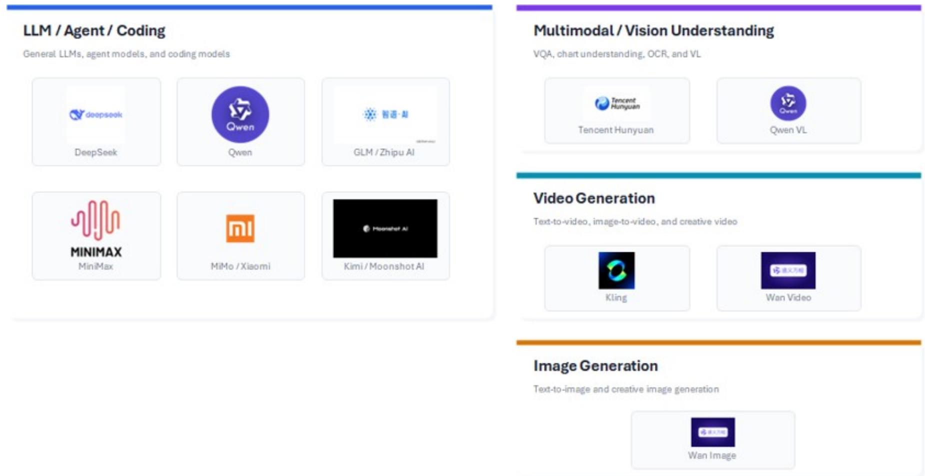

Source: Morgan Stanley Research.

Exhibit 40: AliCloud vs. Kingsoft Cloud EBITA margin, 2022-29e

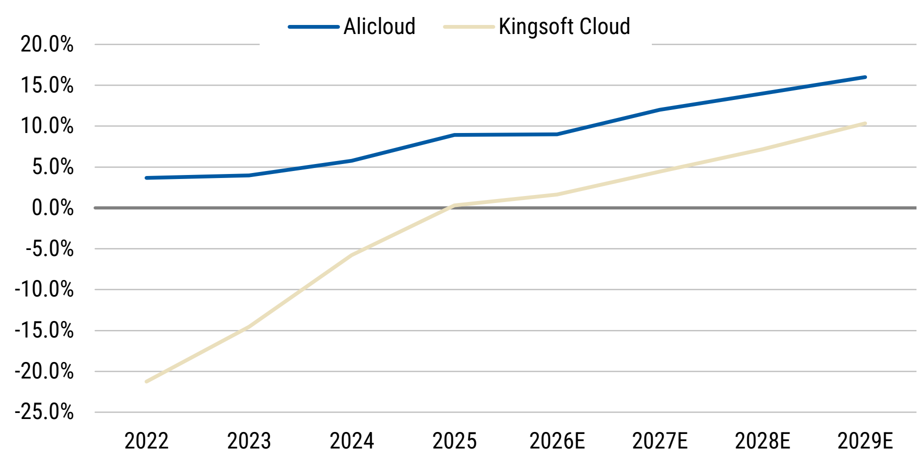

Source: Company data, Morgan Stanley Research estimates.

Unit Economics Analysis Under Longer Lifecycles, Financing Leases, and Customer Prepayments

The market has concerns over the neocloud business model due to its heavy capital requirements and views it as "balance sheet as a service". We would note, however, that recent industry developments are improving the upside risks of this model, as there is a fundamental shortage of computing power combined with a low interest rate environment in China. This supports easier access to financing leases, and charging customer prepayments can materially ease KC's balance sheet constraints.

Our unit economic analysis suggests 4-6ppt IRR upside from each of the following changes (they are not mutually exclusive):

- Servers can operate and collect rental revenue longer (six years) than expected (five years).
- Neoclouds get access to financing leases to cover 50% of initial capex.
- Neoclouds charge customers a one-year prepayment for the rental.

Most neoclouds offer hourly-priced short-term computing power leasing to mid to small customers, and investors can monitor the real time price to gauge the market dynamics. This will be quite elastic to earnings given most of the costs are fixed by nature. However, that is only a small part of the neocloud business (a single-digit %) and the bulk of neocloud capacity is leased to big customers under long-term contracts and predetermined pricing schemes. That is why we believe our unit economics analysis is more relevant to the underlying business.

Revenue and profit structure

KC mainly generates AI revenue by procuring the AI servers and building them into a network to provide AI-related training/inference services to customers (mainly bare metal leasing to large customers under long term contracts). The primary cost of the service is server procurement (depreciation), followed by data center costs (hosting, power and network) and a relatively small amount of personnel costs. As AI server prices have been rising in China due to severe computing power shortages, KC prices its rental based on the server procurement price to lock in a certain gross margin (20-25% GPM based on our estimates).

For KC, bare metal leasing contracts with big customers are usually 4-5 years long (to cover the lifecycle of GPUs) with a gross asset yield of 30%+ or a payback period of about 3-4 years (cash inflow covers the initial capex). Typically, ecosystem customers provide a prepayment to KC. This was very evident in the 2025 result, as its OCF rose 5x to Rmb3.8bn, mainly driven by a Rmb2.2bn prepayment from Xiaomi. In 2026, this scheme gradually expanded to other customers due to the capacity shortage, with big customers prepaying a percentage of the rental fee and small customers even prepaying for the equipment's whole lifecycle. Put another way, KC is transferring part of its pricing power to cash flow terms. After 4-5 years of rental, the server should be disposed of at minimal residual value (accounting assumption is 0-5% of initial cost) after full depreciation. Server equipment operating longer or shorter than the 5-year depreciative life will provide upside and downside risks to the business, respectively.

Exhibit 41: Business model

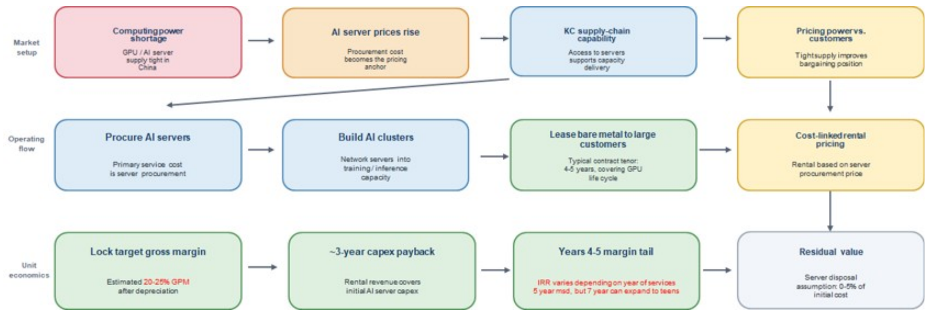

Source: Morgan Stanley Research.

Source of funding – financing leases and customer prepayments becoming more important

The neocloud business is capital-intensive due to the high price of GPU servers, and is frequently viewed as "balance sheet as a service". With limited balance sheet capability via financing leases (initially from Xiaomi Finance, but later other independent financial leasing partners) and customer prepayments (varying from some percentage of the annual rental to the whole contract lifecycle). Both approaches require customers to sign long-term leasing contracts, which may not be attractive to customers who purely train LLMs (they only want short-term usage of the most cutting-edge GPUs), but is quite doable if customers have diversified workloads (e.g., LLM training in the first 2-3 years, and inference or small model training when the next generation GPU becomes available).

- **For large customers**, KC engages in financing leases to better manage its cash outflow and leverage cheap funding costs in the onshore debt market. This model usually requires dedicated capacity for certain investment grade customers. We note that half of KC's server procurement is conducted through financing leases to smooth the cash outflow. The financing lease contract is usually 5 years in tenure, matching the length of its leasing contract with customers. In 2025, KC spent Rmb4.7bn in cash capex while also adding a Rmb3.5bn financing lease contract, making its effective capex reach Rmb8.2bn. We forecast KC's cash capex to reach Rmb10bn while financing lease contracts should contribute another Rmb10bn server procurement in 2026, totaling Rmb20bn for server procurement to drive AI revenue growth.
- **For ecosystem customers and small customers**, KC charges a substantial prepayment (sometimes even covering the whole contract lifecycle). Customers can get attractive pricing (in exchange for the time value of the capital) in a sellers' market. This effectively transitions the business model from upfront cash outflow and gradual inflow/revenue to a business model with upfront cash collection and gradual revenue recognition, reducing the balance sheet/cash flow burden significantly.

Our analysis shows that the different structure does not impact the reported operating profit margin, but has a significant impact on the cash cycle and IRR of the project. We think the overall risk and volatility of KC will be reduced significantly under this model through,

1. **Better duration matching** – 5-year financing leases (depreciation on the income statement) will be matched by the rental contract duration with a pre-specified margin. This improves the overall profitability of the income statement and results in better cash flow management (more smooth outflow vs a one-off lump sum).
2. **Faster cash recycling** – given the contract value is collected at day 1, KC can better utilize that cash for further investment if needed or just improve the overall cash and leverage structure of the firm.

Another key assumptions to make is the terminal value. We note that accounting wise, companies usually assume a 0-5% disposal value for the AI servers. However, recent industry practice is that the AI servers can function longer than 5 years (like early stage Nvidia A and H series of GPU can continue to generate rental revenue even after being fully depreciated). US peers Oracle, Amazon, Alphabet, Microsoft, and CoreWeave have already announced they will extend the GPU server depreciation cycle from 4-5 years to 6 years, while Chinese cloud players have yet to do so. This will lead to relatively low margins for Chinese clouds in the first 4-5 years, followed by a significant step up afterwards. Our recent channel checks suggested a 5-year old NV A series GPU secondary market price is almost 70% of its initial cost in China.

Financial statement impacts: As illustrated in our UE calculation, a base case 5-year rental contract could yield a pre-tax unlevered IRR of 9%, which we think is decent but not exciting. However, with the help of leases, prepayment and extending operating life, the return can be enhanced significantly.

- **Financing leases:** Financing leases reduce the cash burden on upfront investments, but the subsequent cash flow and operating profit are lowered due to the requirement to pay lease fees and additional interest costs to the finance provider. Under a lease structure, the annual cash flow is reduced but compensated by lower upfront spending. Our scenario analysis indicates that a 50% lease financing could increase unlevered IRR by 4ppt. We also note that local banks are supportive of higher leverage ratios.
- **Prepayments:** Prepayments are positive for both cash flow and net margins. Cash IRR is improved given upfront cash collection and could be further enhanced if we account for investment into other projects. Prepayments also support net margins by reducing financing requirements and the associated interest burden. Our scenario analysis suggests that a one-year prepayment arrangement could increase IRR by approximately 6ppt relative to our base case.
- **Extending useful life:** As discussed earlier, servers can provide services beyond their 5-year depreciation period, which could materially improve profitability. Once a server is fully depreciated, net margins could increase significantly, to 64% from 8% before depreciation. Our analysis also suggests that extending the useful life by one additional year could increase unlevered IRR by approximately 5ppt.

In reality, KC adopts a combination of the above-mentioned factors to drive a pre-tax IRR that is much higher than the base case against a cost of debt of ~3% in China.

Why is this important? We believe KC's capital structure has been a key investor concern over the past two years. In the absence of sustainable net cash inflows, the company's financing needs could continue to grow, whether through equity issuance, which would dilute existing shareholders, or additional debt, which would place further pressure on the balance sheet and cash flows. We believe the improvements in KC's business model have made its capital structure more manageable. While we cannot rule out the need for further financing, that is likely to come with a more visible project return and payback period. Leveraging financing leases has already contributed to a significant improvement in leverage, with net debt/EBITDA declining from 6.6x in 2024 to 2.0x in 2025.

Exhibit 42: KC unit metrics on neoclouds

Year	0	1	2	3	4	5	6
Capex	100						
LTC	80%						
Interest rate	3.0%						
Debt	80						
Equity	20						
Tax rate	15.0%						
Revenue		**30.0**	**30.0**	**30.0**	**30.0**	**30.0**	**30.0**
Opex		4.5	4.5	4.5	4.5	4.5	4.5
EBITDA		**25.5**	**25.5**	**25.5**	**25.5**	**25.5**	**25.5**
Margin		*85%*	*85%*	*85%*	*85%*	*85%*	*85%*
Depreciation		20.0	20.0	20.0	20.0	20.0	-
EBIT		**5.5**	**5.5**	**5.5**	**5.5**	**5.5**	**25.5**
Margin		*18%*	*18%*	*18%*	*18%*	*18%*	*85%*
Interest expense		2.4	2.4	2.4	2.4	2.4	2.4
Tax expense		0.83	0.83	0.83	0.83	0.83	3.83
Net profit		**2.28**	**2.28**	**2.28**	**2.28**	**2.28**	**19.28**
Margin		8%	8%	8%	8%	8%	64%
Gross asset yield		30%	30%	30%	30%	30%	30%
Payback years	3.92						
If 5 years							
Unleveraged cashflow	(100)	25.5	25.5	25.5	25.5	25.5	
IRR (pre-tax)	9%						
If 6 years							
Unleveraged cashflow	(100)	25.5	25.5	25.5	25.5	25.5	25.5
IRR (pre-tax)	14%						
If prepaid 1 year							
Unleveraged cashflow	(74.5)	25.5	25.5	25.5	25.5		
IRR (pre-tax)	14%						
If 50% leased							
Lease payment	-	(11.2)	(11.2)	(11.2)	(11.2)	(11.2)	
Unleveraged cashflow	(50.0)	14.3	14.3	14.3	14.3	14.3	
IRR (pre-tax)	13%						

Source: Company data, Morgan Stanley Research estimates.

Exhibit 43: KC total investment and cash expense mix

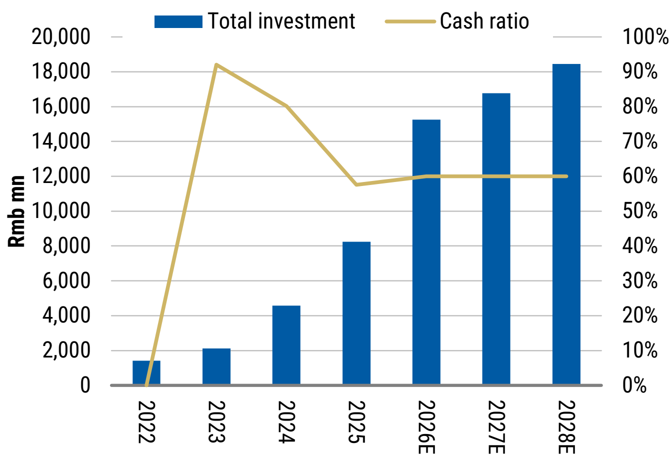

Source: Company data, Morgan Stanley Research estimates.

Exhibit 44: Prepayment forecasts

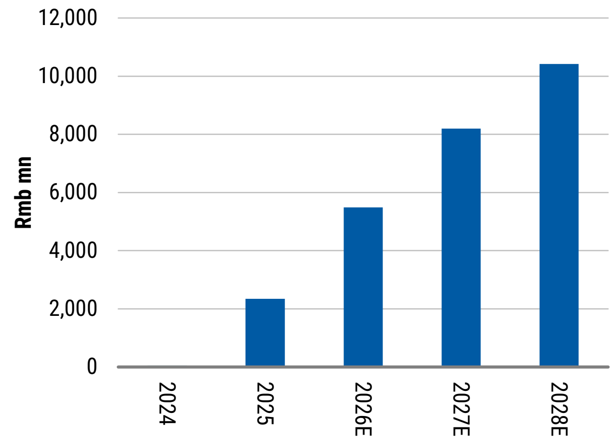

Source: Company data, Morgan Stanley Research estimates.

Cooperation with Xiaomi's Ecosystem

Xiaomi's AI ambition and Kingsoft Cloud's role in the ecosystem

Mr. Jun Lei is the final controller of both Xiaomi and Kingsoft Group (including Kingsoft Corp, Kingsoft Office, and Kingsoft Cloud). They have jointly formed a vibrant technology ecosystem with Kingsoft Cloud acting as the underlying computing power platform supporting various models and applications, including AI models (MiMo), smartphones, home appliances, cars, games, office software, and other emerging applications. Except for the games business under Kingsoft Corp and office software under Kingsoft Office, the rest is all coming from Xiaomi, which drives the computing demand.

Xiaomi's AI strategy, as articulated by Lei Jun in Xiaomi's 2026 Investor Day, is best understood as "AI everywhere" embedded across its ecosystem rather than a standalone business line. The company is committing substantial resources to this direction**, with roughly Rmb200bn in R&D planned for 2026–230**, a meaningful portion allocated to AI and chips, and multi-billion-dollar dedicated AI investments over the next few years. AI already represents a growing share of R&D spend and is being deployed across Xiaomi's full technology stack, from in-house large models (MiMo series) to operating systems, chips, and edge devices. Strategically, this supports Xiaomi's transition toward a "hardcore technology company," with AI acting as the intelligence layer across its "Human × Car × Home" ecosystem. The emphasis is pragmatic: scale deployment and product integration over frontier model leadership, using Xiaomi's large installed base as the primary monetization engine.

In smartphones, AI is positioned as a key driver of user experience upgrades and premiumization. Xiaomi is integrating AI into its operating system and on-device capabilities, enabling features such as smarter assistants, imaging enhancements, and personalized interactions. The goal is to increase device stickiness and support higher ASPs, while also driving greater engagement with Xiaomi's internet services layered on top of its handset base.

In electric vehicles, AI plays a central role in both intelligent driving and in-car user experience. Xiaomi is applying AI to autonomous driving systems, perception, and decision-making, as well as to smart cockpit functionality. This aligns the EV business closely with the broader ecosystem, where the car becomes another connected node, with AI enabling seamless interaction between vehicle, user, and home devices.

In AIoT and smart home, AI enhances automation, interoperability, and user convenience across Xiaomi's extensive portfolio of connected devices. By embedding AI into these products, Xiaomi aims to create a more intuitive and responsive home environment, strengthening ecosystem cohesion and increasing cross-device usage. This segment is particularly important as a scale advantage, where Xiaomi's large installed base can amplify the impact of AI-driven features.

For AI models, Xiaomi is a late entrant but the debut of AI model MiMo surprised the market on the positive side. Xiaomi gained a lot of industry insights after its investment in multiple AI start-ups. Together with Xiaomi's proprietary ecosystem data, unique convergence use cases, strong talent pool and capital support, this late-mover advantage is playing out again, similar to its journey in smartphone and cars. Xiaomi's AI model focuses on the integration with its hardware ecosystem, different from other AI models focusing on intelligence ranking or token usage. MiMo has three versions: 1) MiMo-Pro (flagship version) serves agents like coding, to prove its capabilities; 2) MiMO-Omni version as a world model serving hardware convergence use cases; and 3) MiMo-TTS/ASR version focuses on voice input, recognition and generation.

Finally, at the platform and infrastructure level, AI underpins Xiaomi's efforts in operating systems and semiconductors. The company is integrating AI capabilities into its OS and investing in chip development to optimize on-device performance and reduce reliance on external suppliers. This layer is critical for long-term differentiation, as it enables tighter vertical integration and more efficient deployment of AI across all end-user products.

How do they collaborate?

According to the connected party transaction agreement, Kingsoft Cloud provides cloud services to Xiaomi with a new cap of Rmb4bn and Rmb6bn in 2026/27 (up from Rmb3.1bn and Rmb4.0bn), respectively. We expect ecosystem revenue to account for 35%/39% of our revenue estimates. KC needs to offer the cloud service to Xiaomi at a fair market price (not higher than the price KC offers to third parties and not lower than the price third-party vendors offer to Xiaomi), and the capacity offered to Xiaomi is dedicated. Xiaomi will finance Kingsoft Cloud with an annual cap of Rmb4.2bn, including a Rmb1.2bn finance lease, Rmb1.2bn factoring and a Rmb2.0bn secured loan.

We note that Xiaomi's ecosystem progress could be crucial to Kingsoft Cloud's revenue growth in the following ways:

- **AIOT x Smartphone:** This is the traditional business for Xiaomi, accounting for more than ⅔ of Xiaomi's revenue contribution to Kingsoft Cloud as of 1H24. Kingsoft Cloud continues to provide cloud support for Xiaomi's ecosystem as well as storage and supporting the relevant internet services. However, we think the proportion of contribution should decline as cloud related services ramp up further.
- **Electric vehicle:** Xiaomi's cumulative EV shipments have reached 411k units in 2025 and Kingsoft Cloud offers assistance through 1) R&D process, 2) storage growth from image processing and sensor data, and 3) logistics and operation of smart EV. Cumulative EV shipments are expected to reach 1.1mn units by the end of 2026E, according to our hardware analyst.
- **Xiaomi's foundational model:** Xiaomi's AI model MiMo has been ranked as one of the leading Chinese models by Artificial Analysis. It has also achieved top tokens usage on Open Reuter in the first 2 weeks of relesae. We think Kingsoft Cloud could benefit from the training and upgrade for future versions of MiMo as well as the internal adoption of the LLM in Xiaomi's traditional and auto businesses. KC could also act as the proprietary API vendor (model as a service) for MiMo's closed source models in the future. This could become a bull case to drive higher MaaS revenue.

We estimate that Kingsoft Cloud could benefit from rising business demand from Xiaomi (from both new device shipment and also the rising usage of existing user base), driving its revenue to grow by 31%/30% in 2026-27. In the first four months of 2026, Xiaomi already contributed Rmb749mn to cloud service revenue, up 81% YoY, according to KC.

Exhibit 45: Shareholding structure of KC (as of December 31, 2025)

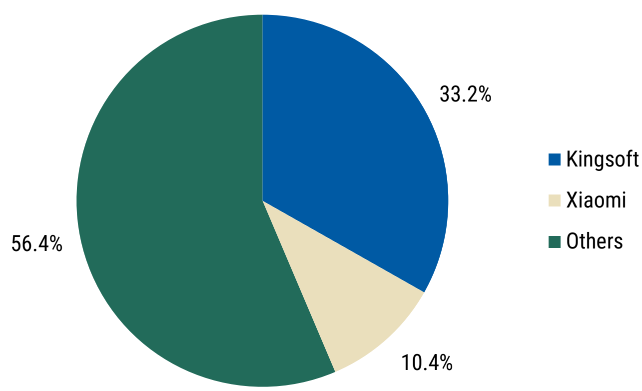

Source: Morgan Stanley Research.

Exhibit 46: KC revenue breakdown (ecosystem vs. other)

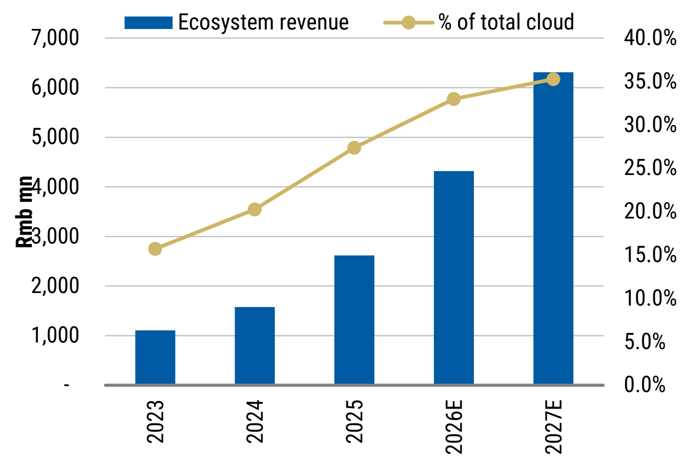

Source: Company data, Morgan Stanley Research estimates.

Exhibit 47: Ecosystem chart

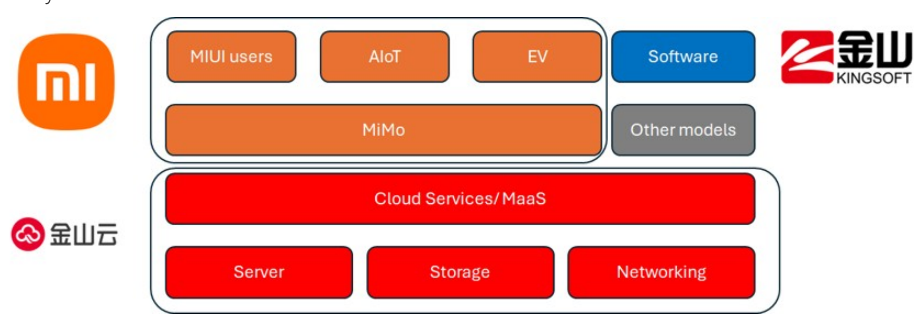

Source: Company data, Morgan Stanley Research.

Exhibit 48: Xiaomi's smartphone shipments

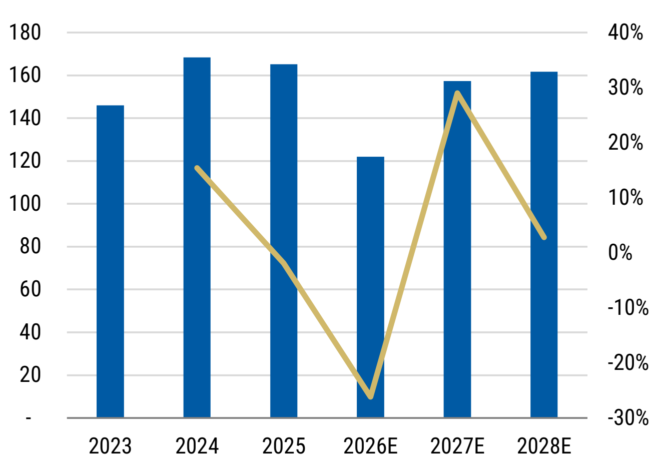

Source: Company data, Morgan Stanley Research estimates.

Exhibit 49: Xiaomi's AIOT shipments

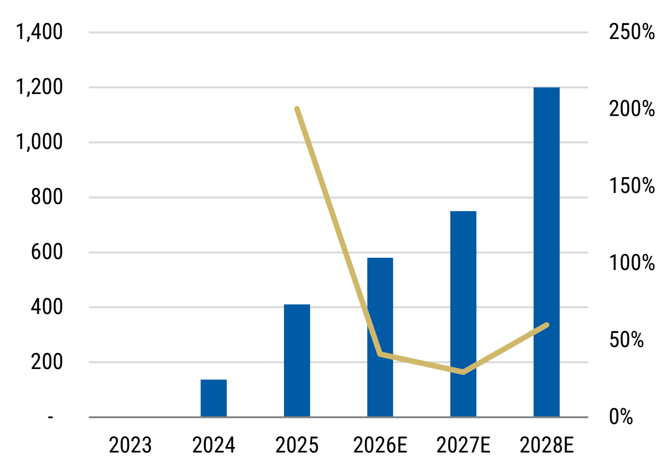

Source: Company data, Morgan Stanley Research estimates.

Exhibit 50: Xiaomi's EV shipments

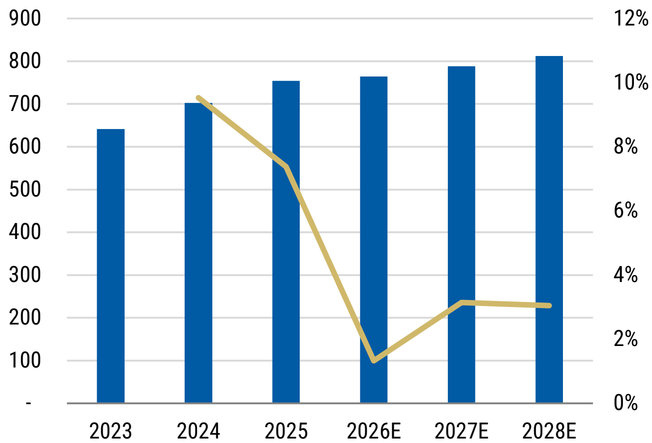

Source: Company data, Morgan Stanley Research estimates.

Exhibit 51:

Xiaomi MiMo's performance highlight

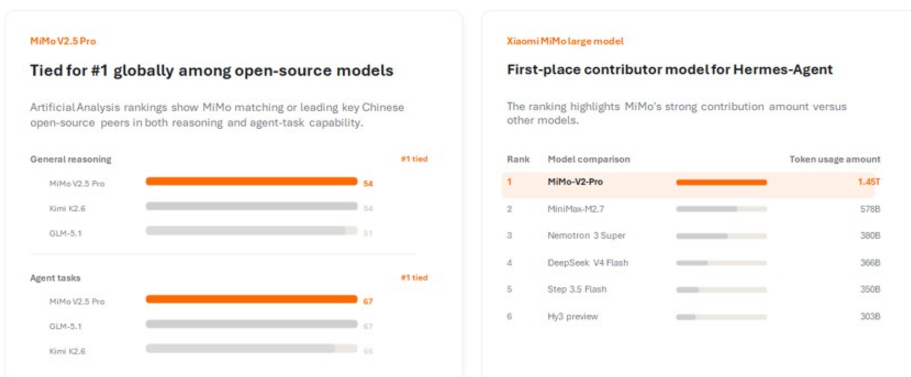

Source: Xiaomi, Artificial Analysis, OpenReuter Morgan Stanley Research.

Earnings Estimates and Valuation

We forecast revenue to grow 37%/37%/31% YoY in 2026-28, driven by a continuous ramp-up in AI-related cloud services.

- Public cloud services is forecast to grow 50%/47%/37% YoY in 2026-28 backed by strong growth in ecosystem revenue (mainly Xiaomi) and AI-related services. AI-related revenue was to grow 109%/74%/49% YoY during the period, mainly driven by GPU procurement with a gross asset yield of 28-31%.
- Enterprise cloud services is expected to remain stable with 8%/5%/4% annual growth in 2026-28.

We forecast GPM of 15-17% in 2026-29, mainly driven by a mix shift towards more AI-related business which has a higher margin profile (20-25%) than the remaining business. Depreciation will be the largest component of COGS due to the large amount of capex and leases needed to support business expansion.

We forecast adjusted EBITDA to expand at a 79% CAGR over 2025-28 off a low base. We expect the margin to expand rapid from 24% in 2025 to 57% in 2028.

We forecast adjusted net profit to break even, mainly driven by a larger scale and higher-margin business.

Exhibit 52: KC total investment size

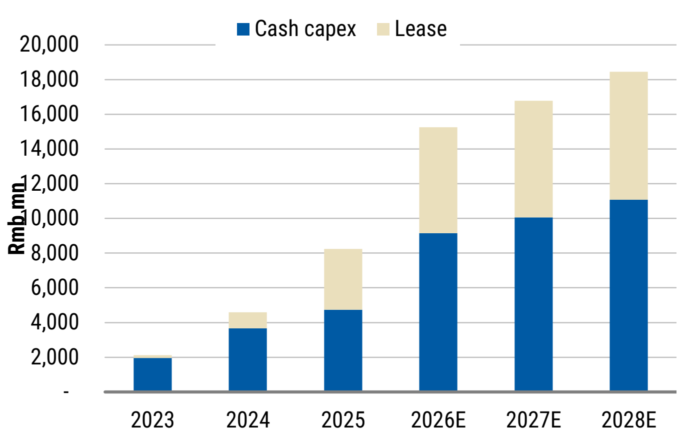

Source: Company data, Morgan Stanley Research estimates.

Exhibit 53: KC new gross asset yield assumptions (bull/base/bear)

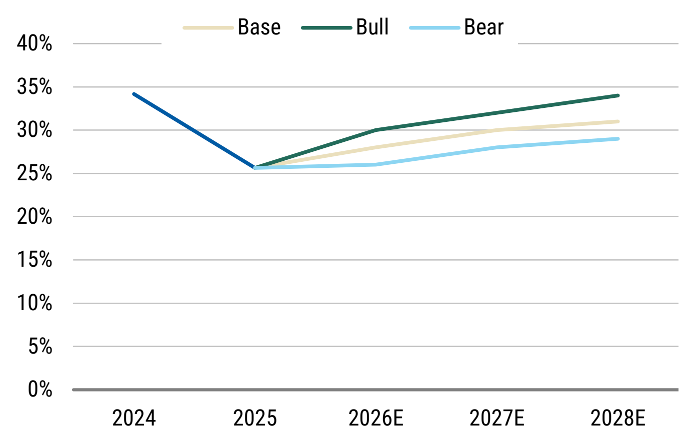

Source: Company data, Morgan Stanley Research estimates.

Exhibit 54: KC revenue forecasts

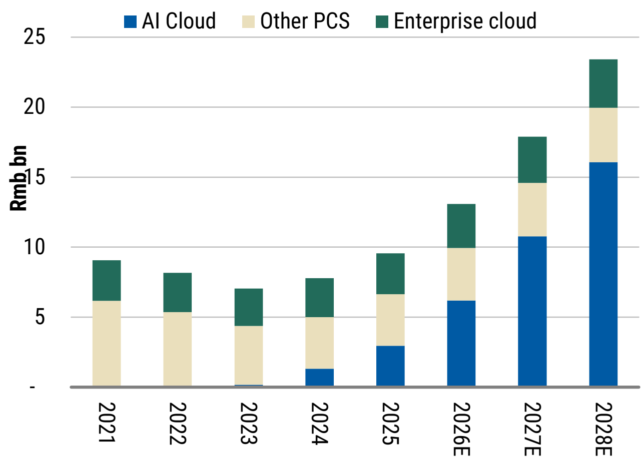

Source: Company data, Morgan Stanley Research estimates.

Exhibit 55: AI revenue growth forecasts

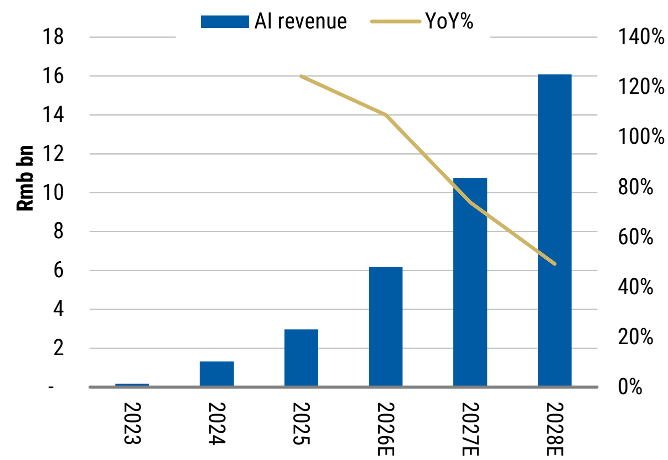

Source: Company data, Morgan Stanley Research estimates.

Exhibit 56: KC gross profit and margin forecasts

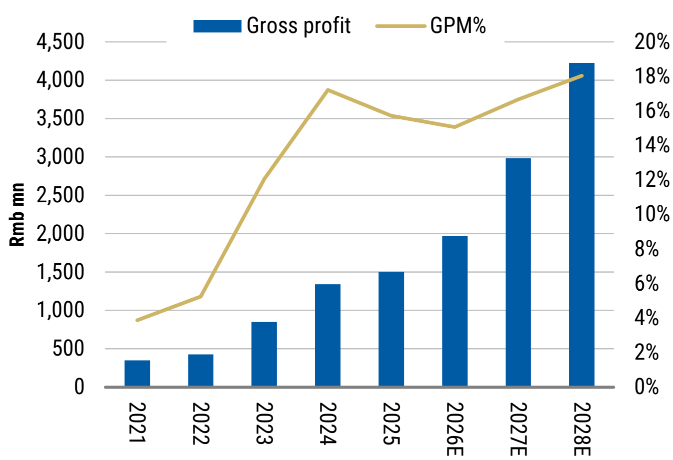

Source: Company data, Morgan Stanley Research estimates.

Exhibit 57: KC adjusted EBITDA and margin forecasts

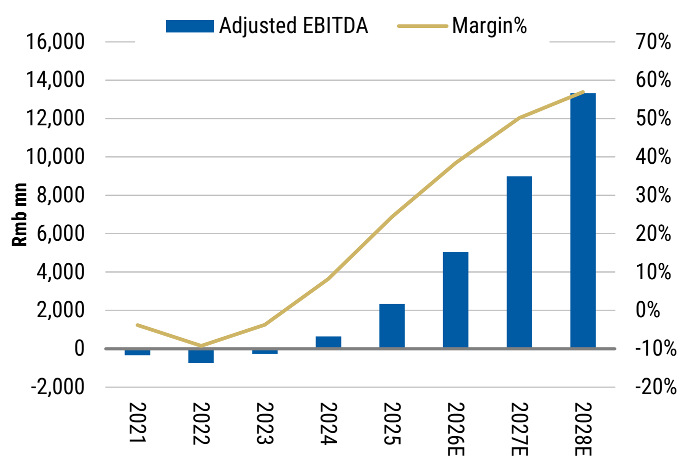

Source: Company data, Morgan Stanley Research estimates.

Exhibit 58: KC adjust operating profit and margin forecasts

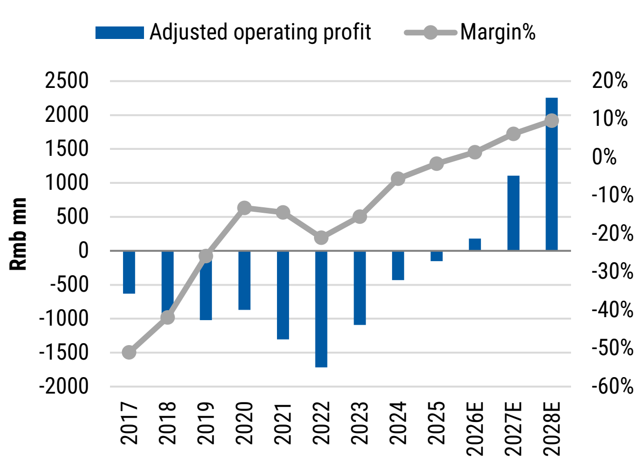

Source: Company data, Morgan Stanley Research estimates.

Exhibit 59: KC adjusted net profit forecasts

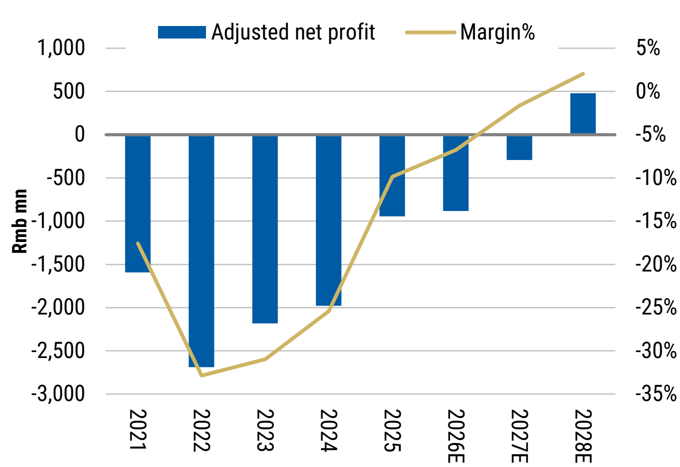

Source: Company data, Morgan Stanley Research estimates.

Valuation methodology

Our price target is US$15 per ADS. We value KC using a targeted EV/EBITDA multiple to account for the large depreciation of the business and in line with global neoclouds such as CoreWeave. Given cloud is an asset-heavy business and operators are adopting different debt leverage ratios and server depreciation policies, we believe EV/EBITDA is better than metrics such as EV/sales or P/E noting that we prefer EBITDA for comparability across companies given differences in depreciation policies and limited EBIT. In our base case, we apply a 2027e EV/EBITDA multiple of 5.5x, at a discount to the US neocloud peer average of 13x and median of 10x. We think the discount is justified by lower asset yields and greater uncertainty around supply chain availability.

Our bull case value of US$27 per ADS is derived from a 2027e EV/EBITDA multiple of 7x.

- We assume a bull case 2027e EBITDA of Rmb10.6bn, 21% above the base case value. This mainly incorporates stronger price hikes on computing power, leading to faster AI revenue growth.

Our bear case value of US$4 per ADS is derived from a 2027e EV/EBITDA multiple of 4x.

- We assume a bear case 2027e EBITDA of Rmb6.7bn, 19% below the base case value. This mainly incorporates an inability to acquire additional computing power and a lower pricing, leading to slower AI revenue growth.

Exhibit 60: KC: EV/EBITDA valuation

Targeted EV/EBITDA			2026E	2027E	2028E
Adjusted EBITDA	Bear	-1.0	4,386	6,697	9,089
	Base	0.0	5,035	8,987	13,333
	Bull	1.0	5,422	10,643	17,438
	Multiple				
Value	Bear	4.0	17,546	26,787	36,356
	Base	5.5	27,695	49,430	73,330
	Bull	7.0	37,954	74,500	122,069
Net cash (debt)			(12,824)	(19,471)	(24,181)
Investment			234	234	234
Value per ADS	Bear		2.0	4.0	6.0
	Base		7.0	**15.0**	24.0
	Bull		13.0	27.0	49.0

Source: Company data, Morgan Stanley Research estimates.

We benchmark KC's valuation to the neocloud players in the US as a primary comp. US neoclouds on average are trading at a EV/S multiple of 15x/7x in 2026-27e and an EV/EBITDA multiple of 36x/13x in 2026-27e. Excluding outlier Digital Ocean, which primarily serves SMB (not very comparable vs. others mainly serving big customers), the average 2027e EV/EBITDA multiple of peers is 6x. KC is trading at a significant discount to the peer average, which we think is partially justified by its US peers having more certainty around acquiring chipsets and a stronger backlog from the hyperscalers.

KC currently trades at a 2027e EV/EBITDA multiple of 4.3x, well below the average of 15x since 2024 (when it started to transit to neocloud). We think this was partially due to a low base in the earlier years, but still looks undervalued compared to its history.

Exhibit 61: Key comps table

Ticker	Name	Rating	CURR	PX	EV/S			EV/EBITDA			3Y Sales	3Y EBITDA
					2026E	2027E	2028E	2026E	2027E	2028E	CAGR	CAGR
KC.O	Kingsoft Cloud	O	USD	9.1	2.4	2.1	1.8	6.3	4.3	3.2	35%	79%
300442.SZ	Range Intelligent	N	CNY	87.8	20.1	16.5	13.7	28.3	23.6	20.0	34%	20%
China - mean					**11.3**	**9.3**	**7.8**	**17.3**	**13.9**	**11.6**	**34%**	**49%**
CRWV.O	CoreWeave	E	USD	81.7	7.2	4.8	3.8	13.5	7.9	5.6	92%	115%
ORCL.N	Oracle	E	USD	140.3	8.6	4.5	NA	15.9	8.0	NA	NA	NA
DOCN.N	Digital Ocean	O	USD	130.1	13.1	8.5	6.1	41.7	24.9	16.3	38%	41%
NBIS.O	Nebius	E	USD	215.6	20.9	8.7	6.4	52.1	14.7	10.0	223%	NA
IREN.O	IREN	N	USD	38.8	22.0	7.3	3.6	55.4	10.2	4.7	125%	155%
US neoclouds					**14.4**	**6.8**	**5.0**	**35.7**	**13.2**	**9.1**	**92%**	**115%**
DLR.N	Digital Realty	E	USD	173.3	11.9	11.2	10.4	23.8	23.3	21.2	10%	44%
EQIX.O	Equinix	O	USD	1002.0	11.7	10.6	9.7	24.4	22.0	20.0	11%	13%
US datacenter					**11.8**	**10.9**	**10.1**	**24.1**	**22.6**	**20.6**	**11%**	**29%**

Note: for non-covered (N) stocks, we use Factset consensus estimate.

Source: FactSet, Morgan Stanley Research estimates. Note: As of the market close on July 2, 2026.

Exhibit 62: Kingsoft Cloud historical EV/EBITDA valuation

Source: FactSet, Morgan Stanley Research.

Exhibit 63: CoreWeave historical EV/EBITDA valuation

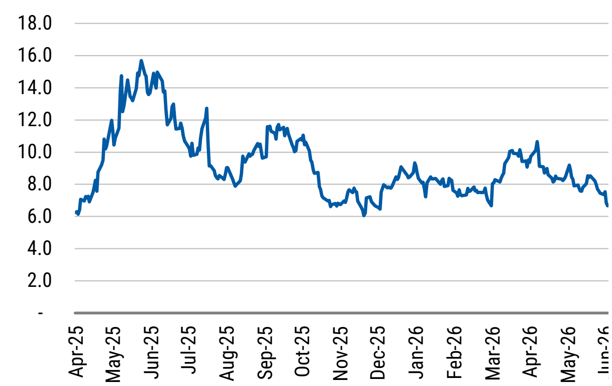

Source: FactSet, Morgan Stanley Research.

Exhibit 64:

Key drivers of Kingsoft Cloud's relative performance since IPO

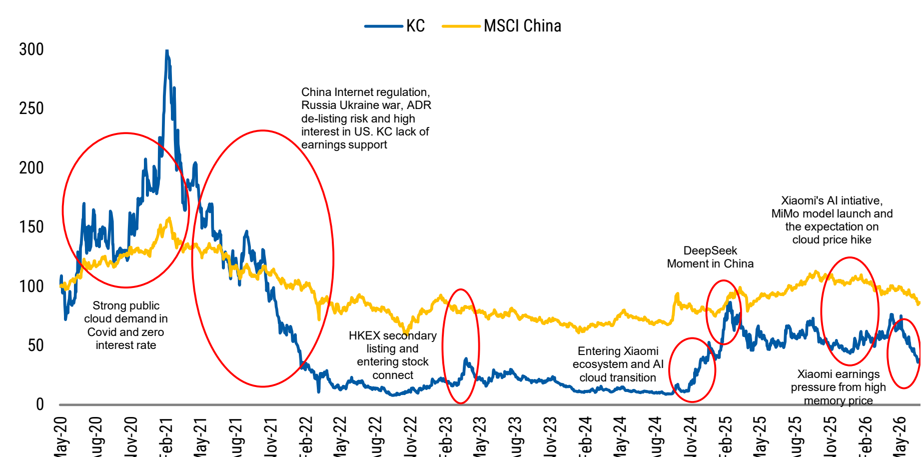

Source: Factset, Morgan Stanley Research

Exhibit 65: KC vs. Xiaomi share price performance: Kingsoft Cloud was widely viewed as a high-beta proxy for Xiaomi

Source: Factset

Our thoughts on the relationship between GPU lifespans and long-term valuation

Following our unit economy analysis above, we assume that in a steady state (no capacity/top-line growth, GPU server passing lifespan, and no more change to the lifespan), neoclouds would need to spend their initial capex divided by GPU lifespan every year to maintain their revenue stream. This will lead to dramatic differences in long-term FCF and EV/EBITDA multiple (assuming different yields). Our conclusion:

- **In the US**, the long-term enterprise value should **expand by 35% and 61%** if the GPU lifespan extends **from five years to six and seven years,** respectively. If staying at **six years, 8.7x steady state EV/EBITDA would be equivalent to 5% yield.**
- **In China,** the long term enterprise value should expand **even more, by 56% and 97%,** if the GPU lifespan extends **from five years to six and seven years,** respectively. If staying at **six years, 7.1x static stage EV/EBITDA would be equivalent to 5% yield.**
- On the contrary, if GPU lifespans were to shorten from five to four years, the US neocloud EV/EBITDA multiple should narrow 53% and the China neocloud EV/EBITDA multiple should narrow 85%. Considering neoclouds' leverage, this would lead to a wiping out of the equity value for neoclouds.
- Due to higher depreciation as % of revenue (GPU accounts for a bigger pie in the value chain), the value change in **China is much more sensitive to GPU lifespan than US. Under the same yield and 6-year GPU lifespan assumptions, China neoclouds should trade at an 18% long-term EV/EBITDA discount vs. US peers.**

The upside risk (longer GPU lifespan) could happen 1) under a computing power shortage scenario, which means legacy GPUs remain popular in the market; and 2) legacy GPU performance improves via software upgrade (*NVIDIA stated in its Q1 FY2026 earnings call: "We've improved Hopper inference performance by 4x in two years. This is the benefit of NVIDIA's programmable CUDA architecture and rich ecosystem."*). The downside risk could happen if the next generation of GPU technology performance, value-for-money and availability are all dozens of times better, and legacy GPUs become obsolete. This is very unlikely to happen, in our view, when inference is taking a bigger share of workload as it can be addressed by lower-end GPUs.

Exhibit 66: US neocloud long-term valuation based on different GPU lifespans

US steady stage				
GPU lifespan	**4 years**	**5 years**	**6 years**	**7 years**
Revenue	20.0	20.0	20.0	20.0
Opex	5.0	5.0	5.0	5.0
EBITDA	15.0	15.0	15.0	15.0
Margin	75%	75%	75%	75%
Recurring capex	12.8	10.2	8.5	7.3
FCF	2.3	4.8	6.5	7.7
Margin	11%	24%	33%	39%
Targeted yield	4.0%	4.0%	4.0%	4.0%
Enterprise value	56.3	120.0	162.5	192.9
EV/EBITDA	3.8	8.0	10.8	12.9
Targeted yield	5.0%	5.0%	5.0%	5.0%
Enterprise value	45.0	96.0	130.0	154.3
EV/EBITDA	3.0	6.4	8.7	10.3
Targeted yield	6.0%	6.0%	6.0%	6.0%
Enterprise value	37.5	80.0	108.3	128.6
EV/EBITDA	2.5	5.3	7.2	8.6

Source: Morgan Stanley Research.

Exhibit 67: China neocloud long-term valuation based on different GPU lifespans

China steady stage				
GPU lifespan	**4 years**	**5 years**	**6 years**	**7 years**
Revenue	22.0	22.0	22.0	22.0
Opex	3.5	3.5	3.5	3.5
EBITDA	18.5	18.5	18.5	18.5
Margin	84%	84%	84%	84%
Recurring capex	17.9	14.3	11.9	10.2
FCF	0.7	4.2	6.6	8.3
Margin	3%	19%	30%	38%
Targeted yield	4.0%	4.0%	4.0%	4.0%
Enterprise value	16.3	105.5	165.0	207.5
EV/EBITDA	0.9	5.7	8.9	11.2
Targeted yield	5.0%	5.0%	5.0%	5.0%
Enterprise value	13.0	84.4	132.0	166.0
EV/EBITDA	0.7	4.6	7.1	9.0
Targeted yield	6.0%	6.0%	6.0%	6.0%
Enterprise value	10.8	70.3	110.0	138.3
EV/EBITDA	0.6	3.8	5.9	7.5

Source: Morgan Stanley Research.

Key risks to our price target

Risks to the upside

- Stronger-than-expected demand for high-end GPU
- Lower-than-expected interest rate driving better return profile
- Better-than-expected development of Xiaomi's AI models

Risks to the downside

- Supply-side restrictions leading to procurement shortfall
- Higher-than-expected interest rates leading to higher cost of funding
- Slower-than-expected AI model developments in China

This report references U.S. Executive Order 14032 and/or entities or securities that are designated thereunder. U.S. persons may be prohibited from buying certain securities of entities named in this report. Readers are solely responsible for ensuring that their investment activities are carried out in compliance with applicable laws.

This report references U.S. Executive Order 14105 and/or entities that may be in scope of such order. U.S. persons may be prohibited from engaging in certain transactions or otherwise require certain other transactions be notified to the U.S. Department of Treasury. Readers are solely responsible for ensuring that their investment or trade activities are carried out in compliance with applicable laws.

This report references export controls and/or entities that may be subject to export control restrictions. Readers are solely responsible for ensuring that their investment or trade activities are carried out in compliance with applicable laws.

Risk Reward Reference links

1. View explanation of Options Probabilities methodology - Options_Probabilities_Exhibit_Link.pdf

2. View descriptions of Risk Rewards Themes - RR_Themes_Exhibit_Link.pdf

3. View explanation of regional hierarchies - GEG_Exhibit_Link.pdf

4. View explanation of Theme/Exposure methodology - ESG_Sustainable_Solutions_External_Link.pdf

5. View explanation of HERS methodology - ESG_HERS_External_Link.pdf

Disclosure Section

The information and opinions in Morgan Stanley Research were prepared or are disseminated by Morgan Stanley Asia Limited (which accepts the responsibility for its contents) and/or Morgan Stanley Asia (Singapore) Pte. (Registration number 199206298Z) and/or Morgan Stanley Asia (Singapore) Securities Pte Ltd (Registration number 200008434H), regulated by the Monetary Authority of Singapore (which accepts legal responsibility for its contents and should be contacted with respect to any matters arising from, or in connection with, Morgan Stanley Research), and/or Morgan Stanley Taiwan Limited and/or Morgan Stanley & Co International plc, Seoul Branch, and/or Morgan Stanley Australia Limited (A.B.N. 67 003 734 576, holder of Australian financial services license No. 233742, which accepts responsibility for its contents), and/or Morgan Stanley Wealth Management Australia Pty Ltd (A.B.N. 19 009 145 555, holder of Australian financial services license No. 240813, which accepts responsibility for its contents), and/or Morgan Stanley India Company Private Limited having Corporate Identification No (CIN) U22990MH1998PTC115305, regulated by the Securities and Exchange Board of India ("SEBI") and holder of licenses as a Research Analyst (SEBI Registration No. INH000001105); Stock Broker (SEBI Stock Broker Registration No. INZ000244438), Merchant Banker (SEBI Registration No. INM000011203), and depository participant with National Securities Depository Limited (SEBI Registration No. IN-DP-NSDL-567-2021) having registered office at Altimus, Level 39 & 40, Pandurang Budhkar Marg, Worli, Mumbai 400018, India; Telephone no. +91-22-61181000; Compliance Officer Details: Mr. Tejarshi Hardas, Tel. No.: +91-22-61181000 or Email: tejarshi.hardas@morganstanley.com; Grievance officer details: Mr. Tejarshi Hardas, Tel. No.: +91-22-61181000 or Email: msic-compliance@morganstanley.com which accepts the responsibility for its contents and should be contacted with respect to any matters arising from, or in connection with, Morgan Stanley Research, and their affiliates (collectively, "Morgan Stanley"). Morgan Stanley India Company Private Limited (MSICPL) may use AI tools in providing research services. All recommendations contained herein are made by the duly qualified research analysts.

For important disclosures, stock price charts and equity rating histories regarding companies that are the subject of this report, please see the Morgan Stanley Research Disclosure Website at www.morganstanley.com/eqr/disclosures/webapp/generalresearch, or contact your investment representative or Morgan Stanley Research at 1585 Broadway, (Attention: Research Management), New York, NY, 10036 USA.

For valuation methodology and risks associated with any recommendation, rating or price target referenced in this research report, please contact the Client Support Team as follows: US/Canada +1 800 303-2495; Hong Kong +852 2848-5999; Latin America +1 718 754-5444 (U.S.); London +44 (0)20-7425-8169; Singapore +65 6834-6860; Sydney +61 (0)2-9770-1505; Tokyo +81 (0)3-6836-9000. Alternatively you may contact your investment representative or Morgan Stanley Research at 1585 Broadway, (Attention: Research Management), New York, NY 10036 USA.

Analyst Certification

The following analysts hereby certify that their views about the companies and their securities discussed in this report are accurately expressed and that they have not received and will not receive direct or indirect compensation in exchange for expressing specific recommendations or views in this report: Yang Liu; Andy Meng, CFA; Tom Tang; Gary Yu.

.

Global Research Conflict Management Policy

Morgan Stanley Research has been published in accordance with our conflict management policy, which is available at www.morganstanley.com/institutional/research/conflictpolicies. A Portuguese version of the policy can be found at www.morganstanley.com.br

Important Regulatory Disclosures on Subject Companies

As of May 29, 2026, Morgan Stanley beneficially owned 1% or more of a class of common equity securities of the following companies covered in Morgan Stanley Research: CoreWeave, Digital Realty Trust Inc., Equinix Inc., Kingdee International Software Group, Meitu Inc, Nebius Group NV, Oracle Corporation.

Within the last 12 months, Morgan Stanley managed or co-managed a public offering (or 144A offering) of securities of CoreWeave, Digital Realty Trust Inc., DigitalOcean Holdings Inc, Equinix Inc., **Kingsoft Cloud Holdings**, Kingsoft Corp Ltd, MiniMax, Nebius Group NV.

Within the last 12 months, Morgan Stanley has received compensation for investment banking services from CoreWeave, Digital Realty Trust Inc., DigitalOcean Holdings Inc, Equinix Inc., **Kingsoft Cloud Holdings**, Kingsoft Corp Ltd, MiniMax, Nebius Group NV.

In the next 3 months, Morgan Stanley expects to receive or intends to seek compensation for investment banking services from Beijing Kingsoft Office Software Inc, Chinasoft International Ltd, CoreWeave, Digital Realty Trust Inc., DigitalOcean Holdings Inc, Equinix Inc., Glodon Co. Ltd., Hundsun Technologies Inc., Kingdee International Software Group, **Kingsoft Cloud Holdings**, Kingsoft Corp Ltd, Knowledge Atlas Technology JSC Ltd, Meitu Inc, MiniMax, Nebius Group NV, Oracle Corporation, Tuya Inc., Yonyou Network Technology Co Ltd.

Within the last 12 months, Morgan Stanley has received compensation for products and services other than investment banking services from CoreWeave, Digital Realty Trust Inc., DigitalOcean Holdings Inc, Equinix Inc., Kingsoft Corp Ltd, Oracle Corporation.

Within the last 12 months, Morgan Stanley has provided or is providing investment banking services to, or has an investment banking client relationship with, the following company: Beijing Kingsoft Office Software Inc, Chinasoft International Ltd, CoreWeave, Digital Realty Trust Inc., DigitalOcean Holdings Inc, Equinix Inc., Glodon Co. Ltd., Hundsun Technologies Inc., Kingdee International Software Group, **Kingsoft Cloud Holdings**, Kingsoft Corp Ltd, Knowledge Atlas Technology JSC Ltd, Meitu Inc, MiniMax, Nebius Group NV, Oracle Corporation, Tuya Inc., Yonyou Network Technology Co Ltd.

Within the last 12 months, Morgan Stanley has either provided or is providing non-investment banking, securities-related services to and/or in the past has entered into an agreement to provide services or has a client relationship with the following company: Beisen Holding Limited, CoreWeave, Digital Realty Trust Inc., DigitalOcean Holdings Inc, Equinix Inc., **Kingsoft Cloud Holdings**, Kingsoft Corp Ltd, Nebius Group NV, Oracle Corporation, Tuya Inc..

The equity research analysts or strategists principally responsible for the preparation of Morgan Stanley Research have received compensation based upon various factors, including quality of research, investor client feedback, stock picking, competitive factors, firm revenues and overall investment banking revenues. Equity Research analysts' or strategists' compensation is not linked to investment banking or capital markets transactions performed by Morgan Stanley or the profitability or revenues of particular trading desks.

Morgan Stanley and its affiliates do business that relates to companies/instruments covered in Morgan Stanley Research, including market making, providing liquidity, fund management, commercial banking, extension of credit, investment services and investment banking. Morgan Stanley sells to and buys from customers the securities/instruments of companies covered in Morgan Stanley Research on a principal basis. Morgan Stanley may have a position in the debt of the Company or instruments discussed in this report. Morgan Stanley trades or may trade as principal in the debt securities (or in related derivatives) that are the subject of the debt research report.

Certain disclosures listed above are also for compliance with applicable regulations in non-US jurisdictions.

STOCK RATINGS

Morgan Stanley uses a relative rating system using terms such as Overweight, Equal-weight, Not-Rated or Underweight (see definitions below). Morgan Stanley does not assign ratings of Buy, Hold or Sell to the stocks we cover. Overweight, Equal-weight, Not-Rated and Underweight are not the equivalent of buy, hold and sell. Investors should carefully read the definitions of all ratings used in Morgan Stanley Research. In addition, since Morgan Stanley Research contains more complete information concerning the analyst's views, investors should carefully read Morgan

Stanley Research, in its entirety, and not infer the contents from the rating alone. In any case, ratings (or research) should not be used or relied upon as investment advice. An investor's decision to buy or sell a stock should depend on individual circumstances (such as the investor's existing holdings) and other considerations.

Global Stock Ratings Distribution

(as of June 30, 2026)

The Stock Ratings described below apply to Morgan Stanley's Fundamental Equity Research and do not apply to Debt Research produced by the Firm.

For disclosure purposes only (in accordance with FINRA requirements), we include the category headings of Buy, Hold, and Sell alongside our ratings of Overweight, Equal-weight, Not-Rated and Underweight. Morgan Stanley does not assign ratings of Buy, Hold or Sell to the stocks we cover. Overweight, Equal-weight, Not-Rated and Underweight are not the equivalent of buy, hold, and sell but represent recommended relative weightings (see definitions below). To satisfy regulatory requirements, we correspond Overweight, our most positive stock rating, with a buy recommendation; we correspond Equal-weight and Not-Rated to hold and Underweight to sell recommendations, respectively.

	Coverage Universe		Investment Banking Clients (IBC)			Other Material Investment Services Clients (MISC)	
Stock Rating Category	Count	% of Total	Count	% of Total IBC	% of Rating Category	Count	% of Total Other MISC
Overweight/Buy	1544	42%	453	49%	29%	757	44%
Equal-weight/Hold	1577	43%	390	42%	25%	769	44%
Not-Rated/Hold	3	0%	1	0%	33%	1	0%
Underweight/Sell	544	15%	89	10%	16%	204	12%
Total	3,668		933			1731	

Data include common stock and ADRs currently assigned ratings. Investment Banking Clients are companies from whom Morgan Stanley received investment banking compensation in the last 12 months. Due to rounding off of decimals, the percentages provided in the "% of total" column may not add up to exactly 100 percent.

Analyst Stock Ratings

Overweight (O). The stock's total return is expected to exceed the average total return of the analyst's industry (or industry team's) coverage universe, on a risk-adjusted basis, over the next 12-18 months.

Equal-weight (E). The stock's total return is expected to be in line with the average total return of the analyst's industry (or industry team's) coverage universe, on a risk-adjusted basis, over the next 12-18 months.

Not-Rated (NR). Currently the analyst does not have adequate conviction about the stock's total return relative to the average total return of the analyst's industry (or industry team's) coverage universe, on a risk-adjusted basis, over the next 12-18 months.

Underweight (U). The stock's total return is expected to be below the average total return of the analyst's industry (or industry team's) coverage universe, on a risk-adjusted basis, over the next 12-18 months.

Unless otherwise specified, the time frame for price targets included in Morgan Stanley Research is 12 to 18 months.

Analyst Industry Views

Attractive (A): The analyst expects the performance of his or her industry coverage universe over the next 12-18 months to be attractive vs. the relevant broad market benchmark, as indicated below.

In-Line (I): The analyst expects the performance of his or her industry coverage universe over the next 12-18 months to be in line with the relevant broad market benchmark, as indicated below.

Cautious (C): The analyst views the performance of his or her industry coverage universe over the next 12-18 months with caution vs. the relevant broad market benchmark, as indicated below.

Benchmarks for each region are as follows: North America - S&P 500; Latin America - relevant MSCI country index or MSCI Latin America Index; Europe - MSCI Europe; Japan - TOPIX; Asia - relevant MSCI country index or MSCI sub-regional index or MSCI AC Asia Pacific ex Japan Index.

Important Disclosures for Morgan Stanley Smith Barney LLC Customers

Important disclosures regarding any material conflict of interest that can reasonably be expected to have influenced Morgan Stanley Smith Barney LLC's choice of a third-party research provider or the subject company of a third-party research report, are available on the Morgan Stanley Wealth Management disclosure website at www.morganstanley.com/online/researchdisclosures.

For Morgan Stanley specific disclosures, you may refer to https://www.morganstanley.com/eqr/disclosures/webapp/generalresearch.

Each Morgan Stanley research report is reviewed and approved on behalf of Morgan Stanley Smith Barney LLC. This review and approval is conducted by the same person who reviews the research report on behalf of Morgan Stanley. This could create a conflict of interest.

Other Important Disclosures

Morgan Stanley Research policy is to update research reports as and when the Research Analyst and Research Management deem appropriate, based on developments with the issuer, the sector, or the market that may have a material impact on the research views or opinions stated therein. In addition, certain Research publications are intended to be updated on a regular periodic basis (weekly/monthly/quarterly/annual) and will ordinarily be updated with that frequency, unless the Research Analyst and Research Management determine that a different publication schedule is appropriate based on current conditions.

Morgan Stanley is not acting as a municipal advisor and the opinions or views contained herein are not intended to be, and do not constitute, advice within the meaning of Section 975 of the Dodd-Frank Wall Street Reform and Consumer Protection Act.

Morgan Stanley produces an equity research product called a "Tactical Idea." Views contained in a "Tactical Idea" on a particular stock may be contrary to the recommendations or views expressed in research on the same stock. This may be the result of differing time horizons, methodologies, market events, or other factors. For all research available on a particular stock, please contact your sales representative or go to Matrix at http://www.morganstanley.com/matrix.

Morgan Stanley Research is provided to our clients through our proprietary research portal on Matrix and also distributed electronically by Morgan Stanley to clients. Certain, but not all, Morgan

Stanley Research products are also made available to clients through third-party vendors or redistributed to clients through alternate electronic means as a convenience. For access to all available Morgan Stanley Research, please contact your sales representative or go to Matrix at http://www.morganstanley.com/matrix.

Any access and/or use of Morgan Stanley Research is subject to Morgan Stanley's Terms of Use (http://www.morganstanley.com/terms.html). By accessing and/or using Morgan Stanley Research, you are indicating that you have read and agree to be bound by our Terms of Use (http://www.morganstanley.com/terms.html). In addition you consent to Morgan Stanley processing your personal data and using cookies in accordance with our Privacy Policy and our Global Cookies Policy (http://www.morganstanley.com/privacy_pledge.html), including for the purposes of setting your preferences and to collect readership data so that we can deliver better and more personalized service and products to you. To find out more information about how Morgan Stanley processes personal data, how we use cookies and how to reject cookies see our Privacy Policy and our Global Cookies Policy (http://www.morganstanley.com/privacy_pledge.html). Please use the provided link to review the Terms and Conditions and Most Important Terms and Conditions for Morgan Stanley India Company Private Limited (https://www.morganstanley.com/assets/pdfs/about-us-global-offices/india/Terms_and_conditions.pdf) and the following link to review the audit report (https://ny.matrix.ms.com/eqr/research/webapp/researchdocs/MSICPL_Morgan_Stanley_Research_Audit_Report.pdf).

If you do not agree to our Terms of Use and/or if you do not wish to provide your consent to Morgan Stanley processing your personal data or using cookies please do not access our research.

Morgan Stanley Research does not provide individually tailored investment advice. Morgan Stanley Research has been prepared without regard to the circumstances and objectives of those who receive it. Morgan Stanley recommends that investors independently evaluate particular investments and strategies, and encourages investors to seek the advice of a financial adviser. The appropriateness of an investment or strategy will depend on an investor's circumstances and objectives. The securities, instruments, or strategies discussed in Morgan Stanley Research may not be suitable for all investors, and certain investors may not be eligible to purchase or participate in some or all of them. Morgan Stanley Research is not an offer to buy or sell or the solicitation of an offer to buy or sell any security/instrument or to participate in any particular trading strategy. The value of and income from your investments may vary because of changes in interest rates, foreign exchange rates, default rates, prepayment rates, securities/instruments prices, market indexes, operational or financial conditions of companies or other factors. There may be time limitations on the exercise of options or other rights in securities/instruments transactions. Past performance is not necessarily a guide to future performance. Estimates of future performance are based on assumptions that may not be realized. If provided, and unless otherwise stated, the closing price on the cover page is that of the primary exchange for the subject company's securities/instruments.

The fixed income research analysts, strategists or economists principally responsible for the preparation of Morgan Stanley Research have received compensation based upon various factors, including quality, accuracy and value of research, firm profitability or revenues (which include fixed income trading and capital markets profitability or revenues), client feedback and competitive factors. Fixed Income Research analysts', strategists' or economists' compensation is not linked to investment banking or capital markets transactions performed by Morgan Stanley or the profitability or revenues of particular trading desks.

The "Important Regulatory Disclosures on Subject Companies" section in Morgan Stanley Research lists all companies mentioned where Morgan Stanley owns 1% or more of a class of common equity securities of the companies. For all other companies mentioned in Morgan Stanley Research, Morgan Stanley may have an investment of less than 1% in securities/instruments or derivatives of securities/instruments of companies and may trade them in ways different from those discussed in Morgan Stanley Research. Employees of Morgan Stanley not involved in the preparation of Morgan Stanley Research may have investments in securities/instruments or derivatives of securities/instruments of companies mentioned and may trade them in ways different from those discussed in Morgan Stanley Research. Derivatives may be issued by Morgan Stanley or associated persons.

With the exception of information regarding Morgan Stanley, Morgan Stanley Research is based on public information. Morgan Stanley makes every effort to use reliable, comprehensive information, but we make no representation that it is accurate or complete. We have no obligation to tell you when opinions or information in Morgan Stanley Research change apart from when we intend to discontinue equity research coverage of a subject company. Facts and views presented in Morgan Stanley Research have not been reviewed by, and may not reflect information known to, professionals in other Morgan Stanley business areas, including investment banking personnel.

Morgan Stanley Research personnel may participate in company events such as site visits and are generally prohibited from accepting payment by the company of associated expenses unless pre-approved by authorized members of Research management.

Morgan Stanley may make investment decisions that are inconsistent with the recommendations or views in this report.

To our readers based in Taiwan or trading in Taiwan securities/instruments: Information on securities/instruments that trade in Taiwan is distributed by Morgan Stanley Taiwan Limited ("MSTL"). Such information is for your reference only. The reader should independently evaluate the investment risks and is solely responsible for their investment decisions. Morgan Stanley Research may not be distributed to the public media or quoted or used by the public media without the express written consent of Morgan Stanley. Any non-customer reader within the scope of Article 7-1 of the Taiwan Stock Exchange Recommendation Regulations accessing and/or receiving Morgan Stanley Research is not permitted to provide Morgan Stanley Research to any third party (including but not limited to related parties, affiliated companies and any other third parties) or engage in any activities regarding Morgan Stanley Research which may create or give the appearance of creating a conflict of interest. Information on securities/instruments that do not trade in Taiwan is for informational purposes only and is not to be construed as a recommendation or a solicitation to trade in such securities/instruments. MSTL may not execute transactions for clients in these securities/instruments.

Certain information in Morgan Stanley Research was sourced by employees of the Shanghai Representative Office of Morgan Stanley Asia Limited for the use of Morgan Stanley Asia Limited. Morgan Stanley is not incorporated under PRC law and the research in relation to this report is conducted outside the PRC. Morgan Stanley Research does not constitute an offer to sell or the solicitation of an offer to buy any securities in the PRC. PRC investors shall have the relevant qualifications to invest in such securities and shall be responsible for obtaining all relevant approvals, licenses, verifications and/or registrations from the relevant governmental authorities themselves. Neither this report nor any part of it is intended as, or shall constitute, provision of any consultancy or advisory service of securities investment as defined under PRC law. Such information is provided for your reference only.

Morgan Stanley Research is disseminated in Brazil by Morgan Stanley C.T.V.M. S.A. located at Av. Brigadeiro Faria Lima, 3600, 6th floor, São Paulo - SP, Brazil; and is regulated by the Comissão de Valores Mobiliários; in Mexico by Morgan Stanley México, Casa de Bolsa, S.A. de C.V which is regulated by Comision Nacional Bancaria y de Valores. Paseo de los Tamarindos 90, Torre 1, Col. Bosques de las Lomas Floor 29, 05120 Mexico City; in Japan by Morgan Stanley MUFG Securities Co., Ltd. and, for Commodities related research reports only, Morgan Stanley Capital Group Japan Co., Ltd; in Hong Kong by Morgan Stanley Asia Limited (which accepts responsibility for its contents) and by Morgan Stanley Bank Asia Limited; in Singapore by Morgan Stanley Asia (Singapore) Pte. (Registration number 199206298Z) and/or Morgan Stanley Asia (Singapore) Securities Pte Ltd (Registration number 200008434H), regulated by the Monetary Authority of Singapore (which accepts legal responsibility for its contents and should be contacted with respect to any matters arising from, or in connection with, Morgan Stanley Research) and by Morgan Stanley Bank Asia Limited, Singapore Branch (Registration number T14FC0118J); in Australia to "wholesale clients" within the meaning of the Australian Corporations Act by Morgan Stanley Australia Limited A.B.N. 67 003 734 576, holder of Australian financial services license No. 233742, which accepts responsibility for its contents; in Australia to "wholesale clients" and "retail clients" within the meaning of the Australian Corporations Act by Morgan Stanley Wealth Management Australia Pty Ltd (A.B.N. 19 009 145 555, holder of Australian financial services license No. 240813, which accepts responsibility for its contents; in Korea by Morgan Stanley & Co International plc, Seoul Branch; in India by Morgan Stanley India Company Private Limited having Corporate Identification No (CIN) U22990MH1998PTC115305, regulated by the Securities and Exchange Board of India ("SEBI") and holder of licenses as a Research Analyst (SEBI Registration No. INH000001105); Stock Broker (SEBI Stock Broker Registration No. INZ000244438), Merchant Banker (SEBI Registration No. INM000011203), and depository participant with National Securities Depository Limited (SEBI Registration No. IN-DP-NSDL-567-2021) having registered office at Altimus, Level 39 & 40, Pandurang Budhkar Marg, Worli, Mumbai 400018, India;

Telephone no. +91-22-61181000; Compliance Officer Details: Mr. Tejarshi Hardas, Tel. No.: +91-22-61181000 or Email: tejarshi.hardas@morganstanley.com; Grievance officer details: Mr. Tejarshi Hardas, Tel. No.: +91-22-61181000 or Email: msic-compliance@morganstanley.com. Morgan Stanley India Company Private Limited (MSICPL) may use AI tools in providing research services. All recommendations contained herein are made by the duly qualified research analysts; in Canada by Morgan Stanley Canada Limited; in Germany and the European Economic Area where required by Morgan Stanley Europe S.E., authorised and regulated by Bundesanstalt fuer Finanzdienstleistungsaufsicht (BaFin) under the reference number 149169; in the US by Morgan Stanley & Co. LLC, which accepts responsibility for its contents. Morgan Stanley & Co. International plc, authorized by the Prudential Regulation Authority and regulated by the Financial Conduct Authority and the Prudential Regulation Authority, disseminates in the UK research that it has prepared, and research which has been prepared by any of its affiliates, only to persons who (i) are investment professionals falling within Article 19(5) of the Financial Services and Markets Act 2000 (Financial Promotion) Order 2005 (as amended, the "Order"); (ii) are persons who are high net worth entities falling within Article 49(2)(a) to (d) of the Order; or (iii) are persons to whom an invitation or inducement to engage in investment activity (within the meaning of section 21 of the Financial Services and Markets Act 2000, as amended) may otherwise lawfully be communicated or caused to be communicated. RMB Morgan Stanley Proprietary Limited is a member of the JSE Limited and A2X (Pty) Ltd. RMB Morgan Stanley Proprietary Limited is a joint venture owned equally by Morgan Stanley International Holdings Inc. and RMB Investment Advisory (Proprietary) Limited, which is wholly owned by FirstRand Limited. The information in Morgan Stanley Research is being disseminated by Morgan Stanley Saudi Arabia, regulated by the Capital Market Authority in the Kingdom of Saudi Arabia, and is directed at Sophisticated investors only.

Morgan Stanley Hong Kong Securities Limited is the liquidity provider/market maker for securities of Chinasoft International Ltd, Kingdee International Software Group, **Kingsoft Cloud Holdings**, Kingsoft Corp Ltd, Knowledge Atlas Technology JSC Ltd, Meitu Inc, MiniMax listed on the Stock Exchange of Hong Kong Limited. An updated list can be found on HKEx website: http://www.hkex.com.hk.

The information in Morgan Stanley Research is being communicated by Morgan Stanley & Co. International plc (DIFC Branch), regulated by the Dubai Financial Services Authority (the DFSA) or by Morgan Stanley & Co. International plc (ADGM Branch), regulated by the Financial Services Regulatory Authority Abu Dhabi (the FSRA), and is directed at Professional Clients only, as defined by the DFSA or the FSRA, respectively. The financial products or financial services to which this research relates will only be made available to a customer who we are satisfied meets the regulatory criteria of a Professional Client. A distribution of the different MS Research ratings or recommendations, in percentage terms for Investments in each sector covered, is available upon request from your sales representative.

The information in Morgan Stanley Research is being communicated by Morgan Stanley & Co. International plc (QFC Branch), regulated by the Qatar Financial Centre Regulatory Authority (the QFCRA), and is directed at business customers and market counterparties only and is not intended for Retail Customers as defined by the QFCRA.

As required by the Capital Markets Board of Turkey, investment information, comments and recommendations stated here, are not within the scope of investment advisory activity. Investment advisory service is provided exclusively to persons based on their risk and income preferences by the authorized firms. Comments and recommendations stated here are general in nature. These opinions may not fit to your financial status, risk and return preferences. For this reason, to make an investment decision by relying solely to this information stated here may not bring about outcomes that fit your expectations.

The trademarks and service marks contained in Morgan Stanley Research are the property of their respective owners. Third-party data providers make no warranties or representations relating to the accuracy, completeness, or timeliness of the data they provide and shall not have liability for any damages relating to such data. The Global Industry Classification Standard (GICS) was developed by and is the exclusive property of MSCI and S&P.

Morgan Stanley Research, or any portion thereof may not be reprinted, sold or redistributed without the written consent of Morgan Stanley.

Indicators and trackers referenced in Morgan Stanley Research may not be used as, or treated as, a benchmark under Regulation EU 2016/1011, or any other similar framework.

The issuers and/or fixed income products recommended or discussed in certain fixed income research reports may not be continuously followed. Accordingly, investors should regard those fixed income research reports as providing stand-alone analysis and should not expect continuing analysis or additional reports relating to such issuers and/or individual fixed income products. Morgan Stanley may hold, from time to time, material financial and commercial interests regarding the company subject to the Research report.

Registration granted by SEBI and certification from the National Institute of Securities Markets (NISM) in no way guarantee performance of the intermediary or provide any assurance of returns to investors. Investment in securities market are subject to market risks. Read all the related documents carefully before investing.

INDUSTRY COVERAGE: Greater China IT Services and Software

COMPANY (TICKER)	RATING (AS OF)	PRICE* (07/03/2026)
Gary Yu		
Knowledge Atlas Technology JSC Ltd (2513.HK)	O (02/20/2026)	HK$1,793.00
MiniMax (0100.HK)	O (02/20/2026)	HK$346.60
Lydia Lin		
Beisen Holding Limited (9669.HK)	O (06/23/2025)	HK$3.53
Chinasoft International Ltd (0354.HK)	U (05/04/2026)	HK$3.08
Glodon Co. Ltd. (002410.SZ)	E (08/23/2023)	Rmb8.38
Hundsun Technologies Inc. (600570.SS)	U (04/07/2025)	Rmb21.23
Longshine Technology Group Co Ltd (300682.SZ)	U (09/17/2025)	Rmb10.06
Meitu Inc (1357.HK)	O (10/01/2024)	HK$4.00
Yang Liu		
Beijing Kingsoft Office Software Inc (688111.SS)	U (08/26/2024)	Rmb211.35
Kingdee International Software Group (0268.HK)	E (03/19/2025)	HK$6.35
Kingsoft Cloud Holdings (KC.O)	O (07/06/2026)	US$9.10
Kingsoft Corp Ltd (3888.HK)	E (07/21/2025)	HK$22.92
Sangfor Technologies Inc (300454.SZ)	E (07/05/2024)	Rmb92.04

Tuya Inc. (TUYA.N)	O (11/29/2023)	US$1.76
Wangsu Science & Technology (300017.SZ)	U (11/12/2024)	Rmb12.73
Yonyou Network Technology Co Ltd (600588.SS)	U (11/12/2024)	Rmb9.31

Stock Ratings are subject to change. Please see latest research for each company.
* Historical prices are not split adjusted.

© 2026 Morgan Stanley