

China Artificial Intelligence

Open-weight commercialization playbook: winner takes more; Raise Zhipu to HK\$2,000 (OW), cut MiniMax to HK\$300 (Neutral)

Investors typically read open-weight releases as monetization leakage, since public weights allow third parties to deploy models outside the first-party API route. We agree, but we believe the **more crucial question is whether LLM providers can turn reduced access control into wider distribution and paid conversion. In our view, the answer depends heavily on model capability** – for competitive models, open weights can scale adoption on external GPU capacity across CSP, inference platforms and private deployments, rather than only on the provider's own compute stack. The LLM provider can still participate and monetize through official-backed routes, licensing, marketplace SaaS, revenue share, deployment support or enterprise co-selling, especially when customers care about model freshness, endpoint quality and reliability. For weaker models, reduced access control can lead to faster price comparison and traffic rerouting. This is positive for Zhipu, given GLM-5.2's globally competitive positioning, while more challenging for MiniMax, as M3 has not yet shown enough differentiation and pricing power. **Open-weight commercialization is therefore becoming a 'winner-takes-more' framework.**

- **Partner channels such as CSPs and inference platforms still have incentives to work with open-weight LLM providers because commercial customers are buying a production-grade model service, rather than just access to weights.** A public release is closer to a starting checkpoint, while the official route can keep improving through post-launch model updates and serving optimization. If third-party deployments lag the official route, the same model name can deliver sizably different user experiences across endpoints. For stronger models, customer demand gives partner channels more reasons to support official-backed access, even when self-deployment offers better gross margins. This allows the model provider to keep a role in the economics of external compute capacity.
- **Recent launches fit this framework.** Zhipu's GLM-5.2 release looks like a more confident open-weight strategy: widen adoption through permissive access, while keeping official route and higher-service variants (e.g. GLM-Turbo series) positioned for quality-sensitive demand. MiniMax adopts a similar strategy by keeping model open-sourced with additional commercialization restrictions, however, the key issue is model competitiveness, in our view. If users view the model as replaceable, open access makes it easier to compare, route and substitute; if the model is strong, open access expands reach and creates more routes back into paid usage. From enterprise users' perspective, open weights can also reduce adoption and accessibility risk, as users have a fallback route if hosted API pricing, access policy or availability changes.
- **We believe this framework provides sizeable optionality for Zhipu, while the revenue path still depends on sustained model leadership.** At the current valuation, we think the market has largely moved beyond the company's US\$1bn year-end ARR guidance, while the broader open-weight monetization path remains less reflected. The key tests are whether GLM-5.2 represents a company-specific capability step-up, how Kimi K3 and DeepSeek V4.1 compare, and whether GLM-5.5/6 can extend the gap. We raise our Dec-26 PT for Zhipu AI to HK\$2,000 and reiterate OW. For MiniMax, the recent launch provides weaker evidence of model-led pricing power and paid conversion under this framework – we lower our Dec-26 PT to HK\$300 and maintain Neutral.

Hong Kong Internet

Olivia Xu ^{AC}

(86-21) 6106 6138

olivia.w.xu@jpmorgan.com

SAC Registration Number: S1730525060001

Alex Yao

(86-21) 6106 6505

alex.yao@jpmorgan.com

SAC Registration Number: S1730523020001

Daniel Chen

(86-21) 6106 6205

daniel.q.chen@jpmorgan.com

SAC Registration Number: S1730521040001

J.P. Morgan Securities (China) Company
Limited

See page 21 for analyst certification and important disclosures, including non-US analyst disclosures.

Equity Ratings and Price Targets

Company	Ticker	Mkt Cap (\$ mn)	Price CCY	Price	Rating		Price Target			
					Cur	Prev	Cur	End Date	Prev	End Date
Zhipu AI	2513 HK	45,431	HKD	1,610.00	OW	n/c	2,000.00	Dec-26	1,800.00	n/c
MiniMax Group Inc - H	100 HK	9,419	HKD	323.80	N	n/c	300.00	Dec-26	400.00	n/c

Source: Company data, Bloomberg Finance L.P., J.P. Morgan estimates. n/c = no change. All prices as of 07 Jul 26.

Open-weight strategy: distribution upside, model-access pressure and a more flexible monetization playbook

Open-weight strategy has been an important part of the China large language model (LLM) market, and we believe the debate for investors is more focused on how these LLM providers use open distribution to build monetizable channels. We believe the commercial logic is to use open access to reach developers, cloud platforms, application programming interface (API) aggregators and overseas users faster, then monetize through official API, cloud partnerships, enterprise deployments and workflow products.

The cost of such open-source strategy is also visible: once a model is open-weight, cloud service providers (CSP), API aggregators and technically capable enterprises can deploy it on their own infrastructure. This creates competition for part of the traffic that may otherwise have gone to the LLM provider's first-party API. The pressure is most visible in simple workloads such as chat, summarization, translation, basic Q&A and other low service-level agreement (SLA) text tasks, where users mainly compare price and basic functionality. These tasks have lower requirements on model freshness, long-context stability, tool calling, prompt caching and enterprise support.

That said, open weights do not make all API endpoints equal. Running an open model is different from operating a high-quality production inference service. Third-party deployments can be cheaper or more widely available, while first-party APIs usually have better access to the latest model version, model-specific serving parameters, prompt caching, long-context execution, rate-limit management, toolchain support and enterprise SLA. For more demanding workloads, proprietary API pricing is tied to stable access to a higher-quality endpoint.

The quality gap can widen after the initial model launch. An open-weight release is usually a released checkpoint, while the official API remains a live product. After launch, the model provider continues to observe real-world traffic and optimize the model, the system layer and the serving stack. These follow-on improvements can include post-training updates, instruction-following refinements, coding and tool-use tuning, refusal-policy changes, routing logic, cache policy, context-window handling, decoding parameters and inference-kernel optimization. Most of these changes are not re-open-sourced in real time. As a result, third-party deployment of the same open-weight model may gradually fall behind the official API in both perceived intelligence and production reliability, even if the original base weights are identical.

This is especially relevant for coding, agentic and long-context workloads. In these tasks, users usually care about task completion, repo-level understanding, debugging quality, multi-step planning, tool reliability and latency consistency. Small differences in the serving stack can create visible differences in user experience. The same model name, therefore, does not guarantee the same product quality.

We therefore think open-weight strategy should be analyzed as a shift in revenue mix. **LLM providers may lose part of the revenue tied to model access alone, while gaining a wider funnel for paid inference, cloud distribution and workflow monetization.** The balance should vary by model quality: stronger models are more likely to attract developers, receive cloud platform priority and convert open-weight adoption into paid usage.

Table 1: Open-weight strategy reshapes the monetization mix

Area	Positive impact	Commercial pressure	What remains monetizable
Model access	Lower barrier for developers, CSPs, aggregators and overseas users	Third parties can deploy and resell the model	Latest versions, premium versions, commercial licenses
API usage	Wider user funnel and more trial demand	Simple workloads face more price comparison from non-proprietary APIs	API quality: latency, throughput, caching, uptime, SLA, model freshness
Cloud distribution	Faster global reach through existing cloud platforms	CSPs may self-deploy and earn 100% of API profit	Revenue share, official endpoints, enterprise co-selling
Workflows	Easier entry into coding, agent and enterprise ecosystems	Generic token pricing remains competitive	Task completion, reliability, integration, support
Enterprise deployment	More awareness and technical validation	Some enterprises can self-host	Private deployment, security, support, customization

Source: J.P. Morgan estimates. (As of Jul 6, 2026)

Open-weight is a dynamic commercial choice for LLM providers

Open-source should be analyzed by model version, license and channel. LLM providers can decide by model generation, capability tier, customer type and commercialization stage which models to open, how permissive the license should be, and which premium versions to keep in paid channels.

The 2025-26 release cycle shows three broad patterns among the major independent China LLM providers: **DeepSeek sits closest to the permissive open-weight end of the spectrum; Zhipu uses permissive open-weight base flagships to drive adoption, while keeping selected Turbo / multimodal Turbo variants in managed API channels; Kimi and MiniMax show how modified or community licenses can preserve different degrees of commercial control.**

For users, the open-weight route creates a fallback option. Once the weights are public, developers and enterprises can self-deploy, use a third-party inference provider or migrate across cloud routes if hosted access changes. This access certainty can support experimentation and ecosystem build-out. It also reduces fear that a model family could disappear from the user's stack after a policy change, access restriction or commercial dispute.

Table 2: China LLM providers are increasingly tiering open weights and monetization control

	License	Conditional / restricted tier	API-led or managed tier
DeepSeek	MIT	Limited commercial restrictions in the MIT release	First-party API with aggressive pricing and cache-based economics
Zhipu / Z.AI	GLM-5 / GLM-5.1 / GLM-5.2 under MIT; GLM-5-Turbo / GLM-5V-Turbo are proprietary	No broad commercial restriction on GLM-5 / GLM-5.1 / GLM-5.2 weights	GLM-5-Turbo, GLM-5V-Turbo and other speed-optimized or multimodal-agent variants are primarily accessed through Z.ai API / managed channels
Moonshot / Kimi	Modified MIT	Modified MIT; attribution required above large-scale thresholds such as 100mn MAU or US\$20mn monthly revenue	Additional restrictions on third-party commercialization
MiniMax	MiniMax community license	M3 uses a community license with commercial-use conditions	Additional restrictions on third-party commercialization

Source: J.P. Morgan estimates. (As of Jul 6, 2026); Note: an open weight Massachusetts Institute of Technology (MIT) license means anyone can download, modify, fine-tune, and commercially deploy the model's neural network weights without paying royalties or facing restrictive terms, provided they include the original copyright notice

API quality: proprietary APIs sell stable access to fresher and higher-quality models

Open-weight release increases model availability, but it does not make API products interchangeable. In production, users are buying a serving route rather than model weights alone. That route includes model freshness, cache policy, throughput, latency, feature support, reliability and provider-level optimization. **We think proprietary APIs remain the most optimal choice to most of users because they offer the most stable route to the newest and highest-quality model variants**, while third-party deployment can widen access with more variable economics and service quality.

Model freshness is increasingly important. An open-weight model should be viewed as a released checkpoint, while the official API remains a continuously optimized product. After launch, the model provider can keep improving instruction following, coding behavior, long-context stability, tool calling, refusal policy, caching strategy and inference efficiency based on live traffic. These optimizations are often absorbed into the official endpoint without a full public re-release of the weights. For users, the official API may behave like a fresher and more capable version of the same model, while third-party endpoints may still run the earlier open-weight snapshot or a less optimized serving stack.

The data points to two forms of first-party API advantage. **DeepSeek V4 Pro is the price-led case:** the official route is roughly **6-12x cheaper** than selected third-party routes in a repeated-context workload, helped by lower list pricing and stronger cache economics. **MiniMax M3 demonstrates the quality-led case:** headline output pricing is broadly similar across providers, while the official route has better cache hit rate, lower realized input cost, higher usage concentration and stronger speed. These cases suggest open-weight distribution can broaden reach without fully commoditizing first-party API revenue.

Case Study #1: DeepSeek V4 Pro – official route wins through list price and cache economics

DeepSeek V4 Pro is the cleanest example of a first-party API using price as a moat. OpenRouter's provider-level data shows the DeepSeek official route has much lower effective input pricing and higher cache hit rate than selected third-party routes, while the model-level price also compares favorably before caching. In a simplified monthly workload with **100M effective input tokens and 20M output tokens**, the official route costs roughly **US\$24-41**, vs **US\$85-196** across selected third-party routes. The gap is especially relevant for coding agents, retrieval-augmented generation (RAG) and long-context automation, where repeated system prompts, repo context, tool schemas and intermediate state make cache efficiency a major driver of realized cost. We believe such effective pricing is more meaningful than list price, as it reflects what customers actually pay after prompt caching, and repeated context can make actual cost meaningfully lower than headline list price.

The same logic applies beyond price. If DeepSeek keeps optimizing the official endpoint after the open-weight release, the official API can retain a quality advantage that is not fully visible from the model name alone. Third-party providers may advertise the same model family, but users may still see differences in answer quality, tool reliability, cache behavior, latency consistency and long-context completion as the official route incorporates more post-launch tuning.

Table 3: Different providers on the same model of DeepSeek V4 Pro offer different API economics

DeepSeek V4 Pro provider	Nominal input / M	Cache-read / M	Nominal output / M	Effective input / M	Cache hit rate	Estimated monthly cost*	Latency	Throughput
DeepSeek official route	US\$0.435/0.870	US\$0.004/0.008	US\$0.870/1.740	US\$0.061	93.5%	US\$23.5-40.9	1.42s	45tps
Bailian Qianfan	US\$0.761	US\$0.063	US\$1.521	US\$0.549	30.2%	US\$85.3	0.72s	24tps
Alibaba Cloud	US\$1.608	US\$0.134	US\$3.216	US\$0.830	53.8%	US\$147.3	1.19s	19tps
GMI Cloud	US\$1.131	US\$0.094	US\$2.262	US\$0.458	64.9%	US\$91.0	1.88s	40tps
SiliconFlow	US\$1.600	US\$0.135	US\$3.135	US\$1.330	18.2%	US\$195.7	2.02s	22tps
NovitaAI	US\$1.600	US\$0.135	US\$3.200	US\$0.277	90.1%	US\$91.7	1.77s	39tps
Together	US\$1.740	US\$0.200	US\$3.480	US\$1.150	38.3%	US\$184.6	0.66s	39tps
Fireworks	US\$1.740	US\$0.145	US\$3.480	US\$0.773	60.6%	US\$146.9	1.59s	32tps


Source: OpenRouter (data as of Jul 6, 2026); * Estimated monthly cost is calculated under the assumption of 100M effective input tokens and 20M output tokens per month

Case Study #2: MiniMax M3 – similar headline price, better official-route quality

MiniMax M3 supports the same thesis from a different angle. The official route does not win mainly by charging a lower output price, as several providers show similar headline output pricing. **The advantage comes from cache efficiency, realized input cost, usage concentration and speed.** OpenRouter shows the MiniMax route with a higher cache hit rate and lower effective input price than selected third-party routes. Artificial Analysis adds the performance angle: MiniMax is among the fastest M3 providers, while lower-cost or quantized third-party routes can be materially slower. For agent and coding workloads, these factors matter because the user pays for task completion, not just raw tokens.

This is also where post-launch optimization matters. If MiniMax continues to tune M3 for coding, long-context and agentic workflows inside its official endpoint, the official route can gradually diverge from third-party deployments even when the underlying open-weight release is the same. The user experience gap may appear in repo retrieval, tool planning, code execution loops, error recovery, latency stability and cache reuse. These differences are difficult to capture through headline token pricing alone.

Figure 1: Different providers on the same model of MiniMax M3 offer different API quality



Source: Artificial Analysis (data as of Jul 6, 2026)

CSP & inference platform distribution: margin retention versus model quality

Open weights give overseas cloud and inference platforms more optionality, but the commercial choice remains the same: they can self-deploy a model and retain more API economics, or work more closely with the model provider and give up part of the economics for better version sync, technical support and endpoint quality. We think this margin-versus-quality trade-off is central to the next stage of LLM distribution, especially as Chinese open-weight models enter global platforms such as AWS, Azure/Foundry, Cloudflare and OpenRouter.

The self-deployment path is economically attractive for CSPs and inference providers. Once weights are open, a platform can deploy the model on its own infrastructure, set its own pricing, route traffic through its own stack and capture more gross margin. This model works best for stable models, generic text tasks and price-sensitive traffic. The cost is that serving quality becomes the platform's responsibility. The platform needs to manage model updates, quantization, batching, context handling, cache behavior, reliability and customer support. A weak deployment can make the same model feel slower, less stable or less capable than the official endpoint.

Self-deployment also carries model-maintenance risk. Once a model is open-weight, the platform can deploy it on its own infrastructure, but it also becomes responsible for quantization, batching, cache policy, kernel optimization, model updates, safety layers and customer-facing reliability. If the model provider keeps improving the official API without releasing every follow-on change to the public weights, the self-deployed endpoint may become stale. This is why the margin-retention decision is not only about graphics processing unit (GPU) cost. The platform is also choosing how much model freshness and quality-control burden it wants to take on.

The partnership path is more attractive for high-value workloads. A platform can work with the model provider through official provider status, licensing, revenue share, technical support, marketplace software as a service (SaaS) distribution or enterprise co-selling. This reduces the platform's take rate, while improving model freshness, serving guidance, enterprise trust and support. We expect this route to be more common for frontier or near-frontier models, coding and agentic workloads, long-context APIs and enterprise production deployments, where customers care more about task completion and reliability than headline token price.

Recent global distribution examples support this framework. On AWS, GLM 5 is available as a Bedrock managed model, while GLM-5.2 is available through AWS Marketplace as a Z.AI SaaS API product. Bedrock gives AWS more control over the managed-model experience, while Marketplace preserves more of Z.AI's first-party API relationship and version control. On Azure, Microsoft Foundry distributes selected Chinese open models (including GLM5.2) through Fireworks AI; the enterprise customer enters through Azure, but Fireworks acts as the inference provider. Fireworks still faces the same trade-off: it can retain more economics by optimizing and serving open models itself, but endpoint quality depends on how well it keeps up with model versions, serving recipes and workload-specific optimization.

The choice between these models should remain dynamic. A model with fast iteration, clear benchmark leadership and strong customer pull is more likely to retain first-party or official-backed monetization, as platforms need fresher versions, official support and better serving recipes to meet user expectations. A model that is older, more stable or less differentiated is more likely to be self-deployed by CSPs and inference providers, as customers may prioritize price and availability over marginal quality differences. Channel strategy can also change as customer feedback evolves. If users report quality gaps on third-party endpoints, traffic may move back towards official APIs or official-backed deployments; if third-party providers deliver comparable quality at lower cost, more volume may shift into CSP, inference platforms or aggregator channels.

Table 4: CSP / inference-platform margin-versus-quality choice

Platform choice	Economics	Advantages	Cost / risk	Best fit
Self-deploy open weights	Platform retains more API revenue	Higher take rate, pricing control, infra utilization	Quality risk, version-sync burden, support responsibility	Generic text tasks, stable older models, price-sensitive traffic
Official-backed deployment	Shared economics, license fee or technical support arrangement	Better version sync, serving guidance, enterprise trust	Lower platform take rate	SOTA or near-SOTA models, coding / agentic tasks, long-context workloads
Marketplace SaaS API	Model provider retains more first-party economics; CSP monetizes procurement channel	Lower purchasing friction, stronger provider control over model versions and pricing	Less CSP control over serving and API experience	Latest models, enterprise buyers needing cloud procurement
Partner-powered inference	Economics shared across CSP, inference provider and model owner	Faster model onboarding, optimized serving stack, enterprise platform access	More parties in the value chain; quality still depends on serving execution	Open models requiring low-latency, high-throughput inference

Source: J.P. Morgan estimates (data as of Jul 6, 2026)

This framework leaves room for LLM-provider revenue even after weights are open. CSPs and inference providers can bypass the model company for some workloads, but quality and support responsibility rise when they do. In workloads where adoption depends on reliability, model freshness, cache economics and task completion, official APIs and official-backed endpoints should remain commercially relevant. Open weights expand distribution, while the best monetization path depends on model quality, iteration speed and customer usage patterns.

For Chinese LLM providers, overseas cloud distribution should still be positive. It reduces customer acquisition cost, expands global reach and allows usage to scale across third-party infrastructure. The trade-off is lower direct control and shared economics in some channels. For strong models, we see this as an acceptable funnel: users may start through a CSP catalogue, marketplace, Fireworks-style inference layer or aggregator, then migrate to first-party APIs, premium endpoints or workflow products when reliability and support become more important.

Workflows: proprietary harnesses can turn open-weight adoption into stickier API usage

API access remains the main commercialization route for LLM providers today.

It is the cleanest way to convert model capability into revenue, especially for developer workloads, enterprise integrations and third-party applications that do not require a full product layer from the model company. We expect API revenue to remain the core line for most LLM providers in the near term, but looking ahead, the more important question is where pricing power can be sustained as model capability diffuses and open-weight alternatives broaden.

We expect API pricing to become increasingly bifurcated. Mature intelligence should deflate over time, simple chat, summarization, translation and basic Q&A are easier to reproduce, easier to evaluate and easier for users to switch across providers – in these workloads, customers compare effective cost, latency and stability, and lower-cost providers can gain share quickly once output quality is good enough. Frontier intelligence follows a different pricing logic: when a model creates a visible gap in coding, agentic execution, long-context reasoning or high-reliability enterprise workflows, users can accept higher API prices because the model changes task completion, not just token cost. This is why frontier improvement can still support pricing power, even as mature capabilities become cheaper (see our previous note [Zhipu AI: Mature intelligence deflates pricing, but GLM-5.2 shows frontier upgrades can do the opposite](#) for a detailed discussion).

This bifurcation also explains why LLM providers are trying to move beyond raw API sales. A company that only sells tokens through third-party channels remains exposed to rerouting. A cloud endpoint, artificial intelligence integrated development environment (AI IDE) or agent platform can shift traffic quickly when another model becomes cheaper or reaches similar capability. LLM providers therefore have a strong incentive to build proprietary harnesses around their own models. These products turn the model from a replaceable backend into part of the user's daily workflow. Zhipu's Z Code, MiniMax's MiniMax Code, Kimi's Kimi Code experiments and DeepSeek's reported hiring around agent harness roles all point in the same direction: **LLM providers are trying to own more of the usage layer, instead of leaving the customer relationship entirely to clouds, IDEs and application platforms.**

This is particularly relevant for open-weight monetization. If users only access an open model through third-party clouds, API routers or external IDEs, the LLM provider may gain awareness without owning the customer relationship. Traffic can be redirected quickly when a peer model becomes cheaper or good enough. A successful first-party harness changes that equation: open weights can broaden adoption at the top of the funnel, while the provider's own coding agent, workflow product or enterprise tool converts part of that adoption into sticky first-party usage. In that scenario, open-source has less negative impact on customer stickiness and proprietary API sales because the user is no longer choosing among comparable endpoints. The user is working inside a product layer designed around the provider's own model.

The strategic value of a proprietary harness is larger than incremental subscription revenue. If users consume the model through the provider's own coding agent or workflow product, the model company gains a direct user relationship, faster feedback loops, more proprietary usage data and stronger retention. It can also make API usage more captive. Instead of competing on generic token pricing, the provider can package the model with repo indexing, context management, tool calling, memory, permissions, debugging loops, enterprise controls and workflow-specific optimization. The customer is then paying for task completion and reliability, with token consumption embedded inside a broader workflow product.

The US coding-agent market offers a useful reference. Cursor became a preferred workflow layer for many developers because it combined a strong coding interface with access to leading models. **The emergence of Claude Code shows that an LLM provider can pull usage back into its own environment when the underlying model is sufficiently differentiated.** For many developers, the original reason to use a third-party coding tool was access to the best coding model available at the time. Once Anthropic paired strong Claude models with a first-party coding harness optimized around those models, some advanced users had a reason to shift heavier agentic coding work into Claude Code. **The broader lesson is that workflow ownership is contestable when the LLM provider can deliver both superior model capability and a better model-specific harness.**

We think two conditions are required for this migration to happen. The first is sustained model leadership. Users need a clear reason to change their existing workflow, and that reason usually starts with model quality. In coding, the gap needs to show up in repo-level understanding, multi-step debugging, long-horizon planning, tool use and completion rate. A model that is only marginally better may win benchmark attention, although it may not change daily developer behavior. A model that is clearly better can become the reason users open a product in the first place. If developers use a third-party IDE mainly to access one model family, the LLM provider has a credible path to pull part of that usage into its own harness.

The second condition is harness-specific optimization. A first-party harness needs to make the provider's own model work better than the same model inside a third-party wrapper. This requires high-frequency iteration across context engineering, repo indexing, memory, tool execution, permissioning, test loops, error recovery and product user experience (UX). The advantage compounds when the model team and product team optimize together. The model can be tuned for the harness, and the harness can be redesigned around the model's strengths and failure modes. Over time, the first-party product can become more than a distribution channel. It becomes part of the model's effective capability.

China's migration path is likely to be more difficult than the US reference case. Tencent WorkBuddy, ByteDance Trae and Alibaba Qoder already care about owning the workflow entrance and have accumulated early users through existing traffic, ecosystem relationships and product distribution. Chinese LLM providers also face a more crowded model landscape, where the capability gap among leading players is often narrower than the gap between the strongest US frontier model and the rest of the market. This makes it harder for any single model company to persuade users to leave an existing workflow product on the basis of modest model improvement.

First-party harness adoption is still possible in China. The key variable is whether an LLM provider can create an absolute capability step-change. If a Chinese model becomes clearly stronger in coding or agentic workflows, and if the provider's own harness consistently converts that model advantage into better task completion than third-party platforms using the same API, user behavior can change. In that scenario, the LLM provider would have a path beyond generic API access. It could own a higher-value workflow layer, improve customer retention, make usage more proprietary, and reduce dependence on price competition in mature token markets.

Company view: SOTA models should capture most of the open-weight monetization upside

Revenue opportunities from open-weight models should remain concentrated. Open weights expand distribution, but monetization still depends on model quality, iteration speed, endpoint reliability, workflow attachment and enterprise trust. **Stronger state of the art (SOTA) or near-SOTA models are more likely to gain developer adoption, receive CSP and third-party inference-platform priority, retain official-backed API economics and convert usage into workflow products.** Weaker models can still generate downloads and experimentation, while paid conversion is harder when the model does not create a visible difference in task completion.

Open weights also reduce user concern around access continuity. A public-weight model gives developers and enterprises a fallback route if a hosted endpoint changes policy, raises price or becomes unavailable. That access certainty can support adoption, experimentation and ecosystem build-out. However, the monetization upside still depends on whether the provider can make its official API, cloud partnership and workflow products meaningfully better than the fallback open-weight route.

This concentration matters more as the market becomes more open. Once weights are available, users and platforms can compare models faster. A model that is merely available or cheap may see usage in generic workloads, but it is more exposed to rerouting when another model becomes cheaper or good enough. A model that is clearly better in coding, agentic execution, long-context reasoning or enterprise reliability has more ways to monetize: first-party API, official-backed deployment, cloud marketplace, enterprise deployment and proprietary harness. This is why we expect the open-weight playbook to reward capability leadership more than participation.

While the Chinese model companies become incrementally competitive on a global basis, we highlight the risks of potential capital-raising and funding requirements. Despite near-term revenue/ARR visibility, both companies remain in a capital-intensive phase, with sustained investment required for model training, inference infrastructure and global go-to-market. We now expect Zhipu to spend Rmb5.7/9.6/14.5bn on R&D (mostly training expense) in 2026-28E, with MiniMax at US\$809mn/1.2bn/1.5bn. Based on disclosed financials, public news and our estimates ([Bloomberg: Zhipu is considering a share sale to raise several billion US dollars in Hong Kong](#); [Zhipu announcement: proposed issue of a shares and listing on the sci-tech board with Rmb15bn targeted proceeds](#); [MiniMax announcement: preliminary proposal for the proposed issue of RMB shares](#)), annual cash burn prior to breakeven remains material, with operating cash outflows likely to persist through 2027/28E for Zhipu/MiniMax – we model both companies to do two more rounds of capital-raising in 2026 and 2027. While current cash balances provide runway under base-case assumptions, accelerated model iteration, larger-scale overseas deployment or higher-than-expected inference costs could necessitate incremental external funding.

Zhipu AI — Maintain OW, raise PT to HK\$2,000

We reiterate OW on Zhipu AI and raise our Dec-26 PT to HK\$2,000, as GLM-5.2 strengthens the case that open-weight commercialization can create sizeable optionality for leading model providers. Based on current valuation, we believe the market has largely moved beyond Zhipu's US\$1bn year-end ARR guidance; we think the remaining upside lies in whether a strong open-weight model can scale through external infrastructure and distribution, rather than only Zhipu's own GPU stack. Public weights allow General Language Models (GLM) to run across clouds, inference providers and enterprise infrastructure, while Zhipu can still participate through official APIs, official-backed endpoints, licensing, deployment support and Z Code-style workflow products. This is optionality rather than confirmed near-term revenue, and it can only scale if Zhipu sustains model leadership. The key watch points are GLM-5.5/6, Kimi K3 and DeepSeek V4.1

Estimate changes

We raise Zhipu's 2026-30E revenue by 3-9% to reflect higher visibility on its ability to scale up globally leveraging its open-sourced strategy. Accordingly, we lower our adjusted net loss forecasts from a loss of Rmb3,750mn, loss of Rmb3,492mn and profit of Rmb1,287mn in 2026, 2027 and 2028 to a loss of Rmb3,711mn, loss of Rmb3,141mn and profit of Rmb2,367mn. We consequently raise our Dec-26 price target to HK\$2,000 from HK\$1,800, still based on 30x 2030E P/E discounted back using a 15% WACC.

Table 5: Estimate changes summary for Zhipu

Rmb mn	2026E	2027E	2028E	2029E	2030E
Revenue					
New	4,715	12,324	33,485	89,984	155,497
Old	4,597	11,430	30,862	82,638	143,700
% change	3%	8%	9%	9%	8%
Operating profit					
New	(5,280)	(6,108)	(2,322)	16,932	42,847
Old	(5,318)	(6,459)	(3,402)	13,757	37,495
% change	1%	5%	32%	23%	14%
Adjusted net profit					
New	(3,711)	(3,141)	2,367	20,141	42,858
Old	(3,750)	(3,492)	1,287	17,601	38,576
% change	1%	10%	84%	14%	11%

Source: J.P. Morgan estimates. (data as of Jul 6, 2026)

MiniMax – Maintain Neutral, lower PT to HK\$300

We maintain our Neutral rating on MiniMax and lower our Dec-26 PT to HK\$300, as open-weight commercialization is becoming a ‘winner-takes-more’ framework and M3 has not yet shown enough evidence of model-led pricing power, in our view. MiniMax remains relevant across multimodal AI, overseas usage and agentic workflows, and M3 improved the product narrative through 1M context, native multimodality, API access and MiniMax Code. However, broader access also increases price comparison and traffic rerouting risk when the model is not clearly differentiated. M3’s permanent 50% discount is an important signal that the model has yet to command a capability premium versus leading domestic peers. We would become more constructive if MiniMax narrows the capability gap, sustains API usage after discount normalization and proves stronger workflow attachment through MiniMax Code.

Estimate changes

We lower MiniMax’s 2027-30E revenue by 2-8% as we believe the open-source license restrictions cap third party reach while the official route lacks demonstrated pricing power. Accordingly, we lower our adjusted net loss forecasts from a loss of US\$940mn/US\$972mn in 2027/2028 to a loss of US\$948mn/US\$1,003mn. We consequently lower our Dec-26 price target to HK\$300 from HK\$400, still based on 30x 2030E P/E discounted back using a 15% WACC.

Table 6: Estimate changes summary for MiniMax

US\$mn	2026E	2027E	2028E	2029E	2030E
Revenue					
New	383	830	1,652	3,194	6,633
Old	383	850	1,734	3,447	7,226
% change	0%	-2%	-5%	-7%	-8%
Operating profit					
New	(932)	(1,217)	(1,353)	(1,148)	(41)
Old	(932)	(1,209)	(1,323)	(1,047)	214
% change	0%	-1%	-2%	-10%	-119%
Adjusted net profit					
New	(432)	(948)	(1,003)	(647)	633
Old	(432)	(940)	(972)	(546)	842
% change	0%	-1%	-3%	18%	-25%

Source: J.P. Morgan estimates. (data as of Jul 6, 2026)

Zhipu AI

Overweight

2513.HK, 2513 HK

Price (07 Jul 26): HK\$1,610.00

▲ **Price Target (Dec-26): HK\$2,000.00**
Prior (Dec-26): HK\$1,800.00

Hong Kong

Internet

Olivia Xu ^{AC}

(86-21) 6106 6138

olivia.w.xu@jpmorgan.com

SAC Registration Number: S1730525060001

J.P. Morgan Securities (China) Company Limited

Key Changes (FYE Dec)

	Prev	Cur	Δ
Adj. EPS - 26E (Rmb)	(8.29)	(8.20)	1.0%
Adj. EPS - 27E (Rmb)	(7.45)	(6.71)	10.0%

Half Yearly Forecasts (FYE Dec)

Adj. EPS (Rmb)

	2025A	2026E	2027E
H1	(11.09)	(5.07)	(4.78)
H2	(9.11)	(3.16)	(1.92)
FY	(19.95)	(8.20)	(6.71)

Style Exposure

Quant Factors	Current	Hist %Rank (1=Top)			
	%Rank	6M	1Y	3Y	5Y
Value	98	94	94	86	
Growth	31	100	100		
Momentum	100				
Quality	86				
Low Vol	100				

Sources for: Style Exposure – J.P. Morgan Global Markets Strategy; all other tables are company data and J.P. Morgan estimates.

Raise PT to HK\$2,000: GLM-5.2 increases open-weight monetization optionality

We reiterate our OW rating on Zhipu AI and raise our Dec-26 PT to HK\$2,000, as GLM-5.2 strengthens the case that open-weight commercialization can create sizeable upside for leading model providers. Based on current valuation, we believe the market has largely moved beyond Zhipu's US\$1bn year-end ARR guidance. We think the remaining upside is less about ARR supported by Zhipu's own GPU stack, and more about whether a strong open-weight model can scale through external infrastructure and distribution. GLM-5.2 improves Zhipu's domestic frontier positioning and raises the probability that open distribution converts into paid usage across official APIs, partner-backed deployment, enterprise support and Z Code-style workflow products.

For near-SOTA models, open weights can broaden adoption while preserving quality-based monetization. Public weights allow model to run on existing GPU capacity across clouds, inference providers and enterprise infrastructure, giving Zhipu a potentially more capital-efficient growth path than a model where all usage must be served by its own compute. Customers may start from first-party APIs, CSP catalogs, API aggregators or private deployment; Zhipu can still monetize through licensing, official-backed endpoints, marketplace SaaS, revenue share, deployment support and enterprise co-selling. We treat this as optionality rather than confirmed near-term revenue, as partnership structures and overseas economics still need to be proven.

In our view, Zhipu is better positioned than peers because open-weight payoff is becoming more concentrated in stronger models. Third-party CSPs and inference platforms still have reasons to work with the model provider, as commercial customers are buying a production-grade model service and update cadence, rather than access to weights alone. A public release is closer to a checkpoint, while the official route can continue to improve after launch based on live traffic and customer feedback. The stronger the model, the more customers care about freshness, reliability and support, and the more likely partners are to support official-backed routes.

Open weights also reduce adoption risk. Once weights are public, developers and enterprises have a fallback route if hosted API pricing, access policy or availability changes, making it easier to test and standardize around GLM. As the model keeps improving, Zhipu has more opportunities to deepen workflow integration through official APIs, managed deployment, support services or Z Code. The key risk is durability of model leadership. GLM-5.2 creates the optionality, while GLM-5.5/6, Kimi K3 and DeepSeek V4.1 will determine whether it can scale into a more visible revenue path.

Price Performance



— 2513.HK Price (HK\$)– MSCI-Cnx (rebased)

	YTD	1m	3m	12m
Abs	-	24.1%	106.7%	-
Rel	-	29.5%	113.0%	-

Company Data

Shares O/S (mn)	221
52-week range (HK\$)	2,980.00-116.10
Market cap (\$ mn)	45,431
Exchange rate	7.84
Free float (%)	56.2%
3M ADV (mn)	2.00
3M ADV (\$ mn)	355.0
Volatility (90 Day)	163
Index	MXCNX
BBG ANR (Buy Hold Sell)	19 4 1

Key Metrics (FYE Dec)

Rmb in millions	FY25A	FY26E	FY27E	FY28E
Financial Estimates				
Revenue	724	4,715	12,324	33,485
Adj. EBIT	(7,015)	(8,973)	(9,285)	(85)
Adj. EBITDA	(3,517)	(4,808)	(5,492)	(1,652)
Adj. net income	(3,182)	(3,711)	(3,141)	2,367
Adj. EPS	(19.95)	(8.20)	(6.71)	5.05
BBG EPS	(8.58)	(10.82)	(10.16)	(5.08)
Cashflow from operations	(2,246)	(1,828)	1,309	8,754
FCFF	(2,195)	(2,225)	768	7,130
Margins and Growth				
Revenue Growth Y/Y (%)	131.9%	550.9%	161.4%	171.7%
EBIT margin	(968.5%)	(190.3%)	(75.3%)	(0.3%)
EBIT Growth Y/Y (%)	38.8%	27.9%	3.5%	(99.1%)
EBITDA margin	(485.5%)	(102.0%)	(44.6%)	(4.9%)
EBITDA Growth Y/Y (%)	55.0%	36.7%	14.2%	(69.9%)
Net margin	(439.3%)	(78.7%)	(25.5%)	7.1%
Adj. EPS growth	29.3%	(58.9%)	(18.2%)	(175.3%)
Ratios				
Adj. tax rate	0.0%	0.0%	0.0%	0.0%
Interest cover	NM	NM	NM	NM
Net debt/Equity	0.1	0.8	NM	NM
Net debt/EBITDA	0.2	0.9	3.7	16.5
ROE	52.8%	53.6%	(2035.9%)	47.5%
Valuation				
FCFF yield	(1.0%)	(0.4%)	0.1%	1.1%
Dividend yield	-	-	-	-
EV/Revenue	471.4	71.6	26.1	9.4
EV/EBITDA	NM	NM	NM	NM
Adj. P/E	NM	NM	NM	276.0

Summary Investment Thesis and Valuation

We rate Z.AI (Zhipu AI) Overweight, as the company better fits our framework of repeated SOTA visibility, high-value workflow exposure and pricing power. Year-to-date, Zhipu has delivered multiple domestic SOTA releases across the GLM series, and GLM-5.1 still ranks as one of the best Chinese models on WebDev Arena even after the later releases of Kimi K2.6 and DeepSeek V4, and the release of GLM-5.2 further enhances its leadership position. In a market where model cycles are compressing, we believe such a track record matters: investors are likely to place more weight on repeated frontier delivery than on one-off usage spikes or single-model success.

Our Dec-26 PT of HK\$2,000 is derived from 30x 2030E P/E, when we expect Zhipu AI to generate normalized earnings, discounted back using a 15% WACC. The 30x multiple represents a valuation premium over tier-1 Chinese internet players, mostly to reflect the company's 100+% revenue CAGR in 2026-30E.

Performance Drivers

Market	4%
Region	4%
Macro	7%
Style	3%
Idiosyn.	81%

Factors	6M Corr	1Y Corr
Market: MSCI Asia Pac ex JP	0.29	0.21
Region: Hong Kong	-0.21	-0.21
Macro:		
JPM China A-shares Sentiment	-0.26	-0.24
JPM EM Currency(EMCI) Fixing	-0.18	-0.17
JPM Global Equity Sentiment	-0.22	-0.14
Quant Styles:		
Value	-0.13	-0.14
Quality	-0.18	-0.13
Momentum	0.11	0.08

Investment Thesis, Valuation and Risks

Knowledge Atlas Technology Joint Stock Co. Ltd.

(Overweight; Price Target: HK\$2,000.00)

Investment Thesis

We rate Z.AI (Zhipu AI) Overweight, as the company better fits our framework of repeated SOTA visibility, high-value workflow exposure and pricing power. Year-to-date, Zhipu has delivered multiple domestic SOTA releases across the GLM series, and GLM-5.1 still ranks as one of the best Chinese models on WebDev Arena even after the later releases of Kimi K2.6 and DeepSeek V4. The release of GLM5.2 further enhances its leadership position. In a market where model cycles are compressing, we believe such a track record matters: investors are likely to place more weight on repeated frontier delivery than on one-off usage spikes or single-model success.

Valuation

Our Dec-26 PT of HK\$2,000 is based on a 30x 2030E P/E, when we expect Zhipu AI to generate normalized earnings, discounted back using a 15% WACC. The 30x multiple represents a valuation premium over tier-1 Chinese internet players, mostly to reflect the company's 100+% revenue CAGR in 2026-30E.

Price target calculation

2030E Revenue (Rmb mn)	155,497
2030E Adjusted net profit (Rmb mn)	42,858
2030E Adjusted EPS (Rmb/share)	91
Target multiple	30
WACC	15%
Dec-26 PT (HK\$)	2,000

Source: J.P. Morgan estimates.

Risks to Rating and Price Target

Risks to our rating and price target include export controls, geopolitical risk, and entity list designation; intensified competition; high and ongoing R&D investment creates execution risk and pressure on profitability; commercialization and customer adoption remain subject to uncertainty; dependence on computing infrastructure and external suppliers exposes Zhipu to cost and availability risks.

Zhipu AI: Summary of Financials

Income Statement						Cash Flow Statement					
	FY24A	FY25A	FY26E	FY27E	FY28E		FY24A	FY25A	FY26E	FY27E	FY28E
Revenue	312	724	4,715	12,324	33,485	Cash flow from operating activities	(2,245)	(2,246)	(1,828)	1,309	8,754
COGS	(137)	(428)	(3,067)	(7,327)	(19,460)	o/w Depreciation & amortization	270	270	472	616	670
Gross profit	176	297	1,648	4,997	14,025	o/w Changes in working capital	72	822	1,325	3,627	9,945
SG&A	(521)	(896)	(1,203)	(1,488)	(1,825)	Cash flow from investing activities	(49)	(376)	(472)	(616)	(1,674)
Adj. EBITDA	(2,268)	(3,517)	(4,808)	(5,492)	(1,652)	o/w Capital expenditure	(133)	(23)	(472)	(616)	(1,674)
D&A	(270)	(270)	(472)	(616)	(670)	as % of sales	42.5%	3.2%	10.0%	5.0%	5.0%
Adj. EBIT	(5,053)	(7,015)	(8,973)	(9,285)	(85)	Cash flow from financing activities	3,312	2,614	6,000	15,000	0
Net Interest	(38)	(74)	(74)	(74)	(50)	o/w Dividends paid	0	0	0	0	0
Adj. PBT	(2,958)	(4,718)	(5,297)	(6,072)	(2,192)	o/w Shares issued/(repurchased)	0	0	6,000	15,000	0
Tax	0	0	0	0	0	o/w Net debt issued/(repaid)	330	1,211	0	0	0
Minority Interest	(2)	(20)	(21)	(24)	(6)	Net change in cash	1,020	(10)	3,700	15,693	7,080
Adj. Net Income	(2,466)	(3,182)	(3,711)	(3,141)	2,367	Adj. Free cash flow to firm	(2,225)	(2,195)	(2,225)	768	7,130
Reported EPS	(8.72)	(12.03)	(11.66)	(12.91)	(4.67)	y/y Growth	96.8%	(1.3%)	1.4%	(134.5%)	828.9%
Adj. EPS	(15.42)	(19.95)	(8.20)	(6.71)	5.05						
DPS	-	-	-	-	-						
Payout ratio	-	-	-	-	-						
Shares outstanding	160	160	453	468	468						
Balance Sheet						Ratio Analysis					
	FY24A	FY25A	FY26E	FY27E	FY28E		FY24A	FY25A	FY26E	FY27E	FY28E
Cash and cash equivalents	2,269	2,259	5,959	21,653	28,733	Gross margin	56.3%	41.0%	34.9%	40.5%	41.9%
Accounts receivable	667	699	5,846	12,223	26,569	EBITDA margin	(726.0%)	(485.5%)	(102.0%)	(44.6%)	(4.9%)
Inventories	32	126	734	1,579	3,773	EBIT margin	(1617.4%)	(968.5%)	(190.3%)	(75.3%)	(0.3%)
Other current assets	704	826	6,621	13,909	30,632	Net profit margin	(789.2%)	(439.3%)	(78.7%)	(25.5%)	7.1%
Current assets	3,016	3,571	13,066	36,047	59,851	ROE	99.8%	52.8%	53.6%	(2035.9%)	47.5%
PP&E	866	656	584	338	1,007	ROA	(68.1%)	(69.0%)	(38.8%)	(12.2%)	4.8%
LT investments	-	-	-	-	-	ROCE	329.7%	152.4%	163.6%	(581.0%)	(1.3%)
Other non current assets	404	536	536	536	536	SG&A/Sales	166.8%	123.7%	25.5%	12.1%	5.4%
Total assets	4,376	4,854	14,277	37,012	61,485	Net debt/Equity	0.2	0.1	0.8	NM	NM
Short term borrowings	269	605	605	605	605	Net debt/EBITDA	0.4	0.2	0.9	3.7	16.5
Payables	603	1,329	7,547	16,767	38,722	Sales/Assets (x)	0.1	0.2	0.5	0.5	0.7
Other short term liabilities	6,965	10,473	11,374	13,069	17,782	Assets/Equity (x)	NM	NM	NM	166.2	9.9
Current liabilities	7,838	12,406	19,526	30,441	57,109	Interest cover (x)	NM	NM	NM	NM	NM
Long-term debt	458	294	294	294	294	Operating leverage	202.6%	29.4%	5.1%	2.2%	(57.7%)
Other long term liabilities	493	559	559	559	559	Tax rate	0.0%	0.0%	0.0%	0.0%	0.0%
Total liabilities	8,331	12,965	20,084	30,999	57,667	Revenue y/y Growth	150.9%	131.9%	550.9%	161.4%	171.7%
Shareholders' equity	(3,957)	(8,093)	(5,768)	6,076	3,881	EBITDA y/y Growth	303.6%	55.0%	36.7%	14.2%	(69.9%)
Minority interests	1	(18)	(40)	(64)	(64)	EPS y/y Growth	275.0%	29.3%	(58.9%)	(18.2%)	(175.3%)
Total liabilities & equity	4,376	4,854	14,277	37,012	61,485						
BVPS	(24.75)	(50.73)	(12.75)	12.97	8.29						
y/y Growth	280.3%	105.0%	(74.9%)	(201.8%)	(36.1%)						
Net debt/(cash)	(870)	(815)	(4,516)	(20,209)	(27,289)						
						Valuation					
	FY24A	FY25A	FY26E	FY27E	FY28E		FY24A	FY25A	FY26E	FY27E	FY28E
P/E (x)	NM	NM	NM	NM	276.0						
P/BV (x)	NM	NM	NM	107.5	168.3						
EV/EBITDA (x)	NM	NM	NM	NM	NM						
Dividend Yield	-	-	-	-	-						

Source: Company reports and J.P. Morgan estimates.

Note: Rmb in millions (except per-share data). Fiscal year ends Dec. o/w - out of which

Neutral

0100.HK, 100 HK

Price (07 Jul 26): HK\$323.80

▼ **Price Target (Dec-26): HK\$300.00**
Prior (Dec-26): HK\$400.00

**Hong Kong
Internet**

Olivia Xu ^{AC}

(86-21) 6106 6138

olivia.w.xu@jpmorgan.com

SAC Registration Number: S1730525060001

J.P. Morgan Securities (China) Company Limited

Half Yearly Forecasts (FYE Dec)

Adj. EPS (\$)

	2025A	2026E	2027E
H1	(1.75)	(0.99)	(1.81)
H2	(0.22)	(0.38)	(1.12)
FY	(0.80)	(1.36)	(2.92)

Style Exposure

Quant Factors	Current	Hist %Rank (1=Top)			
	%Rank	6M	1Y	3Y	5Y
Value	97	90	89	85	
Growth	16	99	99		
Momentum					
Quality	89				
Low Vol	99				

Sources for: Style Exposure – J.P. Morgan Global Markets Strategy; all other tables are company data and J.P. Morgan estimates.

MiniMax Group Inc - H

Lower PT to HK\$300: open-weight pressure rises without clear model-led pricing power

We maintain Neutral on MiniMax and lower our Dec-26 PT to HK\$300. MiniMax remains relevant across multimodal AI, overseas usage and agentic workflows, and M3 improved the product narrative through 1M context, native multimodality, API access and MiniMax Code. However, open-weight commercialization is becoming a ‘winner-takes-more’ framework, and M3 has not yet shown enough evidence that broader access can convert into durable, high-quality paid usage.

The key issue is pricing power. Strong models can use public weights to broaden adoption, then monetize through official routes, partner-backed endpoints, enterprise deployment and workflow products. Less differentiated models face faster price comparison and easier traffic rerouting across first-party APIs, aggregators, CSP routes and private deployments. M3’s permanent 50% discount is therefore an important signal that the model has not yet commanded a capability premium versus leading domestic peers. Discounting may support near-term usage, while lowering confidence in model-led monetization.

MiniMax’s open-weight strategy still creates distribution value, especially given its multimodal portfolio and overseas developer exposure. MiniMax Code is directionally aligned with the move towards workflow ownership, where a first-party product can build a direct customer relationship and make API usage stickier. The challenge is that workflow monetization requires stronger model pull. A coding or agentic product needs to improve task completion enough to change user habits, rather than simply offer another access route to a broadly available model.

The next model release is the key catalyst. We would become more constructive if MiniMax narrows the capability gap versus domestic leaders, sustains API usage after discount normalization, improves workflow adoption and shows clearer evidence of model-led pricing. Until then, we believe the same open-weight forces that create upside optionality for Zhipu create more risk for MiniMax: broader distribution, faster price comparison and weaker ability to retain premium economics

Price Performance



— 0100.HK Price (HK\$) - MSCI-Cnx (rebased)

	YTD	1m	3m	12m
Abs	-	-41.4%	-65.9%	-
Rel	-	-36.1%	-59.6%	-

Company Data

Shares O/S (mn)	228
52-week range (HK\$)	1,330.00-220.00
Market cap (\$ mn)	9,419
Exchange rate	7.84
Free float (%)	51.7%
3M ADV (mn)	2.00
3M ADV (\$ mn)	171.5
Volatility (90 Day)	140
Index	MXCNX
BBG ANR (Buy Hold Sell)	25 3 1

Key Metrics (FYE Dec)

\$ in millions	FY25A	FY26E	FY27E	FY28E
Financial Estimates				
Revenue	79	383	830	1,652
Adj. EBIT	(257)	(431)	(947)	(1,002)
Adj. EBITDA	(280)	(926)	(1,204)	(1,329)
Adj. net income	(251)	(432)	(948)	(1,003)
Adj. EPS	(0.80)	(1.36)	(2.92)	(3.09)
BBG EPS	(1.12)	(1.90)	(1.97)	(1.61)
Cashflow from operations	(280)	(302)	(823)	(666)
FCFF	(280)	(321)	(864)	(748)
Margins and Growth				
Revenue Growth Y/Y (%)	158.9%	384.9%	116.5%	99.1%
EBIT margin	(325.2%)	(112.4%)	(114.1%)	(60.7%)
EBIT Growth Y/Y (%)	5.5%	67.7%	119.7%	5.9%
EBITDA margin	(354.6%)	(241.6%)	(145.1%)	(80.4%)
EBITDA Growth Y/Y (%)	12.1%	230.4%	30.1%	10.3%
Net margin	(317.4%)	(112.6%)	(114.2%)	(60.7%)
Adj. EPS growth	(64.4%)	70.4%	114.6%	5.8%
Ratios				
Adj. tax rate	0.0%	0.0%	0.0%	0.0%
Interest cover	NM	NM	NM	NM
Net debt/Equity	0.2	0.5	0.5	0.1
Net debt/EBITDA	1.7	1.1	0.9	0.3
ROE	14.6%	17.8%	42.7%	36.5%
Valuation				
FCFF yield	(2.2%)	(2.5%)	(6.5%)	(5.6%)
Dividend yield	-	-	-	-
EV/Revenue	122.0	23.6	10.9	5.9
EV/EBITDA	NM	NM	NM	NM
Adj. P/E	NM	NM	NM	NM

Summary Investment Thesis and Valuation

We are Neutral on MiniMax. MiniMax remains strategically well positioned, with a differentiated multimodal + LLM portfolio, 2C + 2B product coverage, international expansion and strong cost-performance traction. The company has built a credible model platform across text, audio, image, video and productivity use cases, and its OpenRouter usage shows meaningful developer adoption. This combination gives MiniMax long-term option value, especially if AI demand expands from text-heavy workflows into multimodal content, interactive agents and embodied AI. Our Neutral rating is driven by model leadership rather than strategic positioning: from a pure model capability perspective, MiniMax remains in catch-up mode.

Our Dec-26 PT of HK\$300 is derived from 30x 2030E P/E, when we expect MiniMax to generate normalized earnings, discounted back using a 15% WACC. The 30x multiple represents a valuation premium over tier-1 Chinese internet players, mostly to reflect the company's 100+% revenue CAGR in 2026-30E.

Investment Thesis, Valuation and Risks

MiniMax Group Inc - H (Neutral; Price Target: HK\$300.00)

Investment Thesis

We are Neutral on MiniMax. MiniMax remains strategically well positioned, with a differentiated multimodal + LLM portfolio, 2C + 2B product coverage, international expansion and strong cost-performance traction. The company has built a credible model platform across text, audio, image, video and productivity use cases, and its OpenRouter usage shows meaningful developer adoption. This combination gives MiniMax long-term option value, especially if AI demand expands from text-heavy workflows into multimodal content, interactive agents and embodied AI. Our Neutral rating is driven by model leadership rather than strategic positioning: from a pure model capability perspective, MiniMax remains in catch-up mode.

Valuation

Our Dec-26 PT of HK\$300 is derived from 30x 2030E P/E, when we expect MiniMax to generate normalized earnings, discounted back using a 15% WACC. The 30x multiple represents a valuation premium over tier-1 Chinese internet players, mostly to reflect the company's 100+% revenue CAGR in 2026-30E.

Price target calculation

2030E Revenue (US\$mn)	6,633
2030E Adjusted net profit (US\$mn)	633
2030E Adjusted EPS (\$/share)	2.0
Target multiple	30
WACC	15%
Dec-26 PT (HK\$)	300

Source: J.P. Morgan estimates.

Risks to Rating and Price Target

Upside risks to our rating and price target include better-than-expected new model iteration, clearer monetization from MiniMax's multimodal portfolio, and more durable API economics.

Downside risks include MiniMax's litigation proceedings with US studios; intensified competition; high and ongoing R&D investment which creates execution risk and pressure on profitability; commercialization and customer adoption remain subject to uncertainty; and dependence on computing infrastructure and external suppliers exposes MiniMax to cost and availability risks.

MiniMax Group Inc - H: Summary of Financials

Income Statement						Cash Flow Statement					
	FY24A	FY25A	FY26E	FY27E	FY28E		FY24A	FY25A	FY26E	FY27E	FY28E
Revenue	31	79	383	830	1,652	Cash flow from operating activities	(258)	(280)	(302)	(823)	(666)
COGS	(27)	(59)	(307)	(613)	(1,149)	o/w Depreciation & amortization	0	1	6	12	24
Gross profit	4	20	76	217	503	o/w Changes in working capital	20	6	121	96	272
SG&A	(101)	(89)	(199)	(262)	(333)	Cash flow from investing activities	(431)	72	(19)	(41)	(83)
Adj. EBITDA	(250)	(280)	(926)	(1,204)	(1,329)	o/w Capital expenditure	(1)	(1)	(19)	(41)	(83)
D&A	(0)	(1)	(6)	(12)	(24)	as % of sales	2.5%	1.2%	5.0%	5.0%	5.0%
Adj. EBIT	(244)	(257)	(431)	(947)	(1,002)	Cash flow from financing activities	771	423	900	900	0
Net Interest	(1)	(1)	(1)	(1)	(1)	o/w Dividends paid	-	-	-	-	-
Adj. PBT	(465)	(1,872)	(932)	(1,218)	(1,354)	o/w Shares issued/(repurchased)	-	-	-	-	-
Tax	0	0	0	0	0	o/w Net debt issued/(repaid)	33	1	900	900	0
Minority Interest	0	0	0	0	0	Net change in cash	83	219	579	36	(749)
Adj. Net Income	(244)	(251)	(432)	(948)	(1,003)	Adj. Free cash flow to firm	(259)	(280)	(321)	(864)	(748)
Reported EPS	(4.28)	(5.97)	(2.94)	(3.76)	(4.18)	y/y Growth	297.5%	8.2%	14.6%	169.2%	(13.3%)
Adj. EPS	(2.25)	(0.80)	(1.36)	(2.92)	(3.09)						
DPS	-	-	-	-	-						
Payout ratio	-	-	-	-	-						
Shares outstanding	109	314	317	324	324						
Balance Sheet						Ratio Analysis					
	FY24A	FY25A	FY26E	FY27E	FY28E		FY24A	FY25A	FY26E	FY27E	FY28E
Cash and cash equivalents	289	508	1,086	1,122	373	Gross margin	12.2%	25.4%	19.9%	26.2%	30.5%
Accounts receivable	7	11	72	160	329	EBITDA margin	(819.1%)	(354.6%)	(241.6%)	(145.1%)	(80.4%)
Inventories	-	-	-	-	-	EBIT margin	(798.2%)	(325.2%)	(112.4%)	(114.1%)	(60.7%)
Other current assets	517	500	664	872	1,250	Net profit margin	(800.2%)	(317.4%)	(112.6%)	(114.2%)	(60.7%)
Current assets	806	1,007	1,751	1,994	1,623	ROE	42.8%	14.6%	17.8%	42.7%	36.5%
PP&E	1	2	4	8	17	ROA	(39.6%)	(25.1%)	(29.5%)	(48.4%)	(52.8%)
LT investments	95	70	70	70	70	ROCE	43.7%	15.2%	18.1%	43.4%	36.9%
Other non current assets	104	79	79	79	79	SG&A/Sales	332.1%	112.2%	51.9%	31.6%	20.2%
Total assets	911	1,088	1,834	2,081	1,719	Net debt/Equity	0.3	0.2	0.5	0.5	0.1
Short term borrowings	19	35	35	35	35	Net debt/EBITDA	1.1	1.7	1.1	0.9	0.3
Payables	51	58	323	594	1,183	Sales/Assets (x)	0.0	0.1	0.3	0.4	0.9
Other short term liabilities	1,637	3,640	3,661	3,694	3,754	Assets/Equity (x)	NM	NM	NM	NM	NM
Current liabilities	1,708	3,734	4,019	4,323	4,973	Interest cover (x)	NM	NM	NM	NM	NM
Long-term debt	-	-	-	-	-	Operating leverage	22.2%	3.4%	17.6%	102.8%	5.9%
Other long term liabilities	2	3	3	3	3	Tax rate	0.0%	0.0%	0.0%	0.0%	0.0%
Total liabilities	1,710	3,737	4,022	4,326	4,976	Revenue y/y Growth	782.2%	158.9%	384.9%	116.5%	99.1%
Shareholders' equity	(799)	(2,648)	(2,189)	(2,245)	(3,257)	EBITDA y/y Growth	171.2%	12.1%	230.4%	30.1%	10.3%
Minority interests	0	0	0	0	0	EPS y/y Growth	165.8%	(64.4%)	70.4%	114.6%	5.8%
Total liabilities & equity	911	1,088	1,834	2,081	1,719						
BVPS	(7.36)	(8.44)	(6.91)	(6.93)	(10.05)						
y/y Growth	126.6%	14.8%	(18.2%)	0.3%	45.1%						
Net debt/(cash)	(266)	(470)	(1,049)	(1,085)	(336)						
Valuation							FY24A	FY25A	FY26E	FY27E	FY28E
P/E (x)							NM	NM	NM	NM	NM
P/BV (x)							NM	NM	NM	NM	NM
EV/EBITDA (x)							NM	NM	NM	NM	NM
Dividend Yield							-	-	-	-	-

Source: Company reports and J.P. Morgan estimates.

Note: \$ in millions (except per-share data). Fiscal year ends Dec. o/w - out of which

Analyst Certification: The Research Analyst(s) denoted by an “AC” on the cover of this report certifies (or, where multiple Research Analysts are primarily responsible for this report, the Research Analyst denoted by an “AC” on the cover or within the document individually certifies, with respect to each security or issuer that the Research Analyst covers in this research) that: (1) all of the views expressed in this report accurately reflect the Research Analyst’s personal views about any and all of the subject securities or issuers; and (2) no part of any of the Research Analyst’s compensation was, is, or will be directly or indirectly related to the specific recommendations or views expressed by the Research Analyst(s) in this report. For all Korea-based Research Analysts listed on the front cover, if applicable, they also certify, as per KOFIA requirements, that the Research Analyst’s analysis was made in good faith and that the views reflect the Research Analyst’s own opinion, without undue influence or intervention.


All authors named within this report are Research Analysts who produce independent research unless otherwise specified. In Europe, Sector Specialists (Sales and Trading) may be shown on this report as contacts but are not authors of the report or part of the Research Department.

Important Disclosures

- **Market Maker/ Liquidity Provider (Hong Kong):** J.P. Morgan Securities (Asia Pacific) Limited and/or J.P. Morgan Broking (Hong Kong) Limited and/or an affiliate is a market maker and/or liquidity provider in the securities of Zhipu AI, MiniMax Group Inc - H or related entities and/or warrants or options thereon, which are listed or traded on The Stock Exchange of Hong Kong Limited.
- **Client:** J.P. Morgan currently has, or had within the past 12 months, the following entity(ies) as clients: MiniMax Group Inc - H or related entities.
- **Client/Non-Investment Banking, Securities-Related:** J.P. Morgan currently has, or had within the past 12 months, the following entity(ies) as clients, and the services provided were non-investment-banking, securities-related: MiniMax Group Inc - H or related entities.
- **Potential Investment Banking Compensation:** J.P. Morgan expects to receive, or intends to seek, compensation for investment banking services in the next three months from MiniMax Group Inc - H or related entities.
- **Non-Investment Banking Compensation Received:** J.P. Morgan has received compensation in the past 12 months for products or services other than investment banking from MiniMax Group Inc - H or related entities.
- **Debt Position:** J.P. Morgan may hold a position in the debt securities of Zhipu AI, MiniMax Group Inc - H or related entities, if any.

Company-Specific Disclosures: J.P. Morgan does and seeks to do business with companies covered in its research reports. As a result, investors should be aware that the firm may have a conflict of interest that could affect the objectivity of this report. Investors should consider this report as only a single factor in making their investment decision. Important disclosures, including price charts and credit opinion history tables (if applicable), are available for compendium reports and all J.P. Morgan-covered companies, and certain non-covered companies, by visiting <https://www.jpmm.com/research/disclosures>, calling 1-800-477-0406, or e-mailing research.disclosure.inquiries@jpmorgan.com with your request.

Zhipu AI (2513.HK, 2513 HK) Price Chart



Date	Rating	Price (HK\$)	Price Target (HK\$)
09-Feb-26	OW	203.20	400
27-Feb-26	OW	556.50	800
01-Apr-26	OW	693.50	950
12-Jun-26	OW	1061.00	1,400
18-Jun-26	OW	1660.00	1,800

Source: Bloomberg Finance L.P. and J.P. Morgan; price data adjusted for stock splits and dividends.
Initiated coverage Feb 09, 2026. All share prices are as of market close on the previous business day.

MiniMax Group Inc - H (0100.HK, 100 HK) Price Chart



Date	Rating	Price (HK\$)	Price Target (HK\$)
09-Feb-26	OW	461.00	700
27-Feb-26	OW	788.50	1,000
03-Mar-26	OW	752.50	1,100
12-Jun-26	N	435.00	400

Source: Bloomberg Finance L.P. and J.P. Morgan; price data adjusted for stock splits and dividends.
Initiated coverage Feb 09, 2026. All share prices are as of market close on the previous business day.

The chart(s) show J.P. Morgan's continuing coverage of the stocks; the current analysts may or may not have covered it over the entire period.

J.P. Morgan ratings or designations: OW = Overweight, N= Neutral, UW = Underweight, NR = Not Rated

Explanation of Equity Research Ratings, Designations and Analyst(s) Coverage Universe:

J.P. Morgan uses the following rating system: Overweight (over the duration of the price target indicated in this report, we expect this stock will outperform the average total return of the stocks in the Research Analyst's, or the Research Analyst's team's, coverage universe); Neutral (over the duration of the price target indicated in this report, we expect this stock will perform in line with the average total return of the stocks in the Research Analyst's, or the Research Analyst's team's, coverage universe); and Underweight (over the duration of the price target indicated in this report, we expect this stock will underperform the average total return of the stocks in the Research Analyst's, or the Research Analyst's team's, coverage universe. NR is Not Rated. In this case, J.P. Morgan has removed the rating and, if applicable, the price target, for this stock because of either a lack of a sufficient fundamental basis or for legal, regulatory or policy reasons. The previous rating and, if applicable, the price target, no longer should be relied upon. An NR designation is not a recommendation or a rating. Some stocks under coverage have a rating but no price target; in these cases, we expect the stock will outperform/perform in line/underperform the average total return of the stocks in the Research Analyst's, or the Research Analyst's team's, coverage universe of the relevant duration of the region. In our Asia (ex-Australia and ex-India) and U.K. small- and mid-cap Equity Research, each stock's expected total return is compared to the expected total return of a benchmark country market index, not to those Research Analysts' coverage universe. If it does not appear in the Important Disclosures section of this report, the certifying Research Analyst's coverage universe can be found on J.P. Morgan's Research website, <https://www.jpmorganmarkets.com>.

Coverage Universe: Xu, Olivia : MiniMax Group Inc - H (0100.HK), Zhipu AI (2513.HK)

J.P. Morgan Equity Research Ratings Distribution, as of July 04, 2026

	Overweight (buy)	Neutral (hold)	Underweight (sell)
J.P. Morgan Global Equity Research Coverage*	53%	36%	12%
IB clients**	83%	80%	73%
JPMS Equity Research Coverage*	51%	37%	12%
IB clients**	95%	92%	87%

*Please note that the percentages may not add to 100% because of rounding.

**Percentage of subject companies within each of the "buy," "hold" and "sell" categories for which J.P. Morgan has provided investment banking services within the previous 12 months.

For purposes of FINRA ratings distribution rules only, our Overweight rating falls into a buy rating category; our Neutral rating falls into a hold rating category; and our Underweight rating falls into a sell rating category. Please note that stocks with an NR designation are not included in the table above. This information is current as of the end of the most recent calendar quarter.

Equity Valuation and Risks: For valuation methodology and risks associated with covered companies or price targets for covered companies, please see the most recent company-specific research report at <http://www.jpmorganmarkets.com>, contact the primary analyst

or your J.P. Morgan representative, or email research.disclosure.inquiries@jpmorgan.com. For material information about the proprietary models used, please see the Summary of Financials in company-specific research reports and the Company Tearsheets, which are available to download on the company pages of our client website, <http://www.jpmorganmarkets.com>. This report also sets out within it the material underlying assumptions used.

History of Investment Recommendations:

A history of J.P. Morgan investment recommendations disseminated during the preceding 12 months can be accessed on the Research & Commentary page of <http://www.jpmorganmarkets.com> where you can also search by analyst name, sector or financial instrument.

Analysts' Compensation: The research analysts responsible for the preparation of this report receive compensation based upon various factors, including the quality and accuracy of research, client feedback, competitive factors, and overall firm revenues.

Registration of non-US Analysts: Unless otherwise noted, the non-US analysts listed on the front of this report are employees of non-US affiliates of J.P. Morgan Securities LLC, may not be registered as research analysts under FINRA rules, may not be associated persons of J.P. Morgan Securities LLC, and may not be subject to FINRA Rule 2241 or 2242 restrictions on communications with covered companies, public appearances, and trading securities held by a research analyst account.

Other Disclosures

J.P. Morgan is a marketing name for investment banking businesses of JPMorgan Chase & Co. and its subsidiaries and affiliates worldwide.

UK MIFID FICC research unbundling exemption: UK clients should refer to [UK MIFID Research Unbundling exemption](#) for details of J.P. Morgan's implementation of the FICC research exemption and guidance on relevant FICC research categorisation.

All research material made available to clients are simultaneously available on our client website, J.P. Morgan Markets, unless specifically permitted by relevant laws. Not all research content is redistributed, e-mailed or made available to third-party aggregators. For all research material available on a particular stock, please contact your sales representative.

Any long form nomenclature for references to China; Hong Kong; Taiwan; and Macau within this research material are Mainland China; Hong Kong SAR (China); Taiwan (China); and Macau SAR (China).

J.P. Morgan Research may, from time to time, write on issuers or securities targeted by economic or financial sanctions imposed or administered by the governmental authorities of the U.S., EU, UK or other relevant jurisdictions (Sanctioned Securities). Nothing in this report is intended to be read or construed as encouraging, facilitating, promoting or otherwise approving investment or dealing in such Sanctioned Securities. Clients should be aware of their own legal and compliance obligations when making investment decisions.

Any digital or crypto assets discussed in this research report are subject to a rapidly changing regulatory landscape. For relevant regulatory advisories on crypto assets, including bitcoin and ether, please see <https://www.jpmorgan.com/disclosures/cryptoasset-disclosure>.

The author(s) of this research report may not be licensed to carry on regulated activities in your jurisdiction and, if not licensed, do not hold themselves out as being able to do so.

Exchange-Traded Funds (ETFs): J.P. Morgan Securities LLC ("JPMS") acts as authorized participant for substantially all U.S.-listed ETFs. To the extent that any ETFs are mentioned in this report, JPMS may earn commissions and transaction-based compensation in connection with the distribution of those ETF shares and may earn fees for performing other trade-related services, such as securities lending to short sellers of the ETF shares. JPMS may also perform services for the ETFs themselves, including acting as a broker or dealer to the ETFs. In addition, affiliates of JPMS may perform services for the ETFs, including trust, custodial, administration, lending, index calculation and/or maintenance and other services.

Options and Futures related research: If the information contained herein regards options- or futures-related research, such information is available only to persons who have received the proper options or futures risk disclosure documents. Please contact your J.P. Morgan Representative or visit <https://www.theocc.com/components/docs/riskstoc.pdf> for a copy of the Option Clearing Corporation's Characteristics and Risks of Standardized Options or https://www.finra.org/sites/default/files/2020-08/Security_Futures_Risk_Disclosure_Statement_2020.pdf for a copy of the Security Futures Risk Disclosure Statement.

Changes to Interbank Offered Rates (IBORs) and other benchmark rates: Certain interest rate benchmarks are, or may in the future become, subject to ongoing international, national and other regulatory guidance, reform and proposals for reform. For more information, please consult: https://www.jpmorgan.com/global/disclosures/interbank_offered_rates

Private Bank Clients: Where you are receiving research as a client of the private banking businesses offered by JPMorgan Chase & Co. and its subsidiaries ("J.P. Morgan Private Bank"), research is provided to you by J.P. Morgan Private Bank and not by any other division of J.P. Morgan, including, but not limited to, the J.P. Morgan Corporate and Investment Bank and its Global Research division.

Legal entity responsible for the production and distribution of research: The legal entity identified below the name of the Reg AC Research Analyst who authored this material is the legal entity responsible for the production of this research. Where multiple Reg AC Research Analysts authored this material with different legal entities identified below their names, these legal entities are jointly responsible for the production of this research. Where more than one legal entity is listed under an analyst's name, the first legal entity is responsible for the production unless stated otherwise. Research Analysts from various J.P. Morgan affiliates may have contributed to the production of this material but may not be licensed to carry out regulated activities in your jurisdiction (and do not hold themselves out as being able to do so). Unless otherwise stated below in the legal entity disclosures, this material has been distributed by the legal entity responsible for production, or where more than one legal entity is listed under the analyst's name, the first legal entity will be responsible for distribution. If you have any queries, please contact the relevant Research Analyst in your jurisdiction or the entity in your jurisdiction that has distributed this research material.

Legal Entities Disclosures and Country-/Region-Specific Disclosures:

Argentina: JPMorgan Chase Bank N.A Sucursal Buenos Aires is regulated by Banco Central de la República Argentina ("BCRA"- Central Bank of Argentina) and Comisión Nacional de Valores ("CNV"- Argentinian Securities Commission - ALYC y AN Integral N°51).

Australia: J.P. Morgan Securities Australia Limited ("JPMSAL") (ABN 61 003 245 234/AFS Licence No: 238066) is regulated by the Australian Securities and Investments Commission and is a Market Participant of ASX Limited, a Clearing and Settlement Participant of ASX Clear Pty Limited and a Clearing Participant of ASX Clear (Futures) Pty Limited. This material is issued and distributed in Australia by or on behalf of JPMSAL only to "wholesale clients" (as defined in section 761G of the Corporations Act 2001). A list of all financial products covered can be found by visiting <https://www.jpmm.com/research/disclosures>. J.P. Morgan seeks to cover companies of relevance to the domestic and international investor base across all Global Industry Classification Standard (GICS) sectors, as well as across a range of market capitalisation sizes. If applicable, in the course of conducting public side due diligence on the subject company(ies), the Research Analyst team may at times perform such diligence through corporate engagements such as site visits, discussions with company representatives, management presentations, etc. Research issued by JPMSAL has been prepared in accordance with J.P. Morgan Australia's Research Independence Policy which can be found at the following link: [J.P. Morgan Australia - Research Independence Policy](#).

Brazil: Banco J.P. Morgan S.A. is regulated by the Comissão de Valores Mobiliários (CVM) and by the Central Bank of Brazil. Ombudsman J.P. Morgan: 0800-7700847 / 0800-7700810 (For Hearing Impaired) / ouvidoria.jp.morgan@jpmchase.com.

Canada: J.P. Morgan Securities Canada Inc. is a registered investment dealer, regulated by the Canadian Investment Regulatory Organization and the Ontario Securities Commission and is the participating member on Canadian exchanges. This material is distributed in Canada by or on behalf of J.P.Morgan Securities Canada Inc.

Chile: Inversiones J.P. Morgan Limitada is an unregulated entity incorporated in Chile.

China: J.P. Morgan Securities (China) Company Limited has been approved by CSRC to conduct the securities investment consultancy business.

Colombia: Banco J.P. Morgan Colombia S.A. is supervised by the Superintendencia Financiera de Colombia (SFC). Any reference in this material to products or services offered abroad by entities other than the Bank in Colombia is included exclusively for descriptive purposes. Such references do not constitute, and should not be construed as, promotional activity or the provision of financial products or services within Colombian territory, as defined under applicable Colombian regulation.

Dubai International Financial Centre (DIFC): JPMorgan Chase Bank, N.A., Dubai Branch is regulated by the Dubai Financial Services Authority (DFSA) and its registered address is Dubai International Financial Centre - The Gate, West Wing, Level 3 and 9 PO Box 506551, Dubai, UAE. This material has been distributed by JP Morgan Chase Bank, N.A., Dubai Branch to persons regarded as professional clients or market counterparties as defined under the DFSA rules.

European Economic Area (EEA): Unless specified to the contrary, research is distributed in the EEA by J.P. Morgan SE ("JPM SE"), which is authorised as a credit institution by the Federal Financial Supervisory Authority (Bundesanstalt für Finanzdienstleistungsaufsicht, BaFin) and jointly supervised by the BaFin, the German Central Bank (Deutsche Bundesbank) and the European Central Bank (ECB). JPM SE is a company headquartered in Frankfurt with registered address at TaunusTurm, Taunustor 1, Frankfurt am Main, 60310, Germany. The material has been distributed in the EEA to persons regarded as professional investors (or equivalent) pursuant to Art. 4 para. 1 no. 10 and Annex II of MiFID II and its respective implementation in their home jurisdictions ("EEA professional investors"). This material must not be acted on or relied on by persons who are not EEA professional investors. Any investment or investment activity to which this material relates is only available to EEA relevant persons and will be engaged in only with EEA relevant persons.

Hong Kong: J.P. Morgan Securities (Asia Pacific) Limited (CE number AAJ321) is regulated by the Hong Kong Monetary Authority and the Securities and Futures Commission in Hong Kong, and J.P. Morgan Broking (Hong Kong) Limited (CE number AAB027) is

regulated by the Securities and Futures Commission in Hong Kong. JP Morgan Chase Bank, N.A., Hong Kong Branch (CE Number AAL996) is regulated by the Hong Kong Monetary Authority and the Securities and Futures Commission, is organized under the laws of the United States with limited liability. Where the distribution of this material is a regulated activity in Hong Kong, the material is distributed in Hong Kong by or through J.P. Morgan Securities (Asia Pacific) Limited and/or J.P. Morgan Broking (Hong Kong) Limited.

India: J.P. Morgan India Private Limited (Corporate Identity Number - U67120MH1992FTC068724), having its registered office at J.P. Morgan Tower, Off. C.S.T. Road, Kalina, Santacruz - East, Mumbai – 400098, is registered with the Securities and Exchange Board of India (SEBI) as a ‘Research Analyst’ having registration number INH000001873. J.P. Morgan India Private Limited is also registered with SEBI as a member of the National Stock Exchange of India Limited and the Bombay Stock Exchange Limited (SEBI Registration Number – INZ000239730) and as a Merchant Banker (SEBI Registration Number - MB/INM000002970). Telephone: 91-22-6157 3000, Facsimile: 91-22-6157 3990 and Website: <http://www.jpmmipl.com>. JPMorgan Chase Bank, N.A. - Mumbai Branch is licensed by the Reserve Bank of India (RBI) (Licence No. 53/ Licence No. BY.4/94; SEBI - IN/CUS/014/ CDSL : IN-DP-CDSL-444-2008/ IN-DP-NSDL-285-2008/ INBI00000984/ INE231311239) as a Scheduled Commercial Bank in India, which is its primary license allowing it to carry on Banking business in India and other activities, which a Bank branch in India are permitted to undertake. For non-local research material, this material is not distributed in India by J.P. Morgan India Private Limited. Compliance Officer: Ashutosh Sharma; ashutosh.j.sharma@jpmchase.com; +912261575002. Grievance Officer: Ramprasad K, jpmipl.research.feedback@jpmorgan.com; +912261573000. Registration granted by SEBI and certification from NISM in no way guarantee performance of the intermediary or provide any assurance of returns to investors. Please visit [Terms and Conditions and Most Important Terms and Conditions \(MITC\)](#). The annual Compliance audit report is available at <http://www.jpmmipl.com/#research>.

Indonesia: PT J.P. Morgan Sekuritas Indonesia is a member of the Indonesia Stock Exchange and is registered and supervised by the Otoritas Jasa Keuangan (OJK).

Korea: J.P. Morgan Securities (Far East) Limited, Seoul Branch, is a member of the Korea Exchange (KRX). JPMorgan Chase Bank, N.A., Seoul Branch, is licensed as a branch office of foreign bank (JPMorgan Chase Bank, N.A.) in Korea. Both entities are regulated by the Financial Services Commission (FSC) and the Financial Supervisory Service (FSS). For non-macro research material, the material is distributed in Korea by or through J.P. Morgan Securities (Far East) Limited, Seoul Branch.

Japan: JPMorgan Securities Japan Co., Ltd. and JPMorgan Chase Bank, N.A., Tokyo Branch are regulated by the Financial Services Agency in Japan.

Malaysia: This material is issued and distributed in Malaysia by JPMorgan Securities (Malaysia) Sdn Bhd (18146-X), which is a Participating Organization of Bursa Malaysia Berhad and holds a Capital Markets Services License issued by the Securities Commission in Malaysia.

Mexico: J.P. Morgan Casa de Bolsa, S.A. de C.V., J.P. Morgan Grupo Financiero is member of the Mexican Stock Exchange (“Bolsa Mexicana de Valores”) and the Institutional Stock Exchange (“Bolsa Institucional de Valores”), and it is authorized to act as a broker dealer by the National Banking and Securities Exchange Commission (“Comisión Nacional Bancaria y de Valores”).

New Zealand: This material is issued and distributed by JPMSAL in New Zealand only to “wholesale clients” (as defined in the Financial Markets Conduct Act 2013). JPMSAL is registered as a Financial Service Provider under the Financial Service providers (Registration and Dispute Resolution) Act of 2008.

Philippines: J.P. Morgan Securities Philippines Inc. is a Trading Participant of the Philippine Stock Exchange and a member of the Securities Clearing Corporation of the Philippines and the Securities Investor Protection Fund. It is regulated by the Securities and Exchange Commission.

Singapore: This material is issued and distributed in Singapore by or through J.P. Morgan Securities Singapore Private Limited (JPMSS) [MDDI (P) 057/08/2025 and Co. Reg. No.: 199405335R], which is a member of the Singapore Exchange Securities Trading Limited, and/or JPMorgan Chase Bank, N.A., Singapore branch (JPMCB Singapore), both of which are regulated by the Monetary Authority of Singapore. This material is issued and distributed in Singapore only to accredited investors, expert investors and institutional investors, as defined in Section 4A of the Securities and Futures Act, Cap. 289 (SFA). This material is not intended to be issued or distributed to any retail investors or any other investors that do not fall into the classes of “accredited investors,” “expert investors” or “institutional investors,” as defined under Section 4A of the SFA. Recipients of this material in Singapore are to contact JPMSS or JPMCB Singapore in respect of any matters arising from, or in connection with, the material.

South Africa: J.P. Morgan Equities South Africa Proprietary Limited and JPMorgan Chase Bank, N.A., Johannesburg Branch are members of the Johannesburg Securities Exchange and are regulated by the Financial Services Conduct Authority (FSCA).

Taiwan: J.P. Morgan Securities (Taiwan) Limited is a participant of the Taiwan Stock Exchange (company-type) and regulated by the Taiwan Securities and Futures Bureau. Material relating to equity securities is issued and distributed in Taiwan by J.P. Morgan Securities (Taiwan) Limited, subject to the license scope and the applicable laws and the regulations in Taiwan. **To the extent that J.P. Morgan Securities (Taiwan) Limited produces research materials on securities not listed on the Taiwan Stock Exchange or Taipei Exchange (“Non-Taiwan Listed Securities”), these materials shall not constitute securities recommendations for the purpose of applicable Taiwan regulations, and, for the avoidance of doubt, J.P. Morgan Securities (Taiwan) Limited does not act as broker for Non-Taiwan Listed Securities.** According to Paragraph 2, Article 7-1 of Operational Regulations Governing Securities Firms Recommending Trades in Securities to Customers (as amended or supplemented) and/or other applicable laws or regulations, please note that the recipient of this material is not permitted to engage in any activities in connection with the material that may give rise to conflicts of interests, unless otherwise disclosed in the “Important Disclosures” in this material.

Thailand: This material is issued and distributed in Thailand by JPMorgan Securities (Thailand) Ltd., which is a member of the Stock Exchange of Thailand and is regulated by the Ministry of Finance and the Securities and Exchange Commission. The registered address is 548 One City Center Building, 50th Floor, Ploenchit Road, Lumpini, Pathum Wan, Bangkok 10330.

UK: Research is produced in the UK by J.P. Morgan Securities plc ("JPMS plc") which is a member of the London Stock Exchange and is authorised by the Prudential Regulation Authority and regulated by the Financial Conduct Authority and the Prudential Regulation Authority or J.P. Morgan Markets Limited ("JPMML Ltd") which is authorised and regulated by the Financial Conduct Authority. Unless specified to the contrary, this material is distributed in the UK by JPMS plc and is directed in the UK only to: (a) persons having professional experience in matters relating to investments falling within article 19(5) of the Financial Services and Markets Act 2000 (Financial Promotion) (Order) 2005 ("the FPO"); (b) persons outlined in article 49 of the FPO (high net worth companies, unincorporated associations or partnerships, the trustees of high value trusts, etc.); or (c) any persons to whom this communication may otherwise lawfully be made; all such persons being referred to as "UK relevant persons". This material must not be acted on or relied on by persons who are not UK relevant persons. Any investment or investment activity to which this material relates is only available to UK relevant persons and will be engaged in only with UK relevant persons. A description of J.P. Morgan EMEA's policy for prevention and avoidance of conflicts of interest related to the production of Research can be found at the following link: [J.P. Morgan EMEA - Research Independence Policy](#).

U.S.: J.P. Morgan Securities LLC ("JPMS") is a member of the NYSE, FINRA, SIPC, and the NFA. JPMorgan Chase Bank, N.A. is a member of the FDIC. Material published by non-U.S. affiliates is distributed in the U.S. by JPMS who accepts responsibility for its content.

General: Additional information is available upon request. The information in this material has been obtained from sources believed to be reliable. While all reasonable care has been taken to ensure that the facts stated in this material are accurate and that the forecasts, opinions and expectations contained herein are fair and reasonable, JPMorgan Chase & Co. or its affiliates and/or subsidiaries (collectively J.P. Morgan) make no representations or warranties whatsoever to the completeness or accuracy of the material provided, except with respect to any disclosures relative to J.P. Morgan and the Research Analyst's involvement with the issuer that is the subject of the material. Accordingly, no reliance should be placed on the accuracy, fairness or completeness of the information contained in this material. There may be certain discrepancies with data and/or limited content in this material as a result of calculations, adjustments, translations to different languages, and/or local regulatory restrictions, as applicable. These discrepancies should not impact the overall investment analysis, views and/or recommendations of the subject company(ies) that may be discussed in the material. Artificial intelligence tools may have been used in the preparation of this material, including assisting in data analysis, pattern recognition, and content drafting for research material. J.P. Morgan accepts no liability whatsoever for any loss arising from any use of this material or its contents, and neither J.P. Morgan nor any of its respective directors, officers or employees, shall be in any way responsible for the contents hereof, apart from the liabilities and responsibilities that may be imposed on them by the relevant regulatory authority in the jurisdiction in question, or the regulatory regime thereunder. Opinions, forecasts or projections contained in this material represent J.P. Morgan's current opinions or judgment as of the date of the material only and are therefore subject to change without notice. Periodic updates may be provided on companies/industries based on company-specific developments or announcements, market conditions or any other publicly available information. There can be no assurance that future results or events will be consistent with any such opinions, forecasts or projections, which represent only one possible outcome. Furthermore, such opinions, forecasts or projections are subject to certain risks, uncertainties and assumptions that have not been verified, and future actual results or events could differ materially. The value of, or income from, any investments referred to in this material may fluctuate and/or be affected by changes in exchange rates. All pricing is indicative as of the close of market for the securities discussed, unless otherwise stated. Past performance is not indicative of future results. Accordingly, investors may receive back less than originally invested. This material is not intended as an offer or solicitation for the purchase or sale of any financial instrument. The opinions and recommendations herein do not take into account individual client circumstances, objectives, or needs and are not intended as recommendations of particular securities, financial instruments or strategies to particular clients. This material may include views on structured securities, options, futures and other derivatives. These are complex instruments, may involve a high degree of risk and may be appropriate investments only for sophisticated investors who are capable of understanding and assuming the risks involved. The recipients of this material must make their own independent decisions regarding any securities or financial instruments mentioned herein and should seek advice from such independent financial, legal, tax or other adviser as they deem necessary. J.P. Morgan may trade as a principal on the basis of the Research Analysts' views and research, and it may also engage in transactions for its own account or for its clients' accounts in a manner inconsistent with the views taken in this material, and J.P. Morgan is under no obligation to ensure that such other communication is brought to the attention of any recipient of this material. Others within J.P. Morgan, including Strategists, Sales staff and other Research Analysts, may take views that are inconsistent with those taken in this material. Employees of J.P. Morgan not involved in the preparation of this material may have investments in the securities (or derivatives of such securities) mentioned in this material and may trade them in ways different from those discussed in this material. This material is not an advertisement for or marketing of any issuer, its products or services, or its securities in any jurisdiction.

Confidentiality and Security Notice: This transmission may contain information that is privileged, confidential, legally privileged, and/or exempt from disclosure under applicable law. If you are not the intended recipient, you are hereby notified that any disclosure, copying, distribution, or use of the information contained herein (including any reliance thereon) is STRICTLY PROHIBITED. Although this transmission and any attachments are believed to be free of any virus or other defect that might affect any computer system into

which it is received and opened, it is the responsibility of the recipient to ensure that it is virus free and no responsibility is accepted by JPMorgan Chase & Co., its subsidiaries and affiliates, as applicable, for any loss or damage arising in any way from its use. If you received this transmission in error, please immediately contact the sender and destroy the material in its entirety, whether in electronic or hard copy format. This message is subject to electronic monitoring: <https://www.jpmorgan.com/disclosures/email>

MSCI: Certain information herein ("Information") is reproduced by permission of MSCI Inc., its affiliates and information providers ("MSCI") ©2026. No reproduction or dissemination of the Information is permitted without an appropriate license. MSCI MAKES NO EXPRESS OR IMPLIED WARRANTIES (INCLUDING MERCHANTABILITY OR FITNESS) AS TO THE INFORMATION AND DISCLAIMS ALL LIABILITY TO THE EXTENT PERMITTED BY LAW. No Information constitutes investment advice, except for any applicable Information from MSCI ESG Research. Subject also to [msci.com/disclaimer](https://www.msci.com/disclaimer)

Sustainalytics: Certain information, data, analyses and opinions contained herein are reproduced by permission of Sustainalytics and: (1) includes the proprietary information of Sustainalytics; (2) may not be copied or redistributed except as specifically authorized; (3) do not constitute investment advice nor an endorsement of any product or project; (4) are provided solely for informational purposes; and (5) are not warranted to be complete, accurate or timely. Sustainalytics is not responsible for any trading decisions, damages or other losses related to it or its use. The use of the data is subject to conditions available at <https://www.sustainalytics.com/legal-disclaimers>. ©2026 Sustainalytics. All Rights Reserved.

"Other Disclosures" last revised July 04, 2026.

Copyright 2026 J.P. Morgan Securities (China) Company Limited. All rights reserved. This material or any portion hereof may not be reprinted, sold or redistributed without the written consent of J.P. Morgan. It is strictly prohibited to use or share without prior written consent from J.P. Morgan any research material received from J.P. Morgan or an authorized third-party ("J.P. Morgan Data") in any third-party artificial intelligence ("AI") systems or models when such J.P. Morgan Data is accessible by a third-party.